

tf.contrib.bayesflow.csiszar_divergence.monte_carlo_csiszar_f_divergence

```
monte_carlo_csiszar_f_divergence(
    f,
    p_log_prob,
    q,
    num_draws,
    use_reparameterization=None,
    seed=None,
    name=None
)
```

Defined in [tensorflow/contrib/bayesflow/python/ops/csiszar_divergence_impl.py](#).

Monte-Carlo approximation of the Csiszar f-Divergence.

A Csiszar-function is a member of,

$$F = \{ f: \mathbb{R}_+ \rightarrow \mathbb{R} : f \text{ convex} \}.$$

The Csiszar f-Divergence for Csiszar-function f is given by:

$$\begin{aligned} D_f[p(X), q(X)] &:= E_{\{q(X)\}}[f(p(X)/q(X))] \\ &\approx m^{-1} \sum_j f(p(x_j)/q(x_j)), \\ &\quad \text{where } x_j \sim \text{iid } q(X) \end{aligned}$$

Tricks: Reparameterization and Score-Gradient

When q is "reparameterized", i.e., a diffeomorphic transformation of a parameterless distribution (e.g., $\text{Normal}(Y; \mathbf{m}, \mathbf{s}) \Leftrightarrow Y = \mathbf{s}X + \mathbf{m}, X \sim \text{Normal}(\mathbf{0}, \mathbf{1})$), we can swap gradient and expectation, i.e., $\text{grad}[\text{Avg}\{s_i : i=1\dots n\}] = \text{Avg}\{\text{grad}[s_i] : i=1\dots n\}$ where $S_n = \text{Avg}\{s_i\}$ and $s_i = f(x_i), x_i \sim \text{iid } q(X)$.

However, if q is not reparameterized, TensorFlow's gradient will be incorrect since the chain-rule stops at samples of unreparameterized distributions. In this circumstance using the Score-Gradient trick results in an unbiased gradient, i.e.,

$$\begin{aligned} &\text{grad}[E_q[f(X)]] \\ &= \text{grad}[\int dx q(x) f(x)] \\ &= \int dx \text{grad}[q(x) f(x)] \\ &= \int dx [q'(x) f(x) + q(x) f'(x)] \\ &= \int dx q(x) [q'(x)/q(x) f(x) + f'(x)] \\ &= \int dx q(x) \text{grad}[f(x) q(x) / \text{stop_grad}[q(x)]] \\ &= E_q[\text{grad}[f(x) q(x) / \text{stop_grad}[q(x)]]] \end{aligned}$$

Unless `q.reparameterization_type != distribution.FULLY_REPARAMETERIZED` it is usually preferable to set `use_reparameterization = True`.

Example Application:

The Csiszar f-Divergence is a useful framework for variational inference. I.e., observe that,

$$\begin{aligned} f(p(x)) &= f(E_{\{q(Z|x)\}}[p(x, Z)/q(Z|x)]) \\ &\leq E_{\{q(Z|x)\}}[f(p(x, Z)/q(Z|x))] \\ &:= D_f[p(x, Z), q(Z|x)] \end{aligned}$$

The inequality follows from the fact that the "perspective" of f , i.e., $(s, t) \mapsto t f(s / t)$, is convex in (s, t) when $s/t \in \text{domain}(f)$ and t is a real. Since the above framework includes the popular Evidence Lower Bound (ELBO) as a special case, i.e., $f(u) = -\log(u)$, we call this framework "Evidence Divergence Bound Optimization" (EDBO).

Args:

- `f`: Python `callable` representing a Csiszar-function in log-space, i.e., takes `p_log_prob(q_samples) - q.log_prob(q_samples)`.
- `p_log_prob`: Python `callable` taking (a batch of) samples from `q` and returning the natural-log of the probability under distribution `p`. (In variational inference `p` is the joint distribution.)
- `q`: `tf.Distribution`-like instance; must implement: `reparameterization_type`, `sample(n, seed)`, and `log_prob(x)`. (In variational inference `q` is the approximate posterior distribution.)
- `num_draws`: Integer scalar number of draws used to approximate the f-Divergence expectation.
- `use_reparametrization`: Python `bool`. When `None` (the default), automatically set to: `q.reparameterization_type == distribution.FULLY_REPARAMETERIZED`. When `True` uses the standard Monte-Carlo average. When `False` uses the score-gradient trick. (See above for details.) When `False`, consider using `csiszar_vimco`.
- `seed`: Python `int` seed for `q.sample`.
- `name`: Python `str` name prefixed to Ops created by this function.

Returns:

- `monte_carlo_csiszar_f_divergence`: `float`-like `Tensor` Monte Carlo approximation of the Csiszar f-Divergence.

Raises:

- `ValueError`: if `q` is not a reparameterized distribution and `use_reparametrization = True`. A distribution `q` is said to be "reparameterized" when its samples are generated by transforming the samples of another distribution which does not depend on the parameterization of `q`. This property ensures the gradient (with respect to parameters) is valid.
- `TypeError`: if `p_log_prob` is not a Python `callable`.

Except as otherwise noted, the content of this page is licensed under the [Creative Commons Attribution 3.0 License](#), and code samples are licensed under the [Apache 2.0 License](#). For details, see our [Site Policies](#). Java is a registered trademark of Oracle and/or its affiliates.

Last updated November 2, 2017.

Stay Connected

[Blog](#)

[GitHub](#)

[Twitter](#)

Support

[Issue Tracker](#)

[Release Notes](#)

[Stack Overflow](#)

