

tf.contrib.distributions.Binomial

Contents

Class Binomial

Properties

allow_nan_stats

batch_shape

Class **Binomial**Inherits From: [Distribution](#)Defined in [tensorflow/contrib/distributions/python/ops/binomial.py](#).See the guide: [Statistical Distributions \(contrib\) > Univariate \(scalar\) distributions](#)

Binomial distribution.

This distribution is parameterized by `probs`, a (batch of) probabilities for drawing a `1` and `total_count`, the number of trials per draw from the Binomial.

Mathematical Details

The Binomial is a distribution over the number of `1`'s in `total_count` independent trials, with each trial having the same probability of `1`, i.e., `probs`.

The probability mass function (pmf) is,

$$\text{pmf}(k; n, p) = p^k (1 - p)^{(n - k)} / Z$$
$$Z = k! (n - k)! / n!$$

where: `total_count = n`, `probs = p`, `Z` is the normalizing constant, and, `n!` is the factorial of `n`.

Examples

Create a single distribution, corresponding to 5 coin flips.

```
dist = Binomial(total_count=5., probs=.5)
```

Create a single distribution (using logits), corresponding to 5 coin flips.

```
dist = Binomial(total_count=5., logits=0.)
```

Creates 3 distributions with the third distribution most likely to have successes.

```
p = [.2, .3, .8]
# n will be broadcast to [4., 4., 4.], to match p.
dist = Binomial(total_count=4., probs=p)
```

The distribution functions can be evaluated on counts.

```
# counts same shape as p.
counts = [1., 2, 3]
dist.prob(counts) # Shape [3]

# p will be broadcast to [[.2, .3, .8], [.2, .3, .8]] to match counts.
counts = [[1., 2, 1], [2, 2, 4]]
dist.prob(counts) # Shape [2, 3]

# p will be broadcast to shape [5, 7, 3] to match counts.
counts = [...] # Shape [5, 7, 3]
dist.prob(counts) # Shape [5, 7, 3]
```

Properties

allow_nan_stats

Python `bool` describing behavior when a stat is undefined.

Stats return +/- infinity when it makes sense. E.g., the variance of a Cauchy distribution is infinity. However, sometimes the statistic is undefined, e.g., if a distribution's pdf does not achieve a maximum within the support of the distribution, the mode is undefined. If the mean is undefined, then by definition the variance is undefined. E.g. the mean for Student's T for $df = 1$ is undefined (no clear way to say it is either + or - infinity), so the variance = $E[(X - \text{mean})^2]$ is also undefined.

Returns:

- `allow_nan_stats`: Python `bool`.

batch_shape

Shape of a single sample from a single event index as a `TensorShape`.

May be partially defined or unknown.

The batch dimensions are indexes into independent, non-identical parameterizations of this distribution.

Returns:

- `batch_shape`: `TensorShape`, possibly unknown.

dtype

The `DType` of `Tensor`s handled by this `Distribution`.

event_shape

Shape of a single sample from a single batch as a `TensorShape`.

May be partially defined or unknown.

Returns:

- `event_shape`: `TensorShape`, possibly unknown.

logits

Log-odds of drawing a `1`.

name

Name prepended to all ops created by this `Distribution`.

parameters

Dictionary of parameters used to instantiate this `Distribution`.

probs

Probability of drawing a `1`.

reparameterization_type

Describes how samples from the distribution are reparameterized.

Currently this is one of the static instances `distributions.FULLY_REPARAMETERIZED` or `distributions.NOT_REPARAMETERIZED`.

Returns:

An instance of `ReparameterizationType`.

total_count

Number of trials.

validate_args

Python `bool` indicating possibly expensive checks are enabled.

Methods

`__init__`

```
__init__(
    total_count,
    logits=None,
    probs=None,
    validate_args=False,
    allow_nan_stats=True,
    name='Binomial'
)
```

Initialize a batch of Binomial distributions.

Args:

- `total_count` : Non-negative floating point tensor with shape broadcastable to `[N1, ..., Nm]` with `m >= 0` and the same dtype as `probs` or `logits` . Defines this as a batch of `N1 x ... x Nm` different Binomial distributions. Its components should be equal to integer values.
- `logits` : Floating point tensor representing the log-odds of a positive event with shape broadcastable to `[N1, ..., Nm]` `m >= 0` , and the same dtype as `total_count` . Each entry represents logits for the probability of success for independent Binomial distributions. Only one of `logits` or `probs` should be passed in.
- `probs` : Positive floating point tensor with shape broadcastable to `[N1, ..., Nm]` `m >= 0` , `probs in [0, 1]` . Each entry represents the probability of success for independent Binomial distributions. Only one of `logits` or `probs` should be passed in.
- `validate_args` : Python `bool` , default `False` . When `True` distribution parameters are checked for validity despite possibly degrading runtime performance. When `False` invalid inputs may silently render incorrect outputs.
- `allow_nan_stats` : Python `bool` , default `True` . When `True` , statistics (e.g., mean, mode, variance) use the value "`NaN`" to indicate the result is undefined. When `False` , an exception is raised if one or more of the statistic's batch members are undefined.
- `name` : Python `str` name prefixed to Ops created by this class.

batch_shape_tensor

```
batch_shape_tensor(name='batch_shape_tensor')
```

Shape of a single sample from a single event index as a 1-D `Tensor` .

The batch dimensions are indexes into independent, non-identical parameterizations of this distribution.

Args:

- `name` : name to give to the op

Returns:

- `batch_shape` : `Tensor` .

cdf

```
cdf(
    value,
    name='cdf'
)
```

Cumulative distribution function.

Given random variable `X` , the cumulative distribution function `cdf` is:

```
cdf(x) := P[X <= x]
```

Args:

- `value` : `float` or `double Tensor` .
- `name` : The name to give this op.

Returns:

- `cdf`: a `Tensor` of shape `sample_shape(x) + self.batch_shape` with values of type `self.dtype`.

copy

```
copy(**override_parameters_kwargs)
```

Creates a deep copy of the distribution.

★ **Note:** the copy distribution may continue to depend on the original initialization arguments.

Args:

- `**override_parameters_kwargs`: String/value dictionary of initialization arguments to override with new values.

Returns:

- `distribution`: A new instance of `type(self)` initialized from the union of `self.parameters` and `override_parameters_kwargs`, i.e., `dict(self.parameters, **override_parameters_kwargs)`.

covariance

```
covariance(name='covariance')
```

Covariance.

Covariance is (possibly) defined only for non-scalar-event distributions.

For example, for a length-`k`, vector-valued distribution, it is calculated as,

$$\text{Cov}[i, j] = \text{Covariance}(X_i, X_j) = E[(X_i - E[X_i]) (X_j - E[X_j])]$$

where `Cov` is a (batch of) `k x k` matrix, $0 \leq (i, j) < k$, and `E` denotes expectation.

Alternatively, for non-vector, multivariate distributions (e.g., matrix-valued, Wishart), `Covariance` shall return a (batch of) matrices under some vectorization of the events, i.e.,

$$\text{Cov}[i, j] = \text{Covariance}(\text{Vec}(X)_i, \text{Vec}(X)_j) = [\text{as above}]$$

where `Cov` is a (batch of) `k' x k'` matrices, $0 \leq (i, j) < k' = \text{reduce_prod}(\text{event_shape})$, and `Vec` is some function mapping indices of this distribution's event dimensions to indices of a length-`k'` vector.

Args:

- `name`: The name to give this op.

Returns:

- `covariance`: Floating-point `Tensor` with shape `[B1, ..., Bn, k', k']` where the first `n` dimensions are batch coordinates and `k' = reduce_prod(self.event_shape)`.

entropy

```
entropy(name='entropy')
```

Shannon entropy in nats.

event_shape_tensor

```
event_shape_tensor(name='event_shape_tensor')
```

Shape of a single sample from a single batch as a 1-D int32 **Tensor** .

Args:

- `name` : name to give to the op

Returns:

- `event_shape` : **Tensor** .

is_scalar_batch

```
is_scalar_batch(name='is_scalar_batch')
```

Indicates that `batch_shape == []` .

Args:

- `name` : The name to give this op.

Returns:

- `is_scalar_batch` : **bool** scalar **Tensor** .

is_scalar_event

```
is_scalar_event(name='is_scalar_event')
```

Indicates that `event_shape == []` .

Args:

- `name` : The name to give this op.

Returns:

- `is_scalar_event` : **bool** scalar **Tensor** .

log_cdf

```
log_cdf(
    value,
    name='log_cdf'
)
```

Log cumulative distribution function.

Given random variable **X**, the cumulative distribution function **cdf** is:

$$\text{log_cdf}(x) := \text{Log}[P[X \leq x]]$$

Often, a numerical approximation can be used for **log_cdf(x)** that yields a more accurate answer than simply taking the logarithm of the **cdf** when $x \ll -1$.

Args:

- **value**: **float** or **double Tensor**.
- **name**: The name to give this op.

Returns:

- **logcdf**: a **Tensor** of shape **sample_shape(x) + self.batch_shape** with values of type **self.dtype**.

log_prob

```
log_prob(
    value,
    name='log_prob'
)
```

Log probability density/mass function.

Additional documentation from **Binomial**:

For each batch member of counts **value**, **P[value]** is the probability that after sampling **self.total_count** draws from this Binomial distribution, the number of successes is **value**. Since different sequences of draws can result in the same counts, the probability includes a combinatorial coefficient.

★ **Note:** **value** must be a non-negative tensor with dtype **dtype** and whose shape can be broadcast with **self.probs** and **self.total_count**. **value** is only legal if it is less than or equal to **self.total_count** and its components are equal to integer values.

Args:

- **value**: **float** or **double Tensor**.
- **name**: The name to give this op.

Returns:

- **log_prob**: a **Tensor** of shape **sample_shape(x) + self.batch_shape** with values of type **self.dtype**.

log_survival_function

```
log_survival_function(
    value,
    name='log_survival_function'
)
```

Log survival function.

Given random variable X , the survival function is defined:

```
log_survival_function(x) = Log[ P[X > x] ]
                        = Log[ 1 - P[X <= x] ]
                        = Log[ 1 - cdf(x) ]
```

Typically, different numerical approximations can be used for the log survival function, which are more accurate than $1 - \text{cdf}(x)$ when $x \gg 1$.

Args:

- `value`: `float` or `double Tensor`.
- `name`: The name to give this op.

Returns:

`Tensor` of shape `sample_shape(x) + self.batch_shape` with values of type `self.dtype`.

mean

```
mean(name='mean')
```

Mean.

mode

```
mode(name='mode')
```

Mode.

Additional documentation from `Binomial`:

Note that when $(1 + \text{total_count}) * \text{probs}$ is an integer, there are actually two modes. Namely, $(1 + \text{total_count}) * \text{probs}$ and $(1 + \text{total_count}) * \text{probs} - 1$ are both modes. Here we return only the larger of the two modes.

param_shapes

```
param_shapes(
    cls,
    sample_shape,
    name='DistributionParamShapes'
)
```

Shapes of parameters given the desired shape of a call to `sample()`.

This is a class method that describes what key/value arguments are required to instantiate the given `Distribution` so that a particular shape is returned for that instance's call to `sample()`.

Subclasses should override class method `_param_shapes` .

Args:

- `sample_shape` : `Tensor` or python list/tuple. Desired shape of a call to `sample()` .
- `name` : name to prepend ops with.

Returns:

`dict` of parameter name to `Tensor` shapes.

`param_static_shapes`

```
param_static_shapes(  
    cls,  
    sample_shape  
)
```

`param_shapes` with static (i.e. `TensorShape`) shapes.

This is a class method that describes what key/value arguments are required to instantiate the given `Distribution` so that a particular shape is returned for that instance's call to `sample()` . Assumes that the sample's shape is known statically.

Subclasses should override class method `_param_shapes` to return constant-valued tensors when constant values are fed.

Args:

- `sample_shape` : `TensorShape` or python list/tuple. Desired shape of a call to `sample()` .

Returns:

`dict` of parameter name to `TensorShape` .

Raises:

- `ValueError` : if `sample_shape` is a `TensorShape` and is not fully defined.

`prob`

```
prob(  
    value,  
    name='prob'  
)
```

Probability density/mass function.

Additional documentation from `Binomial` :

For each batch member of counts `value` , `P[value]` is the probability that after sampling `self.total_count` draws from this Binomial distribution, the number of successes is `value` . Since different sequences of draws can result in the same counts, the probability includes a combinatorial coefficient.

★ **Note:** `value` must be a non-negative tensor with dtype `dtype` and whose shape can be broadcast with `self.probs` and `self.total_count`. `value` is only legal if it is less than or equal to `self.total_count` and its components are equal to integer values.

Args:

- `value`: `float` or `double Tensor`.
- `name`: The name to give this op.

Returns:

- `prob`: a `Tensor` of shape `sample_shape(x) + self.batch_shape` with values of type `self.dtype`.

quantile

```
quantile(  
    value,  
    name='quantile'  
)
```

Quantile function. Aka "inverse cdf" or "percent point function".

Given random variable `X` and `p` in `[0, 1]`, the `quantile` is:

```
quantile(p) := x such that P[X <= x] == p
```

Args:

- `value`: `float` or `double Tensor`.
- `name`: The name to give this op.

Returns:

- `quantile`: a `Tensor` of shape `sample_shape(x) + self.batch_shape` with values of type `self.dtype`.

sample

```
sample(  
    sample_shape=(),  
    seed=None,  
    name='sample'  
)
```

Generate samples of the specified shape.

Note that a call to `sample()` without arguments will generate a single sample.

Args:

- `sample_shape`: 0D or 1D `int32 Tensor`. Shape of the generated samples.
- `seed`: Python integer seed for RNG

- `name` : name to give to the op.

Returns:

- `samples` : a `Tensor` with prepended dimensions `sample_shape` .

stddev

```
stddev(name='stddev')
```

Standard deviation.

Standard deviation is defined as,

$$\text{stddev} = E[(X - E[X])**2]**0.5$$

where `X` is the random variable associated with this distribution, `E` denotes expectation, and `stddev.shape = batch_shape + event_shape` .

Args:

- `name` : The name to give this op.

Returns:

- `stddev` : Floating-point `Tensor` with shape identical to `batch_shape + event_shape` , i.e., the same shape as `self.mean()` .

survival_function

```
survival_function(
    value,
    name='survival_function'
)
```

Survival function.

Given random variable `X` , the survival function is defined:

$$\begin{aligned} \text{survival_function}(x) &= P[X > x] \\ &= 1 - P[X \leq x] \\ &= 1 - \text{cdf}(x). \end{aligned}$$

Args:

- `value` : `float` or `double Tensor` .
- `name` : The name to give this op.

Returns:

`Tensor` of shape `sample_shape(x) + self.batch_shape` with values of type `self.dtype` .

variance

```
variance(name='variance')
```

Variance.

Variance is defined as,

$$\text{Var} = E[(X - E[X])**2]$$

where **X** is the random variable associated with this distribution, **E** denotes expectation, and **Var.shape = batch_shape + event_shape**.

Args:

- **name**: The name to give this op.

Returns:

- **variance**: Floating-point **Tensor** with shape identical to **batch_shape + event_shape**, i.e., the same shape as **self.mean()**.

Except as otherwise noted, the content of this page is licensed under the [Creative Commons Attribution 3.0 License](#), and code samples are licensed under the [Apache 2.0 License](#). For details, see our [Site Policies](#). Java is a registered trademark of Oracle and/or its affiliates.

Last updated November 2, 2017.

Stay Connected

[Blog](#)

[GitHub](#)

[Twitter](#)

Support

[Issue Tracker](#)

[Release Notes](#)

[Stack Overflow](#)

English

[Terms](#) | [Privacy](#)