tf.contrib.distributions.estimator_head_distribution_regression

```
estimator_head_distribution_regression(
    make_distribution_fn,
    label_dimension=1,
    logits_dimension=None,
    label_name=None,
    weight_column_name=None,
    enable_centered_bias=False,
    head_name=None
)
```

Defined in tensorflow/contrib/distributions/python/ops/estimator.py .

Creates a `Head` for regression under a generic distribution.

Args:

- `make_distribution_fn` : Python `callable` which returns a `tf.Distribution` instance created using only logits.
- `label_dimension` : Number of regression labels per example. This is the size of the last dimension of the labels `Tensor` (typically, this has shape `[batch_size, label_dimension]`).
- `logits_dimension` : Number of logits per example. This is the size of the last dimension of the logits `Tensor` (typically, this has shape `[batch_size, logits_dimension]`). Default value: `label_dimension` .
- `label_name` : Python `str` , name of the key in label `dict` . Can be `None` if label is a `Tensor` (single headed models).
- `weight_column_name` : Python `str` defining feature column name representing weights. It is used to down weight or boost examples during training. It will be multiplied by the loss of the example.
- `enable_centered_bias` : Python `bool` . If `True` , estimator will learn a centered bias variable for each class. Rest of the model structure learns the residual after centered bias.
- `head_name` : Python `str` , name of the head. Predictions, summary and metrics keys are suffixed by `"/" + head_name` and the default variable scope is `head_name` .

Returns:

An instance of `Head` for generic regression.

Stay Connected

Blog

GitHub

Twitter

Support

Issue Tracker

Release Notes

Stack Overflow

English

Terms | **Privacy**