

tf.contrib.distributions.QuantizedDistribution

Contents

Class QuantizedDistribution

Properties

allow_nan_stats

batch_shape

Class **QuantizedDistribution**Inherits From: [Distribution](#)Defined in [tensorflow/contrib/distributions/python/ops/quantized_distribution.py](#).See the guide: [Statistical Distributions \(contrib\) > Transformed distributions](#)Distribution representing the quantization $Y = \text{ceiling}(X)$.

Definition in terms of sampling.

```

1. Draw X
2. Set Y <-- ceiling(X)
3. If Y < low, reset Y <-- low
4. If Y > high, reset Y <-- high
5. Return Y

```

Definition in terms of the probability mass function.

Given scalar random variable X , we define a discrete random variable Y supported on the integers as follows:
$$\begin{aligned}
 P[Y = j] &:= P[X \leq \text{low}], & \text{if } j == \text{low}, \\
 &:= P[X > \text{high} - 1], & j == \text{high}, \\
 &:= 0, & \text{if } j < \text{low} \text{ or } j > \text{high}, \\
 &:= P[j - 1 < X \leq j], & \text{all other } j.
 \end{aligned}$$
Conceptually, without cutoffs, the quantization process partitions the real line \mathbf{R} into half open intervals, and identifies an integer j with the right endpoints:
$$\begin{array}{cccccccc}
 \mathbf{R} = \dots & (-2, -1] & (-1, 0] & (0, 1] & (1, 2] & (2, 3] & (3, 4] & \dots \\
 j = \dots & -1 & 0 & 1 & 2 & 3 & 4 & \dots
 \end{array}$$
 $P[Y = j]$ is the mass of X within the j th interval. If $\text{low} = 0$, and $\text{high} = 2$, then the intervals are redrawn and j is re-assigned:
$$\begin{array}{cccc}
 \mathbf{R} = & (-\text{infty}, 0] & (0, 1] & (1, \text{infty}) \\
 j = & 0 & 1 & 2
 \end{array}$$
 $P[Y = j]$ is still the mass of X within the j th interval.

Caveats

Since evaluation of each $P[Y = j]$ involves a cdf evaluation (rather than a closed form function such as for a Poisson), computations such as mean and entropy are better done with samples or approximations, and are not implemented by this class.

Properties

allow_nan_stats

Python `bool` describing behavior when a stat is undefined.

Stats return +/- infinity when it makes sense. E.g., the variance of a Cauchy distribution is infinity. However, sometimes the statistic is undefined, e.g., if a distribution's pdf does not achieve a maximum within the support of the distribution, the mode is undefined. If the mean is undefined, then by definition the variance is undefined. E.g. the mean for Student's T for $df = 1$ is undefined (no clear way to say it is either + or - infinity), so the variance = $E[(X - \text{mean})^2]$ is also undefined.

Returns:

- `allow_nan_stats`: Python `bool`.

batch_shape

Shape of a single sample from a single event index as a `TensorShape`.

May be partially defined or unknown.

The batch dimensions are indexes into independent, non-identical parameterizations of this distribution.

Returns:

- `batch_shape`: `TensorShape`, possibly unknown.

distribution

Base distribution, $p(x)$.

dtype

The `DType` of `Tensor`s handled by this `Distribution`.

event_shape

Shape of a single sample from a single batch as a `TensorShape`.

May be partially defined or unknown.

Returns:

- `event_shape`: `TensorShape`, possibly unknown.

name

Name prepended to all ops created by this `Distribution`.

parameters

Dictionary of parameters used to instantiate this `Distribution`.

reparameterization_type

Describes how samples from the distribution are reparameterized.

Currently this is one of the static instances `distributions.FULLY_REPARAMETERIZED` or `distributions.NOT_REPARAMETERIZED`.

Returns:

An instance of `ReparameterizationType`.

validate_args

Python `bool` indicating possibly expensive checks are enabled.

Methods

`__init__`

```
__init__(
    distribution,
    low=None,
    high=None,
    validate_args=False,
    name='QuantizedDistribution'
)
```

Construct a Quantized Distribution representing $Y = \text{ceiling}(X)$.

Some properties are inherited from the distribution defining `X`. Example: `allow_nan_stats` is determined for this `QuantizedDistribution` by reading the `distribution`.

Args:

- `distribution`: The base distribution class to transform. Typically an instance of `Distribution`.
- `low`: `Tensor` with same `dtype` as this distribution and shape able to be added to samples. Should be a whole number. Default `None`. If provided, base distribution's `prob` should be defined at `low`.
- `high`: `Tensor` with same `dtype` as this distribution and shape able to be added to samples. Should be a whole number. Default `None`. If provided, base distribution's `prob` should be defined at `high - 1`. `high` must be strictly greater than `low`.
- `validate_args`: Python `bool`, default `False`. When `True` distribution parameters are checked for validity despite possibly degrading runtime performance. When `False` invalid inputs may silently render incorrect outputs.
- `name`: Python `str` name prefixed to Ops created by this class.

Raises:

- `TypeError` : If `dist_cls` is not a subclass of `Distribution` or continuous.
- `NotImplementedError` : If the base distribution does not implement `cdf`.

`batch_shape_tensor`

```
batch_shape_tensor(name='batch_shape_tensor')
```

Shape of a single sample from a single event index as a 1-D `Tensor`.

The batch dimensions are indexes into independent, non-identical parameterizations of this distribution.

Args:

- `name` : name to give to the op

Returns:

- `batch_shape` : `Tensor`.

`cdf`

```
cdf(  
    value,  
    name='cdf'  
)
```

Cumulative distribution function.

Given random variable `X`, the cumulative distribution function `cdf` is:

```
cdf(x) := P[X <= x]
```

Additional documentation from `QuantizedDistribution` :

For whole numbers `y`,

```
cdf(y) := P[Y <= y]  
        = 1, if y >= high,  
        = 0, if y < low,  
        = P[X <= y], otherwise.
```

Since `Y` only has mass at whole numbers, $P[Y \leq y] = P[Y \leq \text{floor}(y)]$. This dictates that fractional `y` are first floored to a whole number, and then above definition applies.

The base distribution's `cdf` method must be defined on `y - 1`.

Args:

- `value` : `float` or `double Tensor`.
- `name` : The name to give this op.

Returns:

- `cdf`: a `Tensor` of shape `sample_shape(x) + self.batch_shape` with values of type `self.dtype`.

copy

```
copy(**override_parameters_kwargs)
```

Creates a deep copy of the distribution.

★ **Note:** the copy distribution may continue to depend on the original initialization arguments.

Args:

- `**override_parameters_kwargs`: String/value dictionary of initialization arguments to override with new values.

Returns:

- `distribution`: A new instance of `type(self)` initialized from the union of `self.parameters` and `override_parameters_kwargs`, i.e., `dict(self.parameters, **override_parameters_kwargs)`.

covariance

```
covariance(name='covariance')
```

Covariance.

Covariance is (possibly) defined only for non-scalar-event distributions.

For example, for a length-`k`, vector-valued distribution, it is calculated as,

$$\text{Cov}[i, j] = \text{Covariance}(X_i, X_j) = E[(X_i - E[X_i]) (X_j - E[X_j])]$$

where `Cov` is a (batch of) `k x k` matrix, $0 \leq (i, j) < k$, and `E` denotes expectation.

Alternatively, for non-vector, multivariate distributions (e.g., matrix-valued, Wishart), `Covariance` shall return a (batch of) matrices under some vectorization of the events, i.e.,

$$\text{Cov}[i, j] = \text{Covariance}(\text{Vec}(X)_i, \text{Vec}(X)_j) = [\text{as above}]$$

where `Cov` is a (batch of) `k' x k'` matrices, $0 \leq (i, j) < k' = \text{reduce_prod}(\text{event_shape})$, and `Vec` is some function mapping indices of this distribution's event dimensions to indices of a length-`k'` vector.

Args:

- `name`: The name to give this op.

Returns:

- `covariance`: Floating-point `Tensor` with shape `[B1, ..., Bn, k', k']` where the first `n` dimensions are batch coordinates and `k' = reduce_prod(self.event_shape)`.

entropy

```
entropy(name='entropy')
```

Shannon entropy in nats.

event_shape_tensor

```
event_shape_tensor(name='event_shape_tensor')
```

Shape of a single sample from a single batch as a 1-D int32 **Tensor** .

Args:

- `name` : name to give to the op

Returns:

- `event_shape` : **Tensor** .

is_scalar_batch

```
is_scalar_batch(name='is_scalar_batch')
```

Indicates that `batch_shape == []` .

Args:

- `name` : The name to give this op.

Returns:

- `is_scalar_batch` : **bool** scalar **Tensor** .

is_scalar_event

```
is_scalar_event(name='is_scalar_event')
```

Indicates that `event_shape == []` .

Args:

- `name` : The name to give this op.

Returns:

- `is_scalar_event` : **bool** scalar **Tensor** .

log_cdf

```
log_cdf(
    value,
    name='log_cdf'
)
```

Log cumulative distribution function.

Given random variable X , the cumulative distribution function **cdf** is:

```
log_cdf(x) := Log[ P[X <= x] ]
```

Often, a numerical approximation can be used for **log_cdf(x)** that yields a more accurate answer than simply taking the logarithm of the **cdf** when $x \ll -1$.

Additional documentation from **QuantizedDistribution**:

For whole numbers y ,

```
cdf(y) := P[Y <= y]
        = 1, if y >= high,
        = 0, if y < low,
        = P[X <= y], otherwise.
```

Since Y only has mass at whole numbers, $P[Y \leq y] = P[Y \leq \text{floor}(y)]$. This dictates that fractional y are first floored to a whole number, and then above definition applies.

The base distribution's **log_cdf** method must be defined on $y - 1$.

Args:

- **value**: **float** or **double Tensor**.
- **name**: The name to give this op.

Returns:

- **logcdf**: a **Tensor** of shape **sample_shape(x) + self.batch_shape** with values of type **self.dtype**.

log_prob

```
log_prob(
    value,
    name='log_prob'
)
```

Log probability density/mass function.

Additional documentation from **QuantizedDistribution**:

For whole numbers y ,

```
P[Y = y] := P[X <= low], if y == low,
           := P[X > high - 1], y == high,
           := 0, if j < low or y > high,
           := P[y - 1 < X <= y], all other y.
```

The base distribution's **log_cdf** method must be defined on $y - 1$. If the base distribution has a **log_survival_function** method results will be more accurate for large values of y , and in this case the

`log_survival_function` must also be defined on `y - 1`.

Args:

- `value`: `float` or `double Tensor`.
- `name`: The name to give this op.

Returns:

- `log_prob`: a `Tensor` of shape `sample_shape(x) + self.batch_shape` with values of type `self.dtype`.

`log_survival_function`

```
log_survival_function(  
    value,  
    name='log_survival_function'  
)
```

Log survival function.

Given random variable `X`, the survival function is defined:

```
log_survival_function(x) = Log[ P[X > x] ]  
                        = Log[ 1 - P[X <= x] ]  
                        = Log[ 1 - cdf(x) ]
```

Typically, different numerical approximations can be used for the log survival function, which are more accurate than `1 - cdf(x)` when `x >> 1`.

Additional documentation from `QuantizedDistribution`:

For whole numbers `y`,

```
survival_function(y) := P[Y > y]  
                    = 0, if y >= high,  
                    = 1, if y < low,  
                    = P[X <= y], otherwise.
```

Since `Y` only has mass at whole numbers, `P[Y <= y] = P[Y <= floor(y)]`. This dictates that fractional `y` are first floored to a whole number, and then above definition applies.

The base distribution's `log_cdf` method must be defined on `y - 1`.

Args:

- `value`: `float` or `double Tensor`.
- `name`: The name to give this op.

Returns:

`Tensor` of shape `sample_shape(x) + self.batch_shape` with values of type `self.dtype`.

`mean`


```
mean(name='mean')
```

Mean.

mode

```
mode(name='mode')
```

Mode.

param_shapes

```
param_shapes(  
    cls,  
    sample_shape,  
    name='DistributionParamShapes'  
)
```

Shapes of parameters given the desired shape of a call to `sample()`.

This is a class method that describes what key/value arguments are required to instantiate the given `Distribution` so that a particular shape is returned for that instance's call to `sample()`.

Subclasses should override class method `_param_shapes`.

Args:

- `sample_shape`: `Tensor` or python list/tuple. Desired shape of a call to `sample()`.
- `name`: name to prepend ops with.

Returns:

`dict` of parameter name to `Tensor` shapes.

param_static_shapes

```
param_static_shapes(  
    cls,  
    sample_shape  
)
```

`param_shapes` with static (i.e. `TensorShape`) shapes.

This is a class method that describes what key/value arguments are required to instantiate the given `Distribution` so that a particular shape is returned for that instance's call to `sample()`. Assumes that the sample's shape is known statically.

Subclasses should override class method `_param_shapes` to return constant-valued tensors when constant values are fed.

Args:

- `sample_shape`: `TensorShape` or python list/tuple. Desired shape of a call to `sample()`.

Returns:

dict of parameter name to **TensorShape** .

Raises:

- **ValueError** : if **sample_shape** is a **TensorShape** and is not fully defined.

prob

```
prob(  
    value,  
    name='prob'  
)
```

Probability density/mass function.

Additional documentation from **QuantizedDistribution** :

For whole numbers **y** ,

```
P[Y = y] := P[X <= low],  if y == low,  
           := P[X > high - 1],  y == high,  
           := 0, if j < low or y > high,  
           := P[y - 1 < X <= y],  all other y.
```

The base distribution's **cdf** method must be defined on **y - 1** . If the base distribution has a **survival_function** method, results will be more accurate for large values of **y** , and in this case the **survival_function** must also be defined on **y - 1** .

Args:

- **value** : **float** or **double Tensor** .
- **name** : The name to give this op.

Returns:

- **prob** : a **Tensor** of shape **sample_shape(x) + self.batch_shape** with values of type **self.dtype** .

quantile

```
quantile(  
    value,  
    name='quantile'  
)
```

Quantile function. Aka "inverse cdf" or "percent point function".

Given random variable **X** and **p in [0, 1]** , the **quantile** is:

```
quantile(p) := x such that P[X <= x] == p
```

Args:

- `value`: `float` or `double Tensor`.
- `name`: The name to give this op.

Returns:

- `quantile`: a `Tensor` of shape `sample_shape(x) + self.batch_shape` with values of type `self.dtype`.

sample

```
sample(
    sample_shape=(),
    seed=None,
    name='sample'
)
```

Generate samples of the specified shape.

Note that a call to `sample()` without arguments will generate a single sample.

Args:

- `sample_shape`: 0D or 1D `int32 Tensor`. Shape of the generated samples.
- `seed`: Python integer seed for RNG
- `name`: name to give to the op.

Returns:

- `samples`: a `Tensor` with prepended dimensions `sample_shape`.

stddev

```
stddev(name='stddev')
```

Standard deviation.

Standard deviation is defined as,

$$\text{stddev} = E[(X - E[X])**2]**0.5$$

where `X` is the random variable associated with this distribution, `E` denotes expectation, and `stddev.shape = batch_shape + event_shape`.

Args:

- `name`: The name to give this op.

Returns:

- `stddev`: Floating-point `Tensor` with shape identical to `batch_shape + event_shape`, i.e., the same shape as `self.mean()`.

survival_function

```
survival_function(  
    value,  
    name='survival_function'  
)
```

Survival function.

Given random variable X , the survival function is defined:

```
survival_function(x) = P[X > x]  
                    = 1 - P[X <= x]  
                    = 1 - cdf(x).
```

Additional documentation from `QuantizedDistribution`:

For whole numbers y ,

```
survival_function(y) := P[Y > y]  
                    = 0, if y >= high,  
                    = 1, if y < low,  
                    = P[X <= y], otherwise.
```

Since Y only has mass at whole numbers, $P[Y \leq y] = P[Y \leq \text{floor}(y)]$. This dictates that fractional y are first floored to a whole number, and then above definition applies.

The base distribution's `cdf` method must be defined on $y - 1$.

Args:

- `value`: `float` or `double Tensor`.
- `name`: The name to give this op.

Returns:

`Tensor` of shape `sample_shape(x) + self.batch_shape` with values of type `self.dtype`.

variance

```
variance(name='variance')
```

Variance.

Variance is defined as,

```
Var = E[(X - E[X])**2]
```

where X is the random variable associated with this distribution, E denotes expectation, and `Var.shape = batch_shape + event_shape`.

Args:

- `name`: The name to give this op.

Returns:

- `variance`: Floating-point `Tensor` with shape identical to `batch_shape + event_shape`, i.e., the same shape as `self.mean()`.

Except as otherwise noted, the content of this page is licensed under the [Creative Commons Attribution 3.0 License](#), and code samples are licensed under the [Apache 2.0 License](#). For details, see our [Site Policies](#). Java is a registered trademark of Oracle and/or its affiliates.

Last updated November 2, 2017.

Stay Connected

[Blog](#)

[GitHub](#)

[Twitter](#)

Support

[Issue Tracker](#)

[Release Notes](#)

[Stack Overflow](#)

English

[Terms](#) | [Privacy](#)