

tf.contrib.distributions.VectorExponentialDiag

Contents

Class VectorExponentialDiag

Properties

allow_nan_stats

batch_shape

Class **VectorExponentialDiag**Defined in [tensorflow/contrib/distributions/python/ops/vector_exponential_diag.py](#).The vectorization of the Exponential distribution on \mathbb{R}^k .

The vector exponential distribution is defined over a subset of \mathbb{R}^k , and parameterized by a (batch of) length- k **loc** vector and a (batch of) $k \times k$ **scale** matrix: **covariance** = **scale** @ **scale.T**, where @ denotes matrix-multiplication.

Mathematical Details

The probability density function (pdf) is defined over the image of the **scale** matrix + **loc**, applied to the positive half-space: **Supp** = {**loc** + **scale** @ **x** : **x** in \mathbb{R}^k , $x_1 > 0$, ..., $x_k > 0$ }. On this set,

```
pdf(y; loc, scale) = exp(-||x||_1) / Z, for y in Supp
x = inv(scale) @ (y - loc),
Z = |det(scale)|,
```

where:

- **loc** is a vector in \mathbb{R}^k ,
- **scale** is a linear operator in $\mathbb{R}^{k \times k}$, **cov** = **scale** @ **scale.T**,
- **Z** denotes the normalization constant, and,
- $||x||_1$ denotes the **l1** norm of **x**, $\sum_i |x_i|$.

The VectorExponential distribution is a member of the [location-scale family](#), i.e., it can be constructed as,

```
X = (X_1, ..., X_k), each X_i ~ Exponential(rate=1)
Y = (Y_1, ..., Y_k) = scale @ X + loc
```

About **VectorExponential** and **Vector** distributions in TensorFlow.

The **VectorExponential** is a non-standard distribution that has useful properties.

The marginals **Y_1**, ..., **Y_k** are *not* Exponential random variables, due to the fact that the sum of Exponential random variables is not Exponential.

Instead, **Y** is a vector whose components are linear combinations of Exponential random variables. Thus, **Y** lives in the vector space generated by **vectors** of Exponential distributions. This allows the user to decide the mean and covariance

(by setting `loc` and `scale`), while preserving some properties of the Exponential distribution. In particular, the tails of \mathbf{Y}_i will be (up to polynomial factors) exponentially decaying.

To see this last statement, note that the pdf of \mathbf{Y}_i is the convolution of the pdf of \mathbf{k} independent Exponential random variables. One can then show by induction that distributions with exponential (up to polynomial factors) tails are closed under convolution.

Examples

```
ds = tf.contrib.distributions
la = tf.contrib.linalg

# Initialize a single 2-variate VectorExponential, supported on
# {(x, y) in R^2 : x > 0, y > 0}.

# The first component has pdf exp{-x}, the second 0.5 exp{-x / 2}
vex = ds.VectorExponentialDiag(scale_diag=[1., 2.])

# Compute the pdf of an `R^2` observation; return a scalar.
vex.prob([3., 4.]).eval() # shape: []

# Initialize a 2-batch of 3-variate Vector Exponential's.
loc = [[1., 2, 3],
       [1., 0, 0]] # shape: [2, 3]
scale_diag = [[1., 2, 3],
              [0.5, 1, 1.5]] # shape: [2, 3]

vex = ds.VectorExponentialDiag(loc, scale_diag)

# Compute the pdf of two `R^3` observations; return a length-2 vector.
x = [[1.9, 2.2, 3.1],
     [10., 1.0, 9.0]] # shape: [2, 3]
vex.prob(x).eval() # shape: [2]
```

Properties

`allow_nan_stats`

Python `bool` describing behavior when a stat is undefined.

Stats return +/- infinity when it makes sense. E.g., the variance of a Cauchy distribution is infinity. However, sometimes the statistic is undefined, e.g., if a distribution's pdf does not achieve a maximum within the support of the distribution, the mode is undefined. If the mean is undefined, then by definition the variance is undefined. E.g. the mean for Student's T for $df = 1$ is undefined (no clear way to say it is either + or - infinity), so the variance = $E[(X - \text{mean})^2]$ is also undefined.

Returns:

- `allow_nan_stats`: Python `bool`.

`batch_shape`

Shape of a single sample from a single event index as a `TensorShape`.

May be partially defined or unknown.

The batch dimensions are indexes into independent, non-identical parameterizations of this distribution.

Returns:

- `batch_shape` : `TensorShape` , possibly unknown.

bijector

Function transforming $x \Rightarrow y$.

distribution

Base distribution, $p(x)$.

dtype

The `DType` of `Tensor` s handled by this `Distribution` .

event_shape

Shape of a single sample from a single batch as a `TensorShape` .

May be partially defined or unknown.

Returns:

- `event_shape` : `TensorShape` , possibly unknown.

loc

The `loc Tensor` in $Y = \text{scale} @ X + \text{loc}$.

name

Name prepended to all ops created by this `Distribution` .

parameters

Dictionary of parameters used to instantiate this `Distribution` .

reparameterization_type

Describes how samples from the distribution are reparameterized.

Currently this is one of the static instances `distributions.FULLY_REPARAMETERIZED` or `distributions.NOT_REPARAMETERIZED` .

Returns:

An instance of `ReparameterizationType` .

scale

The `scale` `LinearOperator` in `Y = scale @ X + loc`.

`validate_args`

Python `bool` indicating possibly expensive checks are enabled.

Methods

`__init__`

```
__init__(
    loc=None,
    scale_diag=None,
    scale_identity_multiplier=None,
    validate_args=False,
    allow_nan_stats=True,
    name='VectorExponentialDiag'
)
```

Construct Vector Exponential distribution supported on a subset of \mathbb{R}^k .

The `batch_shape` is the broadcast shape between `loc` and `scale` arguments.

The `event_shape` is given by last dimension of the matrix implied by `scale`. The last dimension of `loc` (if provided) must broadcast with this.

Recall that `covariance = scale @ scale.T`.

```
scale = diag(scale_diag + scale_identity_multiplier * ones(k))
```

where:

- `scale_diag.shape = [k]`, and,
- `scale_identity_multiplier.shape = []`.

Additional leading dimensions (if any) will index batches.

If both `scale_diag` and `scale_identity_multiplier` are `None`, then `scale` is the Identity matrix.

Args:

- `loc`: Floating-point `Tensor`. If this is set to `None`, `loc` is implicitly `0`. When specified, may have shape `[B1, ..., Bb, k]` where `b >= 0` and `k` is the event size.
- `scale_diag`: Non-zero, floating-point `Tensor` representing a diagonal matrix added to `scale`. May have shape `[B1, ..., Bb, k]`, `b >= 0`, and characterizes `b`-batches of `k x k` diagonal matrices added to `scale`. When both `scale_identity_multiplier` and `scale_diag` are `None` then `scale` is the `Identity`.
- `scale_identity_multiplier`: Non-zero, floating-point `Tensor` representing a scaled-identity-matrix added to `scale`. May have shape `[B1, ..., Bb]`, `b >= 0`, and characterizes `b`-batches of scaled `k x k` identity matrices added to `scale`. When both `scale_identity_multiplier` and `scale_diag` are `None` then `scale` is the `Identity`.
- `validate_args`: Python `bool`, default `False`. When `True` distribution parameters are checked for validity despite possibly degrading runtime performance. When `False` invalid inputs may silently render incorrect outputs.
- `allow_nan_stats`: Python `bool`, default `True`. When `True`, statistics (e.g., mean, mode, variance) use the value "`NaN`" to indicate the result is undefined. When `False`, an exception is raised if one or more of the statistic's batch members are undefined.

- `name` : Python `str` name prefixed to Ops created by this class.

Raises:

- `ValueError` : if at most `scale_identity_multiplier` is specified.

batch_shape_tensor

```
batch_shape_tensor(name='batch_shape_tensor')
```

Shape of a single sample from a single event index as a 1-D `Tensor` .

The batch dimensions are indexes into independent, non-identical parameterizations of this distribution.

Args:

- `name` : name to give to the op

Returns:

- `batch_shape` : `Tensor` .

cdf

```
cdf(
    value,
    name='cdf'
)
```

Cumulative distribution function.

Given random variable `X`, the cumulative distribution function `cdf` is:

```
cdf(x) := P[X <= x]
```

Args:

- `value` : `float` or `double Tensor` .
- `name` : The name to give this op.

Returns:

- `cdf` : a `Tensor` of shape `sample_shape(x) + self.batch_shape` with values of type `self.dtype` .

copy

```
copy(**override_parameters_kwargs)
```

Creates a deep copy of the distribution.

★ **Note:** the copy distribution may continue to depend on the original initialization arguments.

Args:

- `**override_parameters_kwargs`: String/value dictionary of initialization arguments to override with new values.

Returns:

- `distribution`: A new instance of `type(self)` initialized from the union of `self.parameters` and `override_parameters_kwargs`, i.e., `dict(self.parameters, **override_parameters_kwargs)`.

covariance

```
covariance(name='covariance')
```

Covariance.

Covariance is (possibly) defined only for non-scalar-event distributions.

For example, for a length-`k`, vector-valued distribution, it is calculated as,

```
Cov[i, j] = Covariance(X_i, X_j) = E[(X_i - E[X_i]) (X_j - E[X_j])]
```

where `Cov` is a (batch of) `k x k` matrix, `0 <= (i, j) < k`, and `E` denotes expectation.

Alternatively, for non-vector, multivariate distributions (e.g., matrix-valued, Wishart), `Covariance` shall return a (batch of) matrices under some vectorization of the events, i.e.,

```
Cov[i, j] = Covariance(Vec(X)_i, Vec(X)_j) = [as above]
```

where `Cov` is a (batch of) `k' x k'` matrices, `0 <= (i, j) < k' = reduce_prod(event_shape)`, and `Vec` is some function mapping indices of this distribution's event dimensions to indices of a length-`k'` vector.

Args:

- `name`: The name to give this op.

Returns:

- `covariance`: Floating-point `Tensor` with shape `[B1, ..., Bn, k', k']` where the first `n` dimensions are batch coordinates and `k' = reduce_prod(self.event_shape)`.

entropy

```
entropy(name='entropy')
```

Shannon entropy in nats.

event_shape_tensor

```
event_shape_tensor(name='event_shape_tensor')
```

Shape of a single sample from a single batch as a 1-D int32 `Tensor`.

Args:

- `name` : name to give to the op

Returns:

- `event_shape` : `Tensor` .

`is_scalar_batch`

```
is_scalar_batch(name='is_scalar_batch')
```

Indicates that `batch_shape == []` .

Args:

- `name` : The name to give this op.

Returns:

- `is_scalar_batch` : `bool` scalar `Tensor` .

`is_scalar_event`

```
is_scalar_event(name='is_scalar_event')
```

Indicates that `event_shape == []` .

Args:

- `name` : The name to give this op.

Returns:

- `is_scalar_event` : `bool` scalar `Tensor` .

`log_cdf`

```
log_cdf(  
    value,  
    name='log_cdf'  
)
```

Log cumulative distribution function.

Given random variable `X` , the cumulative distribution function `cdf` is:

```
log_cdf(x) := Log[ P[X <= x] ]
```

Often, a numerical approximation can be used for `log_cdf(x)` that yields a more accurate answer than simply taking the logarithm of the `cdf` when `x << -1` .

Args:

- `value`: `float` or `double Tensor`.
- `name`: The name to give this op.

Returns:

- `logcdf`: a `Tensor` of shape `sample_shape(x) + self.batch_shape` with values of type `self.dtype`.

log_prob

```
log_prob(  
    value,  
    name='log_prob'  
)
```

Log probability density/mass function.

Additional documentation from `VectorExponentialLinearOperator`:

`value` is a batch vector with compatible shape if `value` is a `Tensor` whose shape can be broadcast up to either:

```
self.batch_shape + self.event_shape
```

or

```
[M1, ..., Mm] + self.batch_shape + self.event_shape
```

Args:

- `value`: `float` or `double Tensor`.
- `name`: The name to give this op.

Returns:

- `log_prob`: a `Tensor` of shape `sample_shape(x) + self.batch_shape` with values of type `self.dtype`.

log_survival_function

```
log_survival_function(  
    value,  
    name='log_survival_function'  
)
```

Log survival function.

Given random variable `X`, the survival function is defined:

```
log_survival_function(x) = Log[ P[X > x] ]  
                        = Log[ 1 - P[X <= x] ]  
                        = Log[ 1 - cdf(x) ]
```

Typically, different numerical approximations can be used for the log survival function, which are more accurate than `1 - cdf(x)` when `x >> 1`.

Args:

- `value`: `float` or `double Tensor`.
- `name`: The name to give this op.

Returns:

`Tensor` of shape `sample_shape(x) + self.batch_shape` with values of type `self.dtype`.

mean

```
mean(name='mean')
```

Mean.

mode

```
mode(name='mode')
```

Mode.

param_shapes

```
param_shapes(  
    cls,  
    sample_shape,  
    name='DistributionParamShapes'  
)
```

Shapes of parameters given the desired shape of a call to `sample()`.

This is a class method that describes what key/value arguments are required to instantiate the given `Distribution` so that a particular shape is returned for that instance's call to `sample()`.

Subclasses should override class method `_param_shapes`.

Args:

- `sample_shape`: `Tensor` or python list/tuple. Desired shape of a call to `sample()`.
- `name`: name to prepend ops with.

Returns:

`dict` of parameter name to `Tensor` shapes.

param_static_shapes

```
param_static_shapes(  
    cls,  
    sample_shape  
)
```

param_shapes with static (i.e. `TensorShape`) shapes.

This is a class method that describes what key/value arguments are required to instantiate the given `Distribution` so that a particular shape is returned for that instance's call to `sample()`. Assumes that the sample's shape is known statically.

Subclasses should override class method `_param_shapes` to return constant-valued tensors when constant values are fed.

Args:

- `sample_shape`: `TensorShape` or python list/tuple. Desired shape of a call to `sample()`.

Returns:

`dict` of parameter name to `TensorShape`.

Raises:

- `ValueError`: if `sample_shape` is a `TensorShape` and is not fully defined.

prob

```
prob(  
    value,  
    name='prob'  
)
```

Probability density/mass function.

Additional documentation from `VectorExponentialLinearOperator`:

`value` is a batch vector with compatible shape if `value` is a `Tensor` whose shape can be broadcast up to either:

```
self.batch_shape + self.event_shape
```

or

```
[M1, ..., Mm] + self.batch_shape + self.event_shape
```

Args:

- `value`: `float` or `double Tensor`.
- `name`: The name to give this op.

Returns:

- `prob`: a `Tensor` of shape `sample_shape(x) + self.batch_shape` with values of type `self.dtype`.

quantile

```
quantile(
    value,
    name='quantile'
)
```

Quantile function. Aka "inverse cdf" or "percent point function".

Given random variable **X** and **p in [0, 1]**, the **quantile** is:

```
quantile(p) := x such that P[X <= x] == p
```

Args:

- **value**: **float** or **double Tensor**.
- **name**: The name to give this op.

Returns:

- **quantile**: a **Tensor** of shape **sample_shape(x) + self.batch_shape** with values of type **self.dtype**.

sample

```
sample(
    sample_shape=(),
    seed=None,
    name='sample'
)
```

Generate samples of the specified shape.

Note that a call to **sample()** without arguments will generate a single sample.

Args:

- **sample_shape**: 0D or 1D **int32 Tensor**. Shape of the generated samples.
- **seed**: Python integer seed for RNG
- **name**: name to give to the op.

Returns:

- **samples**: a **Tensor** with prepended dimensions **sample_shape**.

stddev

```
stddev(name='stddev')
```

Standard deviation.

Standard deviation is defined as,

```
stddev = E[(X - E[X])**2]**0.5
```

where **X** is the random variable associated with this distribution, **E** denotes expectation, and **stddev.shape =**

`batch_shape + event_shape` .

Args:

- `name` : The name to give this op.

Returns:

- `stddev` : Floating-point `Tensor` with shape identical to `batch_shape + event_shape` , i.e., the same shape as `self.mean()` .

survival_function

```
survival_function(  
    value,  
    name='survival_function'  
)
```

Survival function.

Given random variable `X` , the survival function is defined:

```
survival_function(x) = P[X > x]  
                    = 1 - P[X <= x]  
                    = 1 - cdf(x) .
```

Args:

- `value` : `float` or `double Tensor` .
- `name` : The name to give this op.

Returns:

`Tensor` of shape `sample_shape(x) + self.batch_shape` with values of type `self.dtype` .

variance

```
variance(name='variance')
```

Variance.

Variance is defined as,

```
Var = E[(X - E[X])**2]
```

where `X` is the random variable associated with this distribution, `E` denotes expectation, and `Var.shape = batch_shape + event_shape` .

Args:

- `name` : The name to give this op.

Returns:

- `variance`: Floating-point `Tensor` with shape identical to `batch_shape + event_shape`, i.e., the same shape as `self.mean()`.

Except as otherwise noted, the content of this page is licensed under the [Creative Commons Attribution 3.0 License](#), and code samples are licensed under the [Apache 2.0 License](#). For details, see our [Site Policies](#). Java is a registered trademark of Oracle and/or its affiliates.

Last updated November 2, 2017.

Stay Connected

[Blog](#)

[GitHub](#)

[Twitter](#)

Support

[Issue Tracker](#)

[Release Notes](#)

[Stack Overflow](#)

English

[Terms](#) | [Privacy](#)