tf.contrib.estimator.binary_classification_head

```
binary_classification_head(
    weight_column=None,
    thresholds=None,
    label_vocabulary=None,
    name=None
)
```

Defined in `tensorflow/contrib/estimator/python/estimator/head.py`.

Creates a `_Head` for single label binary classification.

This head uses `sigmoid_cross_entropy_with_logits` loss.

This head expects to be fed float labels of shape `(batch_size, 1)`.

Args:

- `weight_column`: A string or a `_NumericColumn` created by `tf.feature_column.numeric_column` defining feature column representing weights. It is used to down weight or boost examples during training. It will be multiplied by the loss of the example.
- `thresholds`: Iterable of floats in the range `(0, 1)`. For binary classification metrics such as precision and recall, an eval metric is generated for each threshold value. This threshold is applied to the logistic values to determine the binary classification (i.e., above the threshold is `true`, below is `false`).
- `label_vocabulary`: A list of strings represents possible label values. If it is not given, that means labels are already encoded within [0, 1]. If given, labels must be string type and have any value in `label_vocabulary`. Also there will be errors if vocabulary is not provided and labels are string.
- `name`: name of the head. If provided, summary and metrics keys will be suffixed by `"/" + name`.

Returns:

An instance of `_Head` for binary classification.

Raises:

- `ValueError`: if `thresholds` contains a value outside of `(0, 1)`.

Last updated November 2, 2017.

Stay Connected

Blog

GitHub

Twitter

Support

Issue Tracker

Release Notes

Stack Overflow

English

Terms | **Privacy**