tf.contrib.gan.losses.wargs.acgan_generator_loss

```
acgan_generator_loss(
    discriminator_gen_classification_logits,
    one_hot_labels,
    weights=1.0,
    scope=None,
    loss_collection=tf.GraphKeys.LOSSES,
    reduction=losses.Reduction.SUM_BY_NONZERO_WEIGHTS,
    add_summaries=False
)
```

Defined in `tensorflow/contrib/gan/python/losses/python/losses_impl.py`.

ACGAN loss for the generator.

The ACGAN loss adds a classification loss to the conditional discriminator. Therefore, the discriminator must output a tuple consisting of (1) the real/fake prediction and (2) the logits for the classification (usually the last conv layer, flattened).

For more details: ACGAN: https://arxiv.org/abs/1610.09585

Args:

- `discriminator_gen_classification_logits`: Classification logits for generated data.
- `one_hot_labels`: A Tensor holding one-hot labels for the batch.
- `weights`: Optional `Tensor` whose rank is either 0, or the same rank as `discriminator_gen_classification_logits`, and must be broadcastable to `discriminator_gen_classification_logits` (i.e., all dimensions must be either `1`, or the same as the corresponding dimension).
- `scope`: The scope for the operations performed in computing the loss.
- `loss_collection`: collection to which this loss will be added.
- `reduction`: A `tf.losses.Reduction` to apply to loss.
- `add_summaries`: Whether or not to add summaries for the loss.

Returns:

A loss Tensor. Shape depends on `reduction`.

Raises:

- `ValueError`: if arg module not either `generator` or `discriminator`
- `TypeError`: if the discriminator does not output a tuple.

Stay Connected

Blog

GitHub

Twitter

Support

Issue Tracker

Release Notes

Stack Overflow

English

Terms | **Privacy**