tf.contrib.layers.bow_encoder

```
bow_encoder(
    ids,
    vocab_size,
    embed_dim,
    sparse_lookup=True,
    initializer=None,
    regularizer=None,
    trainable=True,
    scope=None,
    reuse=None
)
```

Defined in `tensorflow/contrib/layers/python/layers/encoders.py`.

Maps a sequence of symbols to a vector per example by averaging embeddings.

Args:

- `ids`: `[batch_size, doc_length]` `Tensor` or `SparseTensor` of type `int32` or `int64` with symbol ids.
- `vocab_size`: Integer number of symbols in vocabulary.
- `embed_dim`: Integer number of dimensions for embedding matrix.
- `sparse_lookup`: `bool`, if `True`, converts ids to a `SparseTensor` and performs a sparse embedding lookup. This is usually faster, but not desirable if padding tokens should have an embedding. Empty rows are assigned a special embedding.
- `initializer`: An initializer for the embeddings, if `None` default for current scope is used.
- `regularizer`: Optional regularizer for the embeddings.
- `trainable`: If `True` also add variables to the graph collection `GraphKeys.TRAINABLE_VARIABLES` (see tf.Variable).
- `scope`: Optional string specifying the variable scope for the op, required if `reuse=True`.
- `reuse`: If `True`, variables inside the op will be reused.

Returns:

Encoding `Tensor` `[batch_size, embed_dim]` produced by averaging embeddings.

Raises:

- `ValueError`: If `embed_dim` or `vocab_size` are not specified.

Stay Connected

Blog

GitHub

Twitter

Support

Issue Tracker

Release Notes

Stack Overflow

English

Terms | **Privacy**