tf.contrib.layers.embed_sequence

```
embed_sequence(
    ids,
    vocab_size=None,
    embed_dim=None,
    unique=False,
    initializer=None,
    regularizer=None,
    trainable=True,
    scope=None,
    reuse=None
)
```

Defined in `tensorflow/contrib/layers/python/layers/encoders.py`.

See the guide: Layers (contrib) > Higher level ops for building neural network layers

Maps a sequence of symbols to a sequence of embeddings.

Typical use case would be reusing embeddings between an encoder and decoder.

Args:

- `ids`: `[batch_size, doc_length]` `Tensor` of type `int32` or `int64` with symbol ids.
- `vocab_size`: Integer number of symbols in vocabulary.
- `embed_dim`: Integer number of dimensions for embedding matrix.
- `unique`: If `True`, will first compute the unique set of indices, and then lookup each embedding once, repeating them in the output as needed.
- `initializer`: An initializer for the embeddings, if `None` default for current scope is used.
- `regularizer`: Optional regularizer for the embeddings.
- `trainable`: If `True` also add variables to the graph collection `GraphKeys.TRAINABLE_VARIABLES` (see `tf.Variable`).
- `scope`: Optional string specifying the variable scope for the op, required if `reuse=True`.
- `reuse`: If `True`, variables inside the op will be reused.

Returns:

`Tensor` of `[batch_size, doc_length, embed_dim]` with embedded sequences.

Raises:

- `ValueError`: if `embed_dim` or `vocab_size` are not specified when `reuse` is `None` or `False`.

Stay Connected

Blog

GitHub

Twitter

Support

Issue Tracker

Release Notes

Stack Overflow

English

Terms | **Privacy**