

tf.contrib.learn.DNNEstimator

Contents

Class DNNEstimator

Properties

config

model_dir

Class **DNNEstimator**Inherits From: [Estimator](#)Defined in [tensorflow/contrib/learn/python/learn/estimators/dnn.py](#).

A Estimator for TensorFlow DNN models with user specified _Head.

Example:

```
sparse_feature_a = sparse_column_with_hash_bucket(...)
sparse_feature_b = sparse_column_with_hash_bucket(...)

sparse_feature_a_emb = embedding_column(sparse_id_column=sparse_feature_a,
                                         ...)
sparse_feature_b_emb = embedding_column(sparse_id_column=sparse_feature_b,
                                         ...)
```

To create a DNNEstimator for binary classification, where

```
estimator = DNNEstimator(
    feature_columns=[sparse_feature_a_emb, sparse_feature_b_emb],
    head=tf.contrib.learn.multi_class_head(n_classes=2),
    hidden_units=[1024, 512, 256])
```

If your label is keyed with "y" in your labels dict, and weights are keyed with "w" in features dict, and you want to enable centered bias,

```
head = tf.contrib.learn.multi_class_head(
    n_classes=2,
    label_name="x",
    weight_column_name="w",
    enable_centered_bias=True)
estimator = DNNEstimator(
    feature_columns=[sparse_feature_a_emb, sparse_feature_b_emb],
    head=head,
    hidden_units=[1024, 512, 256])
```

Input builders

```
def input_fn_train: # returns x, y (where y represents label's class index).
    pass
estimator.fit(input_fn=input_fn_train)
```

```
def input_fn_eval: # returns x, y (where y represents label's class index).
    pass
estimator.evaluate(input_fn=input_fn_eval)
estimator.predict(x=x) # returns predicted labels (i.e. label's class index).
```

Input of `fit` and `evaluate` should have following features, otherwise there will be a `KeyError` :

- if `weight_column_name` is not `None`, a feature with `key=weight_column_name` whose value is a `Tensor` .
- for each `column` in `feature_columns` :
- if `column` is a `SparseColumn` , a feature with `key=column.name` whose `value` is a `SparseTensor` .
- if `column` is a `WeightedSparseColumn` , two features: the first with `key` the id column name, the second with `key` the weight column name. Both features' `value` must be a `SparseTensor` .
- if `column` is a `RealValuedColumn` , a feature with `key=column.name` whose `value` is a `Tensor` .

Properties

`config`

`model_dir`

Methods

`__init__`

```
__init__(
    head,
    hidden_units,
    feature_columns,
    model_dir=None,
    optimizer=None,
    activation_fn=tf.nn.relu,
    dropout=None,
    gradient_clip_norm=None,
    config=None,
    feature_engineering_fn=None,
    embedding_lr_multipliers=None,
    input_layer_min_slice_size=None
)
```

Initializes a `DNNEstimator` instance.

Args:

- `head` : `Head` instance.
- `hidden_units` : List of hidden units per layer. All layers are fully connected. Ex. `[64, 32]` means first layer has 64 nodes and second one has 32.
- `feature_columns` : An iterable containing all the feature columns used by the model. All items in the set should be instances of classes derived from `FeatureColumn` .
- `model_dir` : Directory to save model parameters, graph and etc. This can also be used to load checkpoints from the directory into a estimator to continue training a previously saved model.
- `optimizer` : An instance of `tf.Optimizer` used to train the model. If `None` , will use an Adagrad optimizer.
- `activation_fn` : Activation function applied to each layer. If `None` , will use `tf.nn.relu` . Note that a string containing the unqualified name of the op may also be provided, e.g., "relu", "tanh", or "sigmoid".
- `dropout` : When not `None` , the probability we will drop out a given coordinate.
- `gradient_clip_norm` : A float > 0. If provided, gradients are clipped to their global norm with this clipping ratio. See `tf.clip_by_global_norm` for more details.

- `config`: `RunConfig` object to configure the runtime settings.
- `feature_engineering_fn`: Feature engineering function. Takes features and labels which are the output of `input_fn` and returns features and labels which will be fed into the model.
- `embedding_lr_multipliers`: Optional. A dictionary from `EmbeddingColumn` to a `float` multiplier. Multiplier will be used to multiply with learning rate for the embedding variables.
- `input_layer_min_slice_size`: Optional. The min slice size of input layer partitions. If not provided, will use the default of 64M.

Returns:

A `DNNEstimator` estimator.

evaluate

```
evaluate(
    x=None,
    y=None,
    input_fn=None,
    feed_fn=None,
    batch_size=None,
    steps=None,
    metrics=None,
    name=None,
    checkpoint_path=None,
    hooks=None,
    log_progress=True
)
```

See `Evaluable`. (deprecated arguments)

SOME ARGUMENTS ARE DEPRECATED. They will be removed after 2016-12-01. Instructions for updating: Estimator is decoupled from Scikit Learn interface by moving into separate class SKCompat. Arguments x, y and batch_size are only available in the SKCompat class, Estimator will only accept input_fn. Example conversion: `est = Estimator(...)` -> `est = SKCompat(Estimator(...))`

Raises:

- `ValueError`: If at least one of `x` or `y` is provided, and at least one of `input_fn` or `feed_fn` is provided. Or if `metrics` is not `None` or `dict`.

export

```
export(
    export_dir,
    input_fn=export._default_input_fn,
    input_feature_key=None,
    use_deprecated_input_fn=True,
    signature_fn=None,
    prediction_key=None,
    default_batch_size=1,
    exports_to_keep=None,
    checkpoint_path=None
)
```

Exports inference graph into given dir. (deprecated)

THIS FUNCTION IS DEPRECATED. It will be removed after 2017-03-25. Instructions for updating: Please use `Estimator.export_savedmodel()` instead.

Args:

- `export_dir`: A string containing a directory to write the exported graph and checkpoints.
- `input_fn`: If `use_deprecated_input_fn` is true, then a function that given `Tensor` of `Example` strings, parses it into features that are then passed to the model. Otherwise, a function that takes no argument and returns a tuple of (features, labels), where features is a dict of string key to `Tensor` and labels is a `Tensor` that's currently not used (and so can be `None`).
- `input_feature_key`: Only used if `use_deprecated_input_fn` is false. String key into the features dict returned by `input_fn` that corresponds to a the raw `Example` strings `Tensor` that the exported model will take as input. Can only be `None` if you're using a custom `signature_fn` that does not use the first arg (examples).
- `use_deprecated_input_fn`: Determines the signature format of `input_fn`.
- `signature_fn`: Function that returns a default signature and a named signature map, given `Tensor` of `Example` strings, `dict` of `Tensor` s for features and `Tensor` or `dict` of `Tensor` s for predictions.
- `prediction_key`: The key for a tensor in the `predictions` dict (output from the `model_fn`) to use as the `predictions` input to the `signature_fn`. Optional. If `None`, predictions will pass to `signature_fn` without filtering.
- `default_batch_size`: Default batch size of the `Example` placeholder.
- `exports_to_keep`: Number of exports to keep.
- `checkpoint_path`: the checkpoint path of the model to be exported. If it is `None` (which is default), will use the latest checkpoint in `export_dir`.

Returns:

The string path to the exported directory. NB: this functionality was added ca. 2016/09/25; clients that depend on the return value may need to handle the case where this function returns `None` because subclasses are not returning a value.

`export_savedmodel`

```
export_savedmodel(  
    export_dir_base,  
    serving_input_fn,  
    default_output_alternative_key=None,  
    assets_extra=None,  
    as_text=False,  
    checkpoint_path=None,  
    graph_rewrite_specs=(GraphRewriteSpec((tag_constants.SERVING,)), ()),  
)
```

Exports inference graph as a `SavedModel` into given dir.

Args:

- `export_dir_base`: A string containing a directory to write the exported graph and checkpoints.
- `serving_input_fn`: A function that takes no argument and returns an `InputFnOps`.
- `default_output_alternative_key`: the name of the head to serve when none is specified. Not needed for single-headed models.
- `assets_extra`: A dict specifying how to populate the `assets.extra` directory within the exported `SavedModel`. Each key should give the destination path (including the filename) relative to the `assets.extra` directory. The corresponding

value gives the full path of the source file to be copied. For example, the simple case of copying a single file without renaming it is specified as `{'my_asset_file.txt': '/path/to/my_asset_file.txt'}`.

- `as_text` : whether to write the SavedModel proto in text format.
- `checkpoint_path` : The checkpoint path to export. If `None` (the default), the most recent checkpoint found within the model directory is chosen.
- `graph_rewrite_specs` : an iterable of `GraphRewriteSpec`. Each element will produce a separate MetaGraphDef within the exported SavedModel, tagged and rewritten as specified. Defaults to a single entry using the default serving tag ("serve") and no rewriting.

Returns:

The string path to the exported directory.

Raises:

- `ValueError` : if an unrecognized export_type is requested.

fit

```
fit(  
    x=None,  
    y=None,  
    input_fn=None,  
    steps=None,  
    batch_size=None,  
    monitors=None,  
    max_steps=None  
)
```

See `Trainable`. (deprecated arguments)

SOME ARGUMENTS ARE DEPRECATED. They will be removed after 2016-12-01. Instructions for updating: Estimator is decoupled from Scikit Learn interface by moving into separate class SKCompat. Arguments x, y and batch_size are only available in the SKCompat class, Estimator will only accept input_fn. Example conversion: `est = Estimator(...)` -> `est = SKCompat(Estimator(...))`

Raises:

- `ValueError` : If `x` or `y` are not `None` while `input_fn` is not `None`.
- `ValueError` : If both `steps` and `max_steps` are not `None`.

get_params

```
get_params(deep=True)
```

Get parameters for this estimator.

Args:

- `deep` : boolean, optional
If `True`, will return the parameters for this estimator and contained subobjects that are estimators.

Returns:

- `params` : mapping of string to any Parameter names mapped to their values.

get_variable_names

```
get_variable_names()
```

Returns list of all variable names in this model.

Returns:

List of names.

get_variable_value

```
get_variable_value(name)
```

Returns value of the variable given by name.

Args:

- `name` : string, name of the tensor.

Returns:

Numpy array - value of the tensor.

partial_fit

```
partial_fit(  
    x=None,  
    y=None,  
    input_fn=None,  
    steps=1,  
    batch_size=None,  
    monitors=None  
)
```

Incremental fit on a batch of samples. (deprecated arguments)

SOME ARGUMENTS ARE DEPRECATED. They will be removed after 2016-12-01. Instructions for updating: Estimator is decoupled from Scikit Learn interface by moving into separate class SKCompat. Arguments x, y and batch_size are only available in the SKCompat class, Estimator will only accept input_fn. Example conversion: `est = Estimator(...)` -> `est = SKCompat(Estimator(...))`

This method is expected to be called several times consecutively on different or the same chunks of the dataset. This either can implement iterative training or out-of-core/online training.

This is especially useful when the whole dataset is too big to fit in memory at the same time. Or when model is taking long time to converge, and you want to split up training into subparts.

Args:

- `x` : Matrix of shape `[n_samples, n_features...]`. Can be iterator that returns arrays of features. The training input samples for fitting the model. If set, `input_fn` must be `None` .
- `y` : Vector or matrix `[n_samples]` or `[n_samples, n_outputs]`. Can be iterator that returns array of labels. The training label values (class labels in classification, real numbers in regression). If set, `input_fn` must be `None` .
- `input_fn` : Input function. If set, `x` , `y` , and `batch_size` must be `None` .
- `steps` : Number of steps for which to train model. If `None` , train forever.
- `batch_size` : minibatch size to use on the input, defaults to first dimension of `x` . Must be `None` if `input_fn` is provided.
- `monitors` : List of `BaseMonitor` subclass instances. Used for callbacks inside the training loop.

Returns:

`self` , for chaining.

Raises:

- `ValueError` : If at least one of `x` and `y` is provided, and `input_fn` is provided.

predict

```
predict(
    x=None,
    input_fn=None,
    batch_size=None,
    outputs=None,
    as_iterable=True
)
```

Returns predictions for given features. (deprecated arguments)

SOME ARGUMENTS ARE DEPRECATED. They will be removed after 2016-12-01. Instructions for updating: Estimator is decoupled from Scikit Learn interface by moving into separate class `SKCompat`. Arguments `x`, `y` and `batch_size` are only available in the `SKCompat` class, Estimator will only accept `input_fn`. Example conversion: `est = Estimator(...)` -> `est = SKCompat(Estimator(...))`

Args:

- `x` : Matrix of shape `[n_samples, n_features...]`. Can be iterator that returns arrays of features. The training input samples for fitting the model. If set, `input_fn` must be `None` .
- `input_fn` : Input function. If set, `x` and 'batch_size' must be `None` .
- `batch_size` : Override default batch size. If set, 'input_fn' must be 'None'.
- `outputs` : list of `str` , name of the output to predict. If `None` , returns all.
- `as_iterable` : If True, return an iterable which keeps yielding predictions for each example until inputs are exhausted. Note: The inputs must terminate if you want the iterable to terminate (e.g. be sure to pass `num_epochs=1` if you are using something like `read_batch_features`).

Returns:

A numpy array of predicted classes or regression values if the constructor's `model_fn` returns a `Tensor` for `predictions` or a `dict` of numpy arrays if `model_fn` returns a `dict` . Returns an iterable of predictions if `as_iterable` is True.

Raises:

- `ValueError` : If `x` and `input_fn` are both provided or both `None` .

set_params

```
set_params(**params)
```

Set the parameters of this estimator.

The method works on simple estimators as well as on nested objects (such as pipelines). The former have parameters of the form `<component>__<parameter>` so that it's possible to update each component of a nested object.

Args:

- `**params` : Parameters.

Returns:

`self`

Raises:

- `ValueError` : If `params` contain invalid names.

Except as otherwise noted, the content of this page is licensed under the [Creative Commons Attribution 3.0 License](#), and code samples are licensed under the [Apache 2.0 License](#). For details, see our [Site Policies](#). Java is a registered trademark of Oracle and/or its affiliates.

Last updated November 2, 2017.

Stay Connected

[Blog](#)

[GitHub](#)

[Twitter](#)

Support

[Issue Tracker](#)

[Release Notes](#)

[Stack Overflow](#)

English

[Terms](#) | [Privacy](#)