tf.contrib.learn.evaluate

```
evaluate(
    graph,
    output_dir,
    checkpoint_path,
    eval_dict,
    update_op=None,
    global_step_tensor=None,
    supervisor_master='',
    log_every_steps=10,
    feed_fn=None,
    max_steps=None
)
```

Defined in `tensorflow/contrib/learn/python/learn/graph_actions.py`.

See the guide: Learn (contrib) > Graph actions

Evaluate a model loaded from a checkpoint. (deprecated)

THIS FUNCTION IS DEPRECATED. It will be removed after 2017-02-15. Instructions for updating: graph_actions.py will be deleted. Use tf.train.* utilities instead. You can use learn/estimators/estimator.py as an example.

Given `graph`, a directory to write summaries to (`output_dir`), a checkpoint to restore variables from, and a `dict` of `Tensor`s to evaluate, run an eval loop for `max_steps` steps, or until an exception (generally, an end-of-input signal from a reader operation) is raised from running `eval_dict`.

In each step of evaluation, all tensors in the `eval_dict` are evaluated, and every `log_every_steps` steps, they are logged. At the very end of evaluation, a summary is evaluated (finding the summary ops using `Supervisor`'s logic) and written to `output_dir`.

Args:

- `graph`: A `Graph` to train. It is expected that this graph is not in use elsewhere.
- `output_dir`: A string containing the directory to write a summary to.
- `checkpoint_path`: A string containing the path to a checkpoint to restore. Can be `None` if the graph doesn't require loading any variables.
- `eval_dict`: A `dict` mapping string names to tensors to evaluate. It is evaluated in every logging step. The result of the final evaluation is returned. If `update_op` is None, then it's evaluated in every step. If `max_steps` is `None`, this should depend on a reader that will raise an end-of-input exception when the inputs are exhausted.
- `update_op`: A `Tensor` which is run in every step.
- `global_step_tensor`: A `Variable` containing the global step. If `None`, one is extracted from the graph using the same logic as in `Supervisor`. Used to place eval summaries on training curves.
- `supervisor_master`: The master string to use when preparing the session.
- `log_every_steps`: Integer. Output logs every `log_every_steps` evaluation steps. The logs contain the `eval_dict` and timing information.
- `feed_fn`: A function that is called every iteration to produce a `feed_dict` passed to `session.run` calls. Optional.

- `max_steps`: Integer. Evaluate `eval_dict` this many times.

Returns:

A tuple `(eval_results, global_step)`: `eval_results`: *A* `dict` *mapping* `string` *to numeric values (*`int`*,* `float`*) that are the result of running eval_dict in the last step.* `None` *if no eval steps were run.* `global_step`: The global step this evaluation corresponds to.

Raises:

- `ValueError`: if `output_dir` is empty.

Last updated November 2, 2017.