

# tf.contrib.learn.KMeansClustering

## Contents

Class `KMeansClustering`

## Properties

`config`

`model_dir`

## Class `KMeansClustering`

Inherits From: [Estimator](#)

Defined in [tensorflow/contrib/learn/python/learn/estimators/kmeans.py](#).

See the guide: [Learn \(contrib\) > Estimators](#)

An Estimator for K-Means clustering.

## Properties

### `config`

### `model_dir`

## Methods

### `__init__`

```
__init__(
    num_clusters,
    model_dir=None,
    initial_clusters=RANDOM_INIT,
    distance_metric=SQUARED_EUCLIDEAN_DISTANCE,
    random_seed=0,
    use_mini_batch=True,
    mini_batch_steps_per_iteration=1,
    kmeans_plus_plus_num_retries=2,
    relative_tolerance=None,
    config=None
)
```

Creates a model for running KMeans training and inference.

### Args:

- `num_clusters`: number of clusters to train.
- `model_dir`: the directory to save the model results and log files.

- `initial_clusters` : specifies how to initialize the clusters for training. See `clustering_ops.kmeans` for the possible values.
- `distance_metric` : the distance metric used for clustering. See `clustering_ops.kmeans` for the possible values.
- `random_seed` : Python integer. Seed for PRNG used to initialize centers.
- `use_mini_batch` : If true, use the mini-batch k-means algorithm. Else assume full batch.
- `mini_batch_steps_per_iteration` : number of steps after which the updated cluster centers are synced back to a master copy. See `clustering_ops.py` for more details.
- `kmeans_plus_plus_num_retries` : For each point that is sampled during kmeans++ initialization, this parameter specifies the number of additional points to draw from the current distribution before selecting the best. If a negative value is specified, a heuristic is used to sample  $O(\log(\text{num\_to\_sample}))$  additional points.
- `relative_tolerance` : A relative tolerance of change in the loss between iterations. Stops learning if the loss changes less than this amount. Note that this may not work correctly if `use_mini_batch=True`.
- `config` : See Estimator

## clusters

```
clusters()
```

Returns cluster centers.

## evaluate

```
evaluate(
    x=None,
    y=None,
    input_fn=None,
    feed_fn=None,
    batch_size=None,
    steps=None,
    metrics=None,
    name=None,
    checkpoint_path=None,
    hooks=None,
    log_progress=True
)
```

See `Evaluable` . (deprecated arguments)

SOME ARGUMENTS ARE DEPRECATED. They will be removed after 2016-12-01. Instructions for updating: Estimator is decoupled from Scikit Learn interface by moving into separate class `SKCompat`. Arguments `x`, `y` and `batch_size` are only available in the `SKCompat` class, Estimator will only accept `input_fn`. Example conversion: `est = Estimator(...)` -> `est = SKCompat(Estimator(...))`

Raises:

- `ValueError` : If at least one of `x` or `y` is provided, and at least one of `input_fn` or `feed_fn` is provided. Or if `metrics` is not `None` or `dict` .

## export

```

export(
    export_dir,
    input_fn=export._default_input_fn,
    input_feature_key=None,
    use_deprecated_input_fn=True,
    signature_fn=None,
    prediction_key=None,
    default_batch_size=1,
    exports_to_keep=None,
    checkpoint_path=None
)

```

Exports inference graph into given dir. (deprecated)

THIS FUNCTION IS DEPRECATED. It will be removed after 2017-03-25. Instructions for updating: Please use `Estimator.export_savedmodel()` instead.

Args:

- `export_dir`: A string containing a directory to write the exported graph and checkpoints.
- `input_fn`: If `use_deprecated_input_fn` is true, then a function that given `Tensor` of `Example` strings, parses it into features that are then passed to the model. Otherwise, a function that takes no argument and returns a tuple of (features, labels), where features is a dict of string key to `Tensor` and labels is a `Tensor` that's currently not used (and so can be `None`).
- `input_feature_key`: Only used if `use_deprecated_input_fn` is false. String key into the features dict returned by `input_fn` that corresponds to a the raw `Example` strings `Tensor` that the exported model will take as input. Can only be `None` if you're using a custom `signature_fn` that does not use the first arg (examples).
- `use_deprecated_input_fn`: Determines the signature format of `input_fn`.
- `signature_fn`: Function that returns a default signature and a named signature map, given `Tensor` of `Example` strings, `dict` of `Tensor`s for features and `Tensor` or `dict` of `Tensor`s for predictions.
- `prediction_key`: The key for a tensor in the `predictions` dict (output from the `model_fn`) to use as the `predictions` input to the `signature_fn`. Optional. If `None`, predictions will pass to `signature_fn` without filtering.
- `default_batch_size`: Default batch size of the `Example` placeholder.
- `exports_to_keep`: Number of exports to keep.
- `checkpoint_path`: the checkpoint path of the model to be exported. If it is `None` (which is default), will use the latest checkpoint in `export_dir`.

Returns:

The string path to the exported directory. NB: this functionality was added ca. 2016/09/25; clients that depend on the return value may need to handle the case where this function returns `None` because subclasses are not returning a value.

## export\_savedmodel

```

export_savedmodel(
    export_dir_base,
    serving_input_fn,
    default_output_alternative_key=None,
    assets_extra=None,
    as_text=False,
    checkpoint_path=None,
    graph_rewrite_specs=(GraphRewriteSpec((tag_constants.SERVING,)), ()),
)

```

Exports inference graph as a SavedModel into given dir.

Args:

- `export_dir_base`: A string containing a directory to write the exported graph and checkpoints.
- `serving_input_fn`: A function that takes no argument and returns an `InputFnOps`.
- `default_output_alternative_key`: the name of the head to serve when none is specified. Not needed for single-headed models.
- `assets_extra`: A dict specifying how to populate the assets.extra directory within the exported SavedModel. Each key should give the destination path (including the filename) relative to the assets.extra directory. The corresponding value gives the full path of the source file to be copied. For example, the simple case of copying a single file without renaming it is specified as `{'my_asset_file.txt': '/path/to/my_asset_file.txt'}`.
- `as_text`: whether to write the SavedModel proto in text format.
- `checkpoint_path`: The checkpoint path to export. If None (the default), the most recent checkpoint found within the model directory is chosen.
- `graph_rewrite_specs`: an iterable of `GraphRewriteSpec`. Each element will produce a separate MetaGraphDef within the exported SavedModel, tagged and rewritten as specified. Defaults to a single entry using the default serving tag ("serve") and no rewriting.

Returns:

The string path to the exported directory.

Raises:

- `ValueError`: if an unrecognized export\_type is requested.

**fit**

```
fit(  
    x=None,  
    y=None,  
    input_fn=None,  
    steps=None,  
    batch_size=None,  
    monitors=None,  
    max_steps=None  
)
```

See `Trainable`. (deprecated arguments)

SOME ARGUMENTS ARE DEPRECATED. They will be removed after 2016-12-01. Instructions for updating: Estimator is decoupled from Scikit Learn interface by moving into separate class SKCompat. Arguments x, y and batch\_size are only available in the SKCompat class, Estimator will only accept input\_fn. Example conversion: `est = Estimator(...)` -> `est = SKCompat(Estimator(...))`

Raises:

- `ValueError`: If `x` or `y` are not `None` while `input_fn` is not `None`.
- `ValueError`: If both `steps` and `max_steps` are not `None`.

## get\_params

```
get_params(deep=True)
```

Get parameters for this estimator.

Args:

- `deep` : boolean, optional

If `True`, will return the parameters for this estimator and contained subobjects that are estimators.

Returns:

- `params` : mapping of string to any Parameter names mapped to their values.

## get\_variable\_names

```
get_variable_names()
```

Returns list of all variable names in this model.

Returns:

List of names.

## get\_variable\_value

```
get_variable_value(name)
```

Returns value of the variable given by name.

Args:

- `name` : string, name of the tensor.

Returns:

Numpy array - value of the tensor.

## partial\_fit

```
partial_fit(  
    x=None,  
    y=None,  
    input_fn=None,  
    steps=1,  
    batch_size=None,  
    monitors=None  
)
```

Incremental fit on a batch of samples. (deprecated arguments)

SOME ARGUMENTS ARE DEPRECATED. They will be removed after 2016-12-01. Instructions for updating: Estimator is decoupled from Scikit Learn interface by moving into separate class SKCompat. Arguments x, y and batch\_size are only available in the SKCompat class, Estimator will only accept input\_fn. Example conversion: est = Estimator(...) -> est = SKCompat(Estimator(...))

This method is expected to be called several times consecutively on different or the same chunks of the dataset. This either can implement iterative training or out-of-core/online training.

This is especially useful when the whole dataset is too big to fit in memory at the same time. Or when model is taking long time to converge, and you want to split up training into subparts.

#### Args:

- **x** : Matrix of shape [n\_samples, n\_features...]. Can be iterator that returns arrays of features. The training input samples for fitting the model. If set, **input\_fn** must be **None** .
- **y** : Vector or matrix [n\_samples] or [n\_samples, n\_outputs]. Can be iterator that returns array of labels. The training label values (class labels in classification, real numbers in regression). If set, **input\_fn** must be **None** .
- **input\_fn** : Input function. If set, **x** , **y** , and **batch\_size** must be **None** .
- **steps** : Number of steps for which to train model. If **None** , train forever.
- **batch\_size** : minibatch size to use on the input, defaults to first dimension of **x** . Must be **None** if **input\_fn** is provided.
- **monitors** : List of **BaseMonitor** subclass instances. Used for callbacks inside the training loop.

#### Returns:

**self** , for chaining.

#### Raises:

- **ValueError** : If at least one of **x** and **y** is provided, and **input\_fn** is provided.

### predict

```
predict(  
    x=None,  
    input_fn=None,  
    batch_size=None,  
    outputs=None,  
    as_iterable=True  
)
```

Returns predictions for given features. (deprecated arguments)

SOME ARGUMENTS ARE DEPRECATED. They will be removed after 2016-12-01. Instructions for updating: Estimator is decoupled from Scikit Learn interface by moving into separate class SKCompat. Arguments x, y and batch\_size are only available in the SKCompat class, Estimator will only accept input\_fn. Example conversion: est = Estimator(...) -> est = SKCompat(Estimator(...))

#### Args:

- **x** : Matrix of shape [n\_samples, n\_features...]. Can be iterator that returns arrays of features. The training input samples for fitting the model. If set, **input\_fn** must be **None** .

- `input_fn`: Input function. If set, `x` and 'batch\_size' must be `None`.
- `batch_size`: Override default batch size. If set, 'input\_fn' must be 'None'.
- `outputs`: list of `str`, name of the output to predict. If `None`, returns all.
- `as_iterable`: If True, return an iterable which keeps yielding predictions for each example until inputs are exhausted.  
Note: The inputs must terminate if you want the iterable to terminate (e.g. be sure to pass num\_epochs=1 if you are using something like read\_batch\_features).

Returns:

A numpy array of predicted classes or regression values if the constructor's `model_fn` returns a `Tensor` for `predictions` or a `dict` of numpy arrays if `model_fn` returns a `dict`. Returns an iterable of predictions if `as_iterable` is True.

Raises:

- `ValueError`: If `x` and `input_fn` are both provided or both `None`.

## predict\_cluster\_idx

```
predict_cluster_idx(input_fn=None)
```

Yields predicted cluster indices.

## score

```
score(
    input_fn=None,
    steps=None
)
```

Predict total sum of distances to nearest clusters.

Note that this function is different from the corresponding one in sklearn which returns the negative of the sum of distances.

Args:

- `input_fn`: see predict.
- `steps`: see predict.

Returns:

Total sum of distances to nearest clusters.

## set\_params

```
set_params(**params)
```

Set the parameters of this estimator.

The method works on simple estimators as well as on nested objects (such as pipelines). The former have parameters of the form `<component>__<parameter>` so that it's possible to update each component of a nested object.

Args:

- `**params` : Parameters.

Returns:

self

Raises:

- `ValueError` : If params contain invalid names.

## transform

```
transform(  
    input_fn=None,  
    as_iterable=False  
)
```

Transforms each element to distances to cluster centers.

Note that this function is different from the corresponding one in sklearn. For SQUARED\_EUCLIDEAN distance metric, sklearn transform returns the EUCLIDEAN distance, while this function returns the SQUARED\_EUCLIDEAN distance.

Args:

- `input_fn` : see predict.
- `as_iterable` : see predict

Returns:

Array with same number of rows as x, and num\_clusters columns, containing distances to the cluster centers.

## Class Members

---

**ALL\_SCORES**

**CLUSTERS**

**CLUSTER\_IDX**

**COSINE\_DISTANCE**

**KMEANS\_PLUS\_PLUS\_INIT**

**LOSS\_OP\_NAME**

**RANDOM\_INIT**

**SCORES**

**SQUARED\_EUCLIDEAN\_DISTANCE**



---

Except as otherwise noted, the content of this page is licensed under the [Creative Commons Attribution 3.0 License](#), and code samples are licensed under the [Apache 2.0 License](#). For details, see our [Site Policies](#). Java is a registered trademark of Oracle and/or its affiliates.

Last updated November 2, 2017.

## Stay Connected

[Blog](#)

[GitHub](#)

[Twitter](#)

## Support

[Issue Tracker](#)

[Release Notes](#)

[Stack Overflow](#)

**English**

[Terms](#) | [Privacy](#)