tf.contrib.learn.LogisticRegressor

```
LogisticRegressor(
    model_fn,
    thresholds=None,
    model_dir=None,
    config=None,
    feature_engineering_fn=None
)
```

Defined in `tensorflow/contrib/learn/python/learn/estimators/logistic_regressor.py`.

See the guide: Learn (contrib) > Estimators

Builds a logistic regression Estimator for binary classification.

This method provides a basic Estimator with some additional metrics for custom binary classification models, including AUC, precision/recall and accuracy.

Example:

```
  # See tf.contrib.learn.Estimator(...) for details on model_fn structure
  def my_model_fn(...):
    pass

  estimator = LogisticRegressor(model_fn=my_model_fn)

  # Input builders
  def input_fn_train:
    pass

  estimator.fit(input_fn=input_fn_train)
  estimator.predict(x=x)
```

Args:

- `model_fn`: Model function with the signature: `(features, labels, mode) -> (predictions, loss, train_op)`. Expects the returned predictions to be probabilities in [0.0, 1.0].
- `thresholds`: List of floating point thresholds to use for accuracy, precision, and recall metrics. If `None`, defaults to `[0.5]`.
- `model_dir`: Directory to save model parameters, graphs, etc. This can also be used to load checkpoints from the directory into a estimator to continue training a previously saved model.
- `config`: A RunConfig configuration object.
- `feature_engineering_fn`: Feature engineering function. Takes features and labels which are the output of `input_fn` and returns features and labels which will be fed into the model.

Returns:

An `Estimator` instance.

Except as otherwise noted, the content of this page is licensed under the Creative Commons Attribution 3.0 License, and code samples are licensed under the Apache 2.0 License. For details, see our Site Policies. Java is a registered trademark of Oracle and/or its affiliates.

Last updated November 2, 2017.