

tf.contrib.seq2seq.embedding_attention_seq2seq

```
embedding_attention_seq2seq(  
    encoder_inputs,  
    decoder_inputs,  
    cell,  
    num_encoder_symbols,  
    num_decoder_symbols,  
    embedding_size,  
    num_heads=1,  
    output_projection=None,  
    feed_previous=False,  
    dtype=None,  
    scope=None,  
    initial_state_attention=False  
)
```

Defined in [tensorflow/contrib/legacy_seq2seq/python/ops/seq2seq.py](#).

Embedding sequence-to-sequence model with attention.

This model first embeds `encoder_inputs` by a newly created embedding (of shape `[num_encoder_symbols x input_size]`). Then it runs an RNN to encode embedded `encoder_inputs` into a state vector. It keeps the outputs of this RNN at every step to use for attention later. Next, it embeds `decoder_inputs` by another newly created embedding (of shape `[num_decoder_symbols x input_size]`). Then it runs attention decoder, initialized with the last encoder state, on embedded `decoder_inputs` and attending to encoder outputs.



Warning: when `output_projection` is `None`, the size of the attention vectors and variables will be made proportional to `num_decoder_symbols`, can be large.

Args:

- `encoder_inputs`: A list of 1D int32 Tensors of shape `[batch_size]`.
- `decoder_inputs`: A list of 1D int32 Tensors of shape `[batch_size]`.
- `cell`: `tf.nn.rnn_cell.RNNCell` defining the cell function and size.
- `num_encoder_symbols`: Integer; number of symbols on the encoder side.
- `num_decoder_symbols`: Integer; number of symbols on the decoder side.
- `embedding_size`: Integer, the length of the embedding vector for each symbol.
- `num_heads`: Number of attention heads that read from `attention_states`.
- `output_projection`: `None` or a pair (W, B) of output projection weights and biases; W has shape `[output_size x num_decoder_symbols]` and B has shape `[num_decoder_symbols]`; if provided and `feed_previous=True`, each fed previous output will first be multiplied by W and added B.
- `feed_previous`: Boolean or scalar Boolean Tensor; if True, only the first of `decoder_inputs` will be used (the "GO" symbol), and all other decoder inputs will be taken from previous outputs (as in `embedding_rnn_decoder`). If False, `decoder_inputs` are used as given (the standard decoder case).
- `dtype`: The dtype of the initial RNN state (default: `tf.float32`).
- `scope`: `VariableScope` for the created subgraph; defaults to "embedding_attention_seq2seq".

- `initial_state_attention` : If False (default), initial attentions are zero. If True, initialize the attentions from the initial state and attention states.

Returns:

A tuple of the form (outputs, state), where: `outputs` : A list of the same length as `decoder_inputs` of 2D Tensors with shape `[batch_size x num_decoder_symbols]` containing the generated outputs. `state` : The state of each decoder cell at the final time-step. It is a 2D Tensor of shape `[batch_size x cell.state_size]`.

Except as otherwise noted, the content of this page is licensed under the [Creative Commons Attribution 3.0 License](#), and code samples are licensed under the [Apache 2.0 License](#). For details, see our [Site Policies](#). Java is a registered trademark of Oracle and/or its affiliates.

Last updated November 2, 2017.

Stay Connected

[Blog](#)
[GitHub](#)
[Twitter](#)

Support

[Issue Tracker](#)
[Release Notes](#)
[Stack Overflow](#)

English

[Terms](#) | [Privacy](#)