tf.contrib.losses.mean_squared_error

```
mean_squared_error(
    predictions,
    labels=None,
    weights=1.0,
    scope=None
)
```

Defined in `tensorflow/contrib/losses/python/losses/loss_ops.py`.

See the guide: Losses (contrib) > Loss operations for use in neural networks.

Adds a Sum-of-Squares loss to the training procedure. (deprecated)

THIS FUNCTION IS DEPRECATED. It will be removed after 2016-12-30. Instructions for updating: Use tf.losses.mean_squared_error instead.

`weights` acts as a coefficient for the loss. If a scalar is provided, then the loss is simply scaled by the given value. If `weights` is a tensor of size [batch_size], then the total loss for each sample of the batch is rescaled by the corresponding element in the `weights` vector. If the shape of `weights` matches the shape of `predictions`, then the loss of each measurable element of `predictions` is scaled by the corresponding value of `weights`.

Args:

- `predictions`: The predicted outputs.
- `labels`: The ground truth output tensor, same dimensions as 'predictions'.
- `weights`: Coefficients for the loss a scalar, a tensor of shape [batch_size] or a tensor whose shape matches `predictions`.
- `scope`: The scope for the operations performed in computing the loss.

Returns:

A scalar `Tensor` representing the loss value.

Raises:

- `ValueError`: If the shape of `predictions` doesn't match that of `labels` or if the shape of `weights` is invalid.

GitHub

Twitter

Support

Issue Tracker

Release Notes

Stack Overflow

English

Terms | **Privacy**