

tf.contrib.metrics.streaming_sparse_average_precision_at_k

```
streaming_sparse_average_precision_at_k(
    predictions,
    labels,
    k,
    weights=None,
    metrics_collections=None,
    updates_collections=None,
    name=None
)
```

Defined in [tensorflow/contrib/metrics/python/ops/metric_ops.py](#).

See the guide: [Metrics \(contrib\)](#) > [Metric Ops](#)

Computes average precision@k of predictions with respect to sparse labels.

See [sparse_average_precision_at_k](#) for details on formula. [weights](#) are applied to the result of [sparse_average_precision_at_k](#)

[streaming_sparse_average_precision_at_k](#) creates two local variables, [average_precision_at_<k>/total](#) and [average_precision_at_<k>/max](#), that are used to compute the frequency. This frequency is ultimately returned as [average_precision_at_<k>](#): an idempotent operation that simply divides [average_precision_at_<k>/total](#) by [average_precision_at_<k>/max](#).

For estimation of the metric over a stream of data, the function creates an [update_op](#) operation that updates these variables and returns the [precision_at_<k>](#). Internally, a [top_k](#) operation computes a [Tensor](#) indicating the top [k predictions](#). Set operations applied to [top_k](#) and [labels](#) calculate the true positives and false positives weighted by [weights](#). Then [update_op](#) increments [true_positive_at_<k>](#) and [false_positive_at_<k>](#) using these values.

If [weights](#) is [None](#), weights default to 1. Use weights of 0 to mask values.

Args:

- [predictions](#): Float [Tensor](#) with shape [D1, ... DN, num_classes] where N >= 1. Commonly, N=1 and [predictions](#) has shape [batch size, num_classes]. The final dimension contains the logit values for each class. [D1, ... DN] must match [labels](#).
- [labels](#): [int64 Tensor](#) or [SparseTensor](#) with shape [D1, ... DN, num_labels], where N >= 1 and num_labels is the number of target classes for the associated prediction. Commonly, N=1 and [labels](#) has shape [batch_size, num_labels]. [D1, ... DN] must match [predictions](#). Values should be in range [0, num_classes), where num_classes is the last dimension of [predictions](#). Values outside this range are ignored.
- [k](#): Integer, k for @k metric. This will calculate an average precision for range [\[1, k\]](#), as documented above.
- [weights](#): [Tensor](#) whose rank is either 0, or n-1, where n is the rank of [labels](#). If the latter, it must be broadcastable to [labels](#) (i.e., all dimensions must be either [1](#), or the same as the corresponding [labels](#) dimension).
- [metrics_collections](#): An optional list of collections that values should be added to.
- [updates_collections](#): An optional list of collections that updates should be added to.
- [name](#): Name of new update operation, and namespace for other dependent ops.

Returns:

- `mean_average_precision`: Scalar `float64 Tensor` with the mean average precision values.
- `update`: **Operation** that increments variables appropriately, and whose value matches `metric`.

Except as otherwise noted, the content of this page is licensed under the [Creative Commons Attribution 3.0 License](#), and code samples are licensed under the [Apache 2.0 License](#). For details, see our [Site Policies](#). Java is a registered trademark of Oracle and/or its affiliates.

Last updated November 2, 2017.

Stay Connected

[Blog](#)

[GitHub](#)

[Twitter](#)

Support

[Issue Tracker](#)

[Release Notes](#)

[Stack Overflow](#)

English

[Terms](#) | [Privacy](#)