

tf.contrib.opt.ExternalOptimizerInterface

Contents

Class ExternalOptimizerInterface

Methods

`__init__``minimize`Class **ExternalOptimizerInterface**Defined in `tensorflow/contrib/opt/python/training/external_optimizer.py`.

Base class for interfaces with external optimization algorithms.

Subclass this and implement `_minimize` in order to wrap a new optimization algorithm.**ExternalOptimizerInterface** should not be instantiated directly; instead use e.g. **ScipyOptimizerInterface**.

Methods

`__init__`

```
__init__(  
    loss,  
    var_list=None,  
    equalities=None,  
    inequalities=None,  
    var_to_bounds=None,  
    **optimizer_kwargs  
)
```

Initialize a new interface instance.

Args:

- `loss`: A scalar **Tensor** to be minimized.
- `var_list`: Optional **list** of **Variable** objects to update to minimize `loss`. Defaults to the list of variables collected in the graph under the key **GraphKeys.TRAINABLE_VARIABLES**.
- `equalities`: Optional **list** of equality constraint scalar **Tensor**s to be held equal to zero.
- `inequalities`: Optional **list** of inequality constraint scalar **Tensor**s to be held nonnegative.
- `var_to_bounds`: Optional **dict** where each key is an optimization **Variable** and each corresponding value is a length-2 tuple of **(low, high)** bounds. Although enforcing this kind of simple constraint could be accomplished with the `inequalities` arg, not all optimization algorithms support general inequality constraints, e.g. L-BFGS-B. Both `low` and `high` can either be numbers or anything convertible to a NumPy array that can be broadcast to the shape of `var` (using `np.broadcast_to`). To indicate that there is no bound, use **None** (or

`+/- np.infty`). For example, if `var` is a 2x3 matrix, then any of the following corresponding `bounds` could be supplied:

- `(0, np.infty)` : Each element of `var` held positive.
- `(-np.infty, [1, 2])` : First column less than 1, second column less than 2.
- `(-np.infty, [[1], [2], [3]])` : First row less than 1, second row less than 2, etc.
- `(-np.infty, [[1, 2, 3], [4, 5, 6]])` : Entry `var[0, 0]` less than 1, `var[0, 1]` less than 2, etc.
- `**optimizer_kwargs` : Other subclass-specific keyword arguments.

minimize

```
minimize(  
    session=None,  
    feed_dict=None,  
    fetches=None,  
    step_callback=None,  
    loss_callback=None,  
    **run_kwargs  
)
```

Minimize a scalar `Tensor`.

Variables subject to optimization are updated in-place at the end of optimization.

Note that this method does *not* just return a minimization `Op`, unlike `Optimizer.minimize()`; instead it actually performs minimization by executing commands to control a `Session`.

Args:

- `session` : A `Session` instance.
- `feed_dict` : A feed dict to be passed to calls to `session.run`.
- `fetches` : A list of `Tensor`s to fetch and supply to `loss_callback` as positional arguments.
- `step_callback` : A function to be called at each optimization step; arguments are the current values of all optimization variables flattened into a single vector.
- `loss_callback` : A function to be called every time the loss and gradients are computed, with evaluated fetches supplied as positional arguments.
- `**run_kwargs` : kwargs to pass to `session.run`.

Except as otherwise noted, the content of this page is licensed under the [Creative Commons Attribution 3.0 License](#), and code samples are licensed under the [Apache 2.0 License](#). For details, see our [Site Policies](#). Java is a registered trademark of Oracle and/or its affiliates.

Last updated November 2, 2017.

Stay Connected

[Blog](#)

[GitHub](#)

[Twitter](#)

Support

[Issue Tracker](#)

[Release Notes](#)

[Stack Overflow](#)

English

[Terms](#) | [Privacy](#)