

## tf.distributions.Beta

## Contents

## Class Beta

## Aliases:

## Properties

## allow\_nan\_stats

Class **Beta**

Inherits From: [Distribution](#)

## Aliases:

- Class `tf.contrib.distributions.Beta`
- Class `tf.distributions.Beta`

Defined in [tensorflow/python/ops/distributions/beta.py](#).

See the guide: [Statistical Distributions \(contrib\) > Univariate \(scalar\) distributions](#)

Beta distribution.

The Beta distribution is defined over the `(0, 1)` interval using parameters `concentration1` (aka "alpha") and `concentration0` (aka "beta").

## Mathematical Details

The probability density function (pdf) is,

```
pdf(x; alpha, beta) = x**(alpha - 1) (1 - x)**(beta - 1) / Z
Z = Gamma(alpha) Gamma(beta) / Gamma(alpha + beta)
```

where:

- `concentration1 = alpha`,
- `concentration0 = beta`,
- `Z` is the normalization constant, and,
- `Gamma` is the [gamma function](#).

The concentration parameters represent mean total counts of a `1` or a `0`, i.e.,

```
concentration1 = alpha = mean * total_concentration
concentration0 = beta = (1. - mean) * total_concentration
```

where `mean` in `(0, 1)` and `total_concentration` is a positive real number representing a mean `total_count = concentration1 + concentration0`.

Distribution parameters are automatically broadcast in all functions; see examples for details.

## Examples

```
# Create a batch of three Beta distributions.
alpha = [1, 2, 3]
beta = [1, 2, 3]
dist = Beta(alpha, beta)

dist.sample([4, 5]) # Shape [4, 5, 3]

# `x` has three batch entries, each with two samples.
x = [[.1, .4, .5],
      [.2, .3, .5]]
# Calculate the probability of each pair of samples under the corresponding
# distribution in `dist`.
dist.prob(x) # Shape [2, 3]
```

```
# Create batch_shape=[2, 3] via parameter broadcast:
alpha = [[1.], [2]] # Shape [2, 1]
beta = [3., 4, 5] # Shape [3]
dist = Beta(alpha, beta)

# alpha broadcast as: [[1., 1, 1, ],
#                      [2, 2, 2]]
# beta broadcast as: [[3., 4, 5],
#                     [3, 4, 5]]
# batch_shape [2, 3]
dist.sample([4, 5]) # Shape [4, 5, 2, 3]

x = [.2, .3, .5]
# x will be broadcast as [[.2, .3, .5],
#                         [.2, .3, .5]],
# thus matching batch_shape [2, 3].
dist.prob(x) # Shape [2, 3]
```

## Properties

### allow\_nan\_stats

Python `bool` describing behavior when a stat is undefined.

Stats return +/- infinity when it makes sense. E.g., the variance of a Cauchy distribution is infinity. However, sometimes the statistic is undefined, e.g., if a distribution's pdf does not achieve a maximum within the support of the distribution, the mode is undefined. If the mean is undefined, then by definition the variance is undefined. E.g. the mean for Student's T for  $df = 1$  is undefined (no clear way to say it is either + or - infinity), so the variance =  $E[(X - \text{mean})^2]$  is also undefined.

Returns:

- `allow_nan_stats`: Python `bool`.

### batch\_shape

Shape of a single sample from a single event index as a `TensorShape`.

May be partially defined or unknown.

The batch dimensions are indexes into independent, non-identical parameterizations of this distribution.

Returns:

- `batch_shape`: `TensorShape`, possibly unknown.

### **concentration0**

Concentration parameter associated with a `0` outcome.

### **concentration1**

Concentration parameter associated with a `1` outcome.

### **dtype**

The `DType` of `Tensor`s handled by this `Distribution`.

### **event\_shape**

Shape of a single sample from a single batch as a `TensorShape`.

May be partially defined or unknown.

Returns:

- `event_shape`: `TensorShape`, possibly unknown.

### **name**

Name prepended to all ops created by this `Distribution`.

### **parameters**

Dictionary of parameters used to instantiate this `Distribution`.

### **reparameterization\_type**

Describes how samples from the distribution are reparameterized.

Currently this is one of the static instances `distributions.FULLY_REPARAMETERIZED` or `distributions.NOT_REPARAMETERIZED`.

Returns:

An instance of `ReparameterizationType`.

### **total\_concentration**

Sum of concentration parameters.

## validate\_args

Python `bool` indicating possibly expensive checks are enabled.

## Methods

---

### `__init__`

```
__init__(
    concentration1=None,
    concentration0=None,
    validate_args=False,
    allow_nan_stats=True,
    name='Beta'
)
```

Initialize a batch of Beta distributions.

Args:

- `concentration1`: Positive floating-point `Tensor` indicating mean number of successes; aka "alpha". Implies `self.dtype` and `self.batch_shape`, i.e., `concentration1.shape = [N1, N2, ..., Nm] = self.batch_shape`.
- `concentration0`: Positive floating-point `Tensor` indicating mean number of failures; aka "beta". Otherwise has same semantics as `concentration1`.
- `validate_args`: Python `bool`, default `False`. When `True` distribution parameters are checked for validity despite possibly degrading runtime performance. When `False` invalid inputs may silently render incorrect outputs.
- `allow_nan_stats`: Python `bool`, default `True`. When `True`, statistics (e.g., mean, mode, variance) use the value "`NaN`" to indicate the result is undefined. When `False`, an exception is raised if one or more of the statistic's batch members are undefined.
- `name`: Python `str` name prefixed to Ops created by this class.

## batch\_shape\_tensor

```
batch_shape_tensor(name='batch_shape_tensor')
```

Shape of a single sample from a single event index as a 1-D `Tensor`.

The batch dimensions are indexes into independent, non-identical parameterizations of this distribution.

Args:

- `name`: name to give to the op

Returns:

- `batch_shape`: `Tensor`.

## cdf

```
cdf(  
    value,  
    name='cdf'  
)
```

Cumulative distribution function.

Given random variable  $X$ , the cumulative distribution function `cdf` is:

```
cdf(x) := P[X <= x]
```

Additional documentation from `Beta` :

★ **Note:**  $x$  must have dtype `self.dtype` and be in `[0, 1]`. It must have a shape compatible with `self.batch_shape()`.

Args:

- `value`: `float` or `double Tensor`.
- `name`: The name to give this op.

Returns:

- `cdf`: a `Tensor` of shape `sample_shape(x) + self.batch_shape` with values of type `self.dtype`.

## copy

```
copy(**override_parameters_kwargs)
```

Creates a deep copy of the distribution.

★ **Note:** the copy distribution may continue to depend on the original initialization arguments.

Args:

- `**override_parameters_kwargs`: String/value dictionary of initialization arguments to override with new values.

Returns:

- `distribution`: A new instance of `type(self)` initialized from the union of `self.parameters` and `override_parameters_kwargs`, i.e., `dict(self.parameters, **override_parameters_kwargs)`.

## covariance

```
covariance(name='covariance')
```

Covariance.

Covariance is (possibly) defined only for non-scalar-event distributions.

For example, for a length-`k`, vector-valued distribution, it is calculated as,

```
Cov[i, j] = Covariance(X_i, X_j) = E[(X_i - E[X_i]) (X_j - E[X_j])]
```

where **Cov** is a (batch of)  $k \times k$  matrix,  $0 \leq (i, j) < k$ , and **E** denotes expectation.

Alternatively, for non-vector, multivariate distributions (e.g., matrix-valued, Wishart), **Covariance** shall return a (batch of) matrices under some vectorization of the events, i.e.,

```
Cov[i, j] = Covariance(Vec(X)_i, Vec(X)_j) = [as above]
```

where **Cov** is a (batch of)  $k' \times k'$  matrices,  $0 \leq (i, j) < k' = \text{reduce\_prod}(\text{event\_shape})$ , and **Vec** is some function mapping indices of this distribution's event dimensions to indices of a length- $k'$  vector.

Args:

- **name**: The name to give this op.

Returns:

- **covariance**: Floating-point **Tensor** with shape  $[B_1, \dots, B_n, k', k']$  where the first  $n$  dimensions are batch coordinates and  $k' = \text{reduce\_prod}(\text{self.event\_shape})$ .

## entropy

```
entropy(name='entropy')
```

Shannon entropy in nats.

## event\_shape\_tensor

```
event_shape_tensor(name='event_shape_tensor')
```

Shape of a single sample from a single batch as a 1-D int32 **Tensor**.

Args:

- **name**: name to give to the op

Returns:

- **event\_shape**: **Tensor**.

## is\_scalar\_batch

```
is_scalar_batch(name='is_scalar_batch')
```

Indicates that **batch\_shape** ==  $[]$ .

Args:

- **name**: The name to give this op.

Returns:

- `is_scalar_batch`: `bool` scalar `Tensor` .

## `is_scalar_event`

```
is_scalar_event(name='is_scalar_event')
```

Indicates that `event_shape == []` .

Args:

- `name` : The name to give this op.

Returns:

- `is_scalar_event`: `bool` scalar `Tensor` .

## `log_cdf`

```
log_cdf(  
    value,  
    name='log_cdf'  
)
```

Log cumulative distribution function.

Given random variable `X` , the cumulative distribution function `cdf` is:

```
log_cdf(x) := Log[ P[X <= x] ]
```

Often, a numerical approximation can be used for `log_cdf(x)` that yields a more accurate answer than simply taking the logarithm of the `cdf` when `x << -1` .

Additional documentation from `Beta` :

★ **Note:** `x` must have dtype `self.dtype` and be in `[0, 1]` . It must have a shape compatible with `self.batch_shape()` .

Args:

- `value`: `float` or `double Tensor` .
- `name` : The name to give this op.

Returns:

- `logcdf`: a `Tensor` of shape `sample_shape(x) + self.batch_shape` with values of type `self.dtype` .

## `log_prob`

```
log_prob(
    value,
    name='log_prob'
)
```

Log probability density/mass function.

Additional documentation from [Beta](#) :

★ **Note:**  $x$  must have dtype `self.dtype` and be in `[0, 1]`. It must have a shape compatible with `self.batch_shape()`.

Args:

- `value`: `float` or `double Tensor`.
- `name`: The name to give this op.

Returns:

- `log_prob`: a `Tensor` of shape `sample_shape(x) + self.batch_shape` with values of type `self.dtype`.

## log\_survival\_function

```
log_survival_function(
    value,
    name='log_survival_function'
)
```

Log survival function.

Given random variable  $X$ , the survival function is defined:

$$\begin{aligned}\text{log\_survival\_function}(x) &= \text{Log}[P[X > x]] \\ &= \text{Log}[1 - P[X \leq x]] \\ &= \text{Log}[1 - \text{cdf}(x)]\end{aligned}$$

Typically, different numerical approximations can be used for the log survival function, which are more accurate than `1 - cdf(x)` when  $x \gg 1$ .

Args:

- `value`: `float` or `double Tensor`.
- `name`: The name to give this op.

Returns:

`Tensor` of shape `sample_shape(x) + self.batch_shape` with values of type `self.dtype`.

## mean

```
mean(name='mean')
```

Mean.



## mode

```
mode(name='mode')
```

Mode.

Additional documentation from **Beta** :

★ **Note:** The mode is undefined when `concentration1 <= 1` or `concentration0 <= 1`. If `self.allow_nan_stats` is `True`, `NaN` is used for undefined modes. If `self.allow_nan_stats` is `False` an exception is raised when one or more modes are undefined.

## param\_shapes

```
param_shapes(  
    cls,  
    sample_shape,  
    name='DistributionParamShapes'  
)
```

Shapes of parameters given the desired shape of a call to `sample()`.

This is a class method that describes what key/value arguments are required to instantiate the given **Distribution** so that a particular shape is returned for that instance's call to `sample()`.

Subclasses should override class method `_param_shapes`.

Args:

- `sample_shape` : **Tensor** or python list/tuple. Desired shape of a call to `sample()`.
- `name` : name to prepend ops with.

Returns:

**dict** of parameter name to **Tensor** shapes.

## param\_static\_shapes

```
param_static_shapes(  
    cls,  
    sample_shape  
)
```

`param_shapes` with static (i.e. **TensorShape**) shapes.

This is a class method that describes what key/value arguments are required to instantiate the given **Distribution** so that a particular shape is returned for that instance's call to `sample()`. Assumes that the sample's shape is known statically.

Subclasses should override class method `_param_shapes` to return constant-valued tensors when constant values are fed.

Args:

- `sample_shape` : **TensorShape** or python list/tuple. Desired shape of a call to `sample()`.

Returns:

**dict** of parameter name to **TensorShape** .

Raises:

- **ValueError** : if **sample\_shape** is a **TensorShape** and is not fully defined.

## prob

```
prob(  
    value,  
    name='prob'  
)
```

Probability density/mass function.

Additional documentation from **Beta** :

★ **Note:** **x** must have dtype **self.dtype** and be in **[0, 1]** . It must have a shape compatible with **self.batch\_shape()** .

Args:

- **value** : **float** or **double Tensor** .
- **name** : The name to give this op.

Returns:

- **prob** : a **Tensor** of shape **sample\_shape(x) + self.batch\_shape** with values of type **self.dtype** .

## quantile

```
quantile(  
    value,  
    name='quantile'  
)
```

Quantile function. Aka "inverse cdf" or "percent point function".

Given random variable **X** and **p in [0, 1]** , the **quantile** is:

```
quantile(p) := x such that P[X <= x] == p
```

Args:

- **value** : **float** or **double Tensor** .
- **name** : The name to give this op.

Returns:

- **quantile** : a **Tensor** of shape **sample\_shape(x) + self.batch\_shape** with values of type **self.dtype** .

## sample

```
sample(  
    sample_shape=(),  
    seed=None,  
    name='sample'  
)
```

Generate samples of the specified shape.

Note that a call to `sample()` without arguments will generate a single sample.

Args:

- `sample_shape`: 0D or 1D `int32 Tensor`. Shape of the generated samples.
- `seed`: Python integer seed for RNG
- `name`: name to give to the op.

Returns:

- `samples`: a `Tensor` with prepended dimensions `sample_shape`.

## stddev

```
stddev(name='stddev')
```

Standard deviation.

Standard deviation is defined as,

$$\text{stddev} = E[(X - E[X])**2]**0.5$$

where `X` is the random variable associated with this distribution, `E` denotes expectation, and `stddev.shape = batch_shape + event_shape`.

Args:

- `name`: The name to give this op.

Returns:

- `stddev`: Floating-point `Tensor` with shape identical to `batch_shape + event_shape`, i.e., the same shape as `self.mean()`.

## survival\_function

```
survival_function(  
    value,  
    name='survival_function'  
)
```

Survival function.

Given random variable `X`, the survival function is defined:

```
survival_function(x) = P[X > x]
                    = 1 - P[X <= x]
                    = 1 - cdf(x).
```

Args:

- `value`: `float` or `double Tensor`.
- `name`: The name to give this op.

Returns:

`Tensor` of shape `sample_shape(x) + self.batch_shape` with values of type `self.dtype`.

## variance

```
variance(name='variance')
```

Variance.

Variance is defined as,

```
Var = E[(X - E[X])**2]
```

where `X` is the random variable associated with this distribution, `E` denotes expectation, and `Var.shape = batch_shape + event_shape`.

Args:

- `name`: The name to give this op.

Returns:

- `variance`: Floating-point `Tensor` with shape identical to `batch_shape + event_shape`, i.e., the same shape as `self.mean()`.

---

Except as otherwise noted, the content of this page is licensed under the [Creative Commons Attribution 3.0 License](#), and code samples are licensed under the [Apache 2.0 License](#). For details, see our [Site Policies](#). Java is a registered trademark of Oracle and/or its affiliates.

Last updated November 2, 2017.

### Stay Connected

[Blog](#)

[GitHub](#)

[Twitter](#)

### Support

[Issue Tracker](#)

[Release Notes](#)

[Stack Overflow](#)

**English**

[Terms](#) | [Privacy](#)