

tf.keras.callbacks.ReduceLROnPlateau

Contents

Class ReduceLROnPlateau

Methods

`__init__``in_cooldown`

Class ReduceLROnPlateau

Inherits From: [Callback](#)Defined in [tensorflow/python/keras/_impl/keras/callbacks.py](#).

Reduce learning rate when a metric has stopped improving.

Models often benefit from reducing the learning rate by a factor of 2-10 once learning stagnates. This callback monitors a quantity and if no improvement is seen for a 'patience' number of epochs, the learning rate is reduced.

Example:

```
reduce_lr = ReduceLROnPlateau(monitor='val_loss', factor=0.2,
                              patience=5, min_lr=0.001)
model.fit(X_train, Y_train, callbacks=[reduce_lr])
```

Arguments:

- `monitor`: quantity to be monitored.
- `factor`: factor by which the learning rate will be reduced. $\text{new_lr} = \text{lr} * \text{factor}$
- `patience`: number of epochs with no improvement after which learning rate will be reduced.
- `verbose`: int. 0: quiet, 1: update messages.
- `mode`: one of {auto, min, max}. In `min` mode, lr will be reduced when the quantity monitored has stopped decreasing; in `max` mode it will be reduced when the quantity monitored has stopped increasing; in `auto` mode, the direction is automatically inferred from the name of the monitored quantity.
- `epsilon`: threshold for measuring the new optimum, to only focus on significant changes.
- `cooldown`: number of epochs to wait before resuming normal operation after lr has been reduced.
- `min_lr`: lower bound on the learning rate.

Methods

`__init__`

```
__init__(
    monitor='val_loss',
    factor=0.1,
    patience=10,
    verbose=0,
    mode='auto',
    epsilon=0.0001,
    cooldown=0,
    min_lr=0
)
```

in_cooldown

```
in_cooldown()
```

on_batch_begin

```
on_batch_begin(
    batch,
    logs=None
)
```

on_batch_end

```
on_batch_end(
    batch,
    logs=None
)
```

on_epoch_begin

```
on_epoch_begin(
    epoch,
    logs=None
)
```

on_epoch_end

```
on_epoch_end(
    epoch,
    logs=None
)
```

on_train_begin

```
on_train_begin(logs=None)
```

on_train_end

```
on_train_end(logs=None)
```

set_model

```
set_model(model)
```

set_params

```
set_params(params)
```

Except as otherwise noted, the content of this page is licensed under the [Creative Commons Attribution 3.0 License](#), and code samples are licensed under the [Apache 2.0 License](#). For details, see our [Site Policies](#). Java is a registered trademark of Oracle and/or its affiliates.

Last updated November 2, 2017.

Stay Connected

[Blog](#)

[GitHub](#)

[Twitter](#)

Support

[Issue Tracker](#)

[Release Notes](#)

[Stack Overflow](#)

English

[Terms](#) | [Privacy](#)