tf.nn.quantized_avg_pool

```
quantized_avg_pool(
    input,
    min_input,
    max_input,
    ksize,
    strides,
    padding,
    name=None
)
```

Defined in `tensorflow/python/ops/gen_nn_ops.py`.

See the guide: Neural Network > Candidate Sampling

Produces the average pool of the input tensor for quantized types.

Args:

- `input`: A `Tensor`. Must be one of the following types: `qint8`, `quint8`, `qint16`, `quint16`, `qint32`. 4-D with shape `[batch, height, width, channels]`.
- `min_input`: A `Tensor` of type `float32`. The float value that the lowest quantized input value represents.
- `max_input`: A `Tensor` of type `float32`. The float value that the highest quantized input value represents.
- `ksize`: A list of `ints`. The size of the window for each dimension of the input tensor. The length must be 4 to match the number of dimensions of the input.
- `strides`: A list of `ints`. The stride of the sliding window for each dimension of the input tensor. The length must be 4 to match the number of dimensions of the input.
- `padding`: A `string` from: `"SAME", "VALID"`. The type of padding algorithm to use.
- `name`: A name for the operation (optional).

Returns:

A tuple of `Tensor` objects (output, min_output, max_output).

- `output`: A `Tensor`. Has the same type as `input`.
- `min_output`: A `Tensor` of type `float32`. The float value that the lowest quantized output value represents.
- `max_output`: A `Tensor` of type `float32`. The float value that the highest quantized output value represents.

GitHub

Twitter

Support

Issue Tracker

Release Notes

Stack Overflow

English

Terms | **Privacy**