

tf.norm

Contents

Aliases:

Aliases:

- `tf.linalg.norm`
- `tf.norm`

```
norm(  
    tensor,  
    ord='euclidean',  
    axis=None,  
    keep_dims=False,  
    name=None  
)
```

Defined in [tensorflow/python/ops/linalg_ops.py](#).

See the guide: [Math > Matrix Math Functions](#)

Computes the norm of vectors, matrices, and tensors.

This function can compute several different vector norms (the 1-norm, the Euclidean or 2-norm, the inf-norm, and in general the p-norm for $p > 0$) and matrix norms (Frobenius, 1-norm, and inf-norm).

Args:

- **tensor**: **Tensor** of types `float32`, `float64`, `complex64`, `complex128`
- **ord**: Order of the norm. Supported values are 'fro', 'euclidean', `1`, `2`, `np.inf` and any positive real number yielding the corresponding p-norm. Default is 'euclidean' which is equivalent to Frobenius norm if **tensor** is a matrix and equivalent to 2-norm for vectors. Some restrictions apply: a) The Frobenius norm **fro** is not defined for vectors, b) If **axis** is a 2-tuple (matrix norm), only 'euclidean', 'fro', `1`, `np.inf` are supported. See the description of **axis** on how to compute norms for a batch of vectors or matrices stored in a tensor.
- **axis**: If **axis** is `None` (the default), the input is considered a vector and a single vector norm is computed over the entire set of values in the tensor, i.e. `norm(tensor, ord=ord)` is equivalent to `norm(reshape(tensor, [-1]), ord=ord)`. If **axis** is a Python integer, the input is considered a batch of vectors, and **axis** determines the axis in **tensor** over which to compute vector norms. If **axis** is a 2-tuple of Python integers it is considered a batch of matrices and **axis** determines the axes in **tensor** over which to compute a matrix norm. Negative indices are supported. Example: If you are passing a tensor that can be either a matrix or a batch of matrices at runtime, pass `axis=[-2, -1]` instead of `axis=None` to make sure that matrix norms are computed.
- **keep_dims**: If True, the axis indicated in **axis** are kept with size 1. Otherwise, the dimensions in **axis** are removed from the output shape.
- **name**: The name of the op.

Returns:

- `output` : A `Tensor` of the same type as `tensor`, containing the vector or matrix norms. If `keep_dims` is `True` then the rank of `output` is equal to the rank of `tensor` . Otherwise, if `axis` is `None` the output is a scalar, if `axis` is an integer, the rank of `output` is one less than the rank of `tensor` , if `axis` is a 2-tuple the rank of `output` is two less than the rank of `tensor` .

Raises:

- `ValueError` : If `ord` or `axis` is invalid.

numpy compatibility

Mostly equivalent to `numpy.linalg.norm`. Not supported: `ord <= 0`, 2-norm for matrices, nuclear norm. Other differences: a) If `axis` is `None` , treats the flattened `tensor` as a vector regardless of rank. b) Explicitly supports 'euclidean' norm as the default, including for higher order tensors.

Except as otherwise noted, the content of this page is licensed under the [Creative Commons Attribution 3.0 License](#), and code samples are licensed under the [Apache 2.0 License](#). For details, see our [Site Policies](#). Java is a registered trademark of Oracle and/or its affiliates.

Last updated November 2, 2017.

Stay Connected

[Blog](#)

[GitHub](#)

[Twitter](#)

Support

[Issue Tracker](#)

[Release Notes](#)

[Stack Overflow](#)

English

[Terms](#) | [Privacy](#)