tf.reduce_mean

```
reduce_mean(
    input_tensor,
    axis=None,
    keep_dims=False,
    name=None,
    reduction_indices=None
)
```

Defined in `tensorflow/python/ops/math_ops.py` .

See the guide: Math > Reduction

Computes the mean of elements across dimensions of a tensor.

Reduces `input_tensor` along the dimensions given in `axis` . Unless `keep_dims` is true, the rank of the tensor is reduced by 1 for each entry in `axis` . If `keep_dims` is true, the reduced dimensions are retained with length 1.

If `axis` has no entries, all dimensions are reduced, and a tensor with a single element is returned.

For example:

```
x = tf.constant([[1., 1.], [2., 2.]])
tf.reduce_mean(x)  # 1.5
tf.reduce_mean(x, 0)  # [1.5, 1.5]
tf.reduce_mean(x, 1)  # [1., 2.]
```

Args:

- `input_tensor` : The tensor to reduce. Should have numeric type.
- `axis` : The dimensions to reduce. If `None` (the default), reduces all dimensions. Must be in the range `[-rank(input_tensor), rank(input_tensor))` .
- `keep_dims` : If true, retains reduced dimensions with length 1.
- `name` : A name for the operation (optional).
- `reduction_indices` : The old (deprecated) name for axis.

Returns:

The reduced tensor.

numpy compatibility

Equivalent to np.mean

Last updated November 2, 2017.

Stay Connected

Blog

GitHub

Twitter

Support

Issue Tracker

Release Notes

Stack Overflow

English

Terms | Privacy