

tf.contrib.bayesflow.entropy.elbo_ratio

```
elbo_ratio(
    log_p,
    q,
    z=None,
    n=None,
    seed=None,
    form=None,
    name='elbo_ratio'
)
```

Defined in [tensorflow/contrib/bayesflow/python/ops/entropy_impl.py](#).

See the guide: [BayesFlow Entropy \(contrib\) > Ops](#)

Estimate of the ratio appearing in the **ELBO** and **KL** divergence.

With $p(z) := \exp\{\log_p(z)\}$, this **Op** returns an approximation of

$$E_q[\log[p(Z) / q(Z)]]$$

The term $E_q[\log[p(Z)]]$ is always computed as a sample mean. The term $E_q[\log[q(z)]]$ can be computed with samples, or an exact formula if **q.entropy()** is defined. This is controlled with the kwarg **form**.

This log-ratio appears in different contexts:

KL[q || p]

If $\log_p(z) = \log[p(z)]$ for distribution **p**, this **Op** approximates the negative Kullback-Leibler divergence.

```
elbo_ratio(log_p, q, n=100) = -1 * KL[q || p],
KL[q || p] = E[log[q(Z)] - log[p(Z)]]
```

Note that if **p** is a **Distribution**, then **distributions.kl_divergence(q, p)** may be defined and available as an exact result.

ELBO

If $\log_p(z) = \log[p(z, x)]$ is the log joint of a distribution **p**, this is the Evidence Lower BOUND (ELBO):

$$\begin{aligned} \text{ELBO} &\sim E[\log[p(Z, x)] - \log[q(Z)]] \\ &= \log[p(x)] - \text{KL}[q || p] \\ &\leq \log[p(x)] \end{aligned}$$

User supplies either **Tensor** of samples **z**, or number of samples to draw **n**

Args:

- log_p**: Callable mapping samples from **q** to **Tensors** with shape broadcastable to **q.batch_shape**. For example, **log_p** works "just like" **q.log_prob**.

- `q`: `tf.contrib.distributions.Distribution`.
- `z`: `Tensor` of samples from `q`, produced by `q.sample(n)` for some `n`.
- `n`: Integer `Tensor`. Number of samples to generate if `z` is not provided.
- `seed`: Python integer to seed the random number generator.
- `form`: Either `ELBOForms.analytic_entropy` (use formula for entropy of `q`) or `ELBOForms.sample` (sample estimate of entropy), or `ELBOForms.default` (attempt analytic entropy, fallback on sample). Default value is `ELBOForms.default`.
- `name`: A name to give this `Op`.

Returns:

Scalar `Tensor` holding sample mean KL divergence. `shape` is the batch shape of `q`, and `dtype` is the same as `q`.

Raises:

- `ValueError`: If `form` is not handled by this function.

Except as otherwise noted, the content of this page is licensed under the [Creative Commons Attribution 3.0 License](#), and code samples are licensed under the [Apache 2.0 License](#). For details, see our [Site Policies](#). Java is a registered trademark of Oracle and/or its affiliates.

Last updated November 2, 2017.

Stay Connected

[Blog](#)

[GitHub](#)

[Twitter](#)

Support

[Issue Tracker](#)

[Release Notes](#)

[Stack Overflow](#)

English

[Terms](#) | [Privacy](#)