

tf.contrib.bayesflow.monte_carlo.expectation

```
expectation(
    f,
    samples,
    log_prob=None,
    use_reparametrization=True,
    axis=0,
    keep_dims=False,
    name=None
)
```

Defined in [tensorflow/contrib/bayesflow/python/ops/monte_carlo_impl.py](#).

See the guide: [BayesFlow Monte Carlo \(contrib\) > Ops](#)

Computes the Monte-Carlo approximation of $E_p[f(X)]$.

This function computes the Monte-Carlo approximation of an expectation, i.e.,

$$E_p[f(X)] \approx \frac{1}{m} \sum_{j=1}^m f(x_j), \quad x_j \sim \text{iid } p(X)$$

where:

- $x_j = \text{samples}[j, \dots]$,
- $\log(p(\text{samples})) = \text{log_prob}(\text{samples})$ and
- $m = \text{prod}(\text{shape}(\text{samples})[\text{axis}])$.

Tricks: Reparameterization and Score-Gradient

When p is "reparameterized", i.e., a diffeomorphic transformation of a parameterless distribution (e.g., $\text{Normal}(Y; \mathbf{m}, \mathbf{s}) \Leftrightarrow Y = \mathbf{s}X + \mathbf{m}, X \sim \text{Normal}(\mathbf{0}, \mathbf{1})$), we can swap gradient and expectation, i.e., $\text{grad}[\text{Avg}\{s_i : i=1\dots n\}] = \text{Avg}\{\text{grad}[s_i] : i=1\dots n\}$ where $S_n = \text{Avg}\{s_i\}$ and $s_i = f(x_i), x_i \sim p$.

However, if p is not reparameterized, TensorFlow's gradient will be incorrect since the chain-rule stops at samples of non-reparameterized distributions. (The non-differentiated result, **approx_expectation**, is the same regardless of **use_reparametrization**.) In this circumstance using the Score-Gradient trick results in an unbiased gradient, i.e.,

```
grad[ E_p[f(X)] ]
= grad[ int dx p(x) f(x) ]
= int dx grad[ p(x) f(x) ]
= int dx [ p'(x) f(x) + p(x) f'(x) ]
= int dx p(x) [p'(x) / p(x) f(x) + f'(x) ]
= int dx p(x) grad[ f(x) p(x) / stop_grad[p(x)] ]
= E_p[ grad[ f(x) p(x) / stop_grad[p(x)] ] ]
```

Unless p is not reparameterized, it is usually preferable to **use_reparametrization = True**.

⚠ Warning: users are responsible for verifying p is a "reparameterized" distribution.

Example Use:

```

bf = tf.contrib.bayesflow
ds = tf.contrib.distributions

# Monte-Carlo approximation of a reparameterized distribution, e.g., Normal.

num_draws = int(1e5)
p = ds.Normal(loc=0., scale=1.)
q = ds.Normal(loc=1., scale=2.)
exact_kl_normal_normal = ds.kl_divergence(p, q)
# ==> 0.44314718
approx_kl_normal_normal = bf.expectation(
    f=lambda x: p.log_prob(x) - q.log_prob(x),
    samples=p.sample(num_draws, seed=42),
    log_prob=p.log_prob,
    use_reparameterization=(p.reparameterization_type
                             == distribution.FULLY_REPARAMETERIZED))
# ==> 0.44632751
# Relative Error: <1%

# Monte-Carlo approximation of non-reparameterized distribution, e.g., Gamma.

num_draws = int(1e5)
p = ds.Gamma(concentration=1., rate=1.)
q = ds.Gamma(concentration=2., rate=3.)
exact_kl_gamma_gamma = ds.kl_divergence(p, q)
# ==> 0.37999129
approx_kl_gamma_gamma = bf.expectation(
    f=lambda x: p.log_prob(x) - q.log_prob(x),
    samples=p.sample(num_draws, seed=42),
    log_prob=p.log_prob,
    use_reparameterization=(p.reparameterization_type
                             == distribution.FULLY_REPARAMETERIZED))
# ==> 0.37696719
# Relative Error: <1%

# For comparing the gradients, see `monte_carlo_test.py`.

```

★ **Note:** The above example is for illustration only. To compute approximate KL-divergence, the following is preferred:

```

approx_kl_p_q = bf.monte_carlo_csiszar_f_divergence(
    f=bf.kl_reverse,
    p_log_prob=q.log_prob,
    q=p,
    num_draws=num_draws)

```

Args:

- `f`: Python callable which can return `f(samples)`.
- `samples`: `Tensor` of samples used to form the Monte-Carlo approximation of $E_p[f(X)]$. A batch of samples should be indexed by `axis` dimensions.
- `log_prob`: Python callable which can return `log_prob(samples)`. Must correspond to the natural-logarithm of the pdf/pmf of each sample. Only required/used if `use_reparameterization=False`. Default value: `None`.
- `use_reparameterization`: Python `bool` indicating that the approximation should use the fact that the gradient of samples is unbiased. Whether `True` or `False`, this arg only affects the gradient of the resulting `approx_expectation`. Default value: `True`.
- `axis`: The dimensions to average. If `None`, averages all dimensions. Default value: `0` (the left-most dimension).
- `keep_dims`: If `True`, retains averaged dimensions using size `1`. Default value: `False`.

- `name` : A `name_scope` for operations created by this function. Default value: `None` (which implies "expectation").

Returns:

- `approx_expectation` : `Tensor` corresponding to the Monte-Carlo approximation of `Ep[f(X)]` .

Raises:

- `ValueError` : if `f` is not a Python `callable` .
- `ValueError` : if `use_reparametrization=False` and `log_prob` is not a Python `callable` .

Except as otherwise noted, the content of this page is licensed under the [Creative Commons Attribution 3.0 License](#), and code samples are licensed under the [Apache 2.0 License](#). For details, see our [Site Policies](#). Java is a registered trademark of Oracle and/or its affiliates.

Last updated November 2, 2017.

Stay Connected

Blog

GitHub

Twitter

Support

Issue Tracker

Release Notes

Stack Overflow

English

[Terms](#) | [Privacy](#)