

## tf.contrib.cloud.BigQueryReader

### Contents

Class `BigQueryReader`

### Properties

`reader_ref`

`supports_serialize`

## Class `BigQueryReader`

Inherits From: `ReaderBase`

Defined in `tensorflow/contrib/cloud/python/ops/bigquery_reader_ops.py`.

A Reader that outputs keys and `tf.Example` values from a BigQuery table.

Example use:

```
# Assume a BigQuery has the following schema,
#   name      STRING,
#   age       INT,
#   state     STRING

# Create the parse_examples list of features.
features = dict(
    name=tf.FixedLenFeature([1], tf.string),
    age=tf.FixedLenFeature([1], tf.int32),
    state=tf.FixedLenFeature([1], dtype=tf.string, default_value="UNK"))

# Create a Reader.
reader = bigquery_reader_ops.BigQueryReader(project_id=PROJECT,
                                             dataset_id=DATASET,
                                             table_id=TABLE,
                                             timestamp_millis=TIME,
                                             num_partitions=NUM_PARTITIONS,
                                             features=features)

# Populate a queue with the BigQuery Table partitions.
queue = tf.train.string_input_producer(reader.partitions())

# Read and parse examples.
row_id, examples_serialized = reader.read(queue)
examples = tf.parse_example(examples_serialized, features=features)

# Process the Tensors examples["name"], examples["age"], etc...
```

Note that to create a reader a snapshot timestamp is necessary. This will enable the reader to look at a consistent snapshot of the table. For more information, see 'Table Decorators' in BigQuery docs.

See `ReaderBase` for supported methods.

## Properties

---

### **reader\_ref**

Op that implements the reader.

### **supports\_serialize**

Whether the Reader implementation can serialize its state.

## Methods

---

### **\_\_init\_\_**

```
__init__(
    project_id,
    dataset_id,
    table_id,
    timestamp_millis,
    num_partitions,
    features=None,
    columns=None,
    test_end_point=None,
    name=None
)
```

Creates a BigQueryReader.

Args:

- `project_id`: GCP project ID.
- `dataset_id`: BigQuery dataset ID.
- `table_id`: BigQuery table ID.
- `timestamp_millis`: timestamp to snapshot the table in milliseconds since the epoch. Relative (negative or zero) snapshot times are not allowed. For more details, see 'Table Decorators' in BigQuery docs.
- `num_partitions`: Number of non-overlapping partitions to read from.
- `features`: parse\_example compatible dict from keys to `VarLenFeature` and `FixedLenFeature` objects. Keys are read as columns from the db.
- `columns`: list of columns to read, can be set iff features is None.
- `test_end_point`: Used only for testing purposes (optional).
- `name`: a name for the operation (optional).

Raises:

- `TypeError`: - If features is neither None nor a dict or - If columns is neither None nor a list or - If both features and columns are None or set.

### **num\_records\_produced**

```
num_records_produced(name=None)
```

Returns the number of records this reader has produced.

This is the same as the number of Read executions that have succeeded.

Args:

- `name` : A name for the operation (optional).

Returns:

An int64 Tensor.

## **num\_work\_units\_completed**

```
num_work_units_completed(name=None)
```

Returns the number of work units this reader has finished processing.

Args:

- `name` : A name for the operation (optional).

Returns:

An int64 Tensor.

## **partitions**

```
partitions(name=None)
```

Returns serialized BigQueryTablePartition messages.

These messages represent a non-overlapping division of a table for a bulk read.

Args:

- `name` : a name for the operation (optional).

Returns:

1-D string **Tensor** of serialized **BigQueryTablePartition** messages.

## **read**

```
read(  
    queue,  
    name=None  
)
```

Returns the next record (key, value) pair produced by a reader.

Will dequeue a work unit from queue if necessary (e.g. when the Reader needs to start reading from a new file since it has

finished with the previous file).

Args:

- `queue` : A Queue or a mutable string Tensor representing a handle to a Queue, with string work items.
- `name` : A name for the operation (optional).

Returns:

A tuple of Tensors (key, value). `key` : A *string scalar Tensor*. `value` : A string scalar Tensor.

## **read\_up\_to**

```
read_up_to(  
    queue,  
    num_records,  
    name=None  
)
```

Returns up to `num_records` (key, value) pairs produced by a reader.

Will dequeue a work unit from queue if necessary (e.g., when the Reader needs to start reading from a new file since it has finished with the previous file). It may return less than `num_records` even before the last batch.

Args:

- `queue` : A Queue or a mutable string Tensor representing a handle to a Queue, with string work items.
- `num_records` : Number of records to read.
- `name` : A name for the operation (optional).

Returns:

A tuple of Tensors (keys, values). `keys` : A *1-D string Tensor*. `values` : A 1-D string Tensor.

## **reset**

```
reset(name=None)
```

Restore a reader to its initial clean state.

Args:

- `name` : A name for the operation (optional).

Returns:

The created Operation.

## **restore\_state**

```
restore_state(  
    state,  
    name=None  
)
```

Restore a reader to a previously saved state.

Not all Readers support being restored, so this can produce an Unimplemented error.

Args:

- `state` : A string Tensor. Result of a `SerializeState` of a Reader with matching type.
- `name` : A name for the operation (optional).

Returns:

The created Operation.

## **serialize\_state**

```
serialize_state(name=None)
```

Produce a string tensor that encodes the state of a reader.

Not all Readers support being serialized, so this can produce an Unimplemented error.

Args:

- `name` : A name for the operation (optional).

Returns:

A string Tensor.

---

*Except as otherwise noted, the content of this page is licensed under the [Creative Commons Attribution 3.0 License](#), and code samples are licensed under the [Apache 2.0 License](#). For details, see our [Site Policies](#). Java is a registered trademark of Oracle and/or its affiliates.*

*Last updated November 2, 2017.*

### **Stay Connected**

Blog

GitHub

Twitter

### **Support**

Issue Tracker

Release Notes

Stack Overflow

English

[Terms](#) | [Privacy](#)