

## tf.contrib.distributions.RelaxedOneHotCategorical

## Contents

Class RelaxedOneHotCategorical

## Properties

allow\_nan\_stats

batch\_shape

## Class RelaxedOneHotCategorical

Inherits From: [TransformedDistribution](#)Defined in [tensorflow/contrib/distributions/python/ops/relaxed\\_onehot\\_categorical.py](#).

RelaxedOneHotCategorical distribution with temperature and logits.

The RelaxedOneHotCategorical is a distribution over random probability vectors, vectors of positive real values that sum to one, which continuously approximates a OneHotCategorical. The degree of approximation is controlled by a temperature: as the temperature goes to 0 the RelaxedOneHotCategorical becomes discrete with a distribution described by the **logits** or **probs** parameters, as the temperature goes to infinity the RelaxedOneHotCategorical becomes the constant distribution that is identically the constant vector of (1/event\_size, ..., 1/event\_size).

The RelaxedOneHotCategorical distribution was concurrently introduced as the Gumbel-Softmax (Jang et al., 2016) and Concrete (Maddison et al., 2016) distributions for use as a reparameterized continuous approximation to the **Categorical** one-hot distribution. If you use this distribution, please cite both papers.

### Examples

Creates a continuous distribution, which approximates a 3-class one-hot categorical distribution. The 2nd class is the most likely to be the largest component in samples drawn from this distribution.

```
temperature = 0.5
p = [0.1, 0.5, 0.4]
dist = RelaxedOneHotCategorical(temperature, probs=p)
```

Creates a continuous distribution, which approximates a 3-class one-hot categorical distribution. The 2nd class is the most likely to be the largest component in samples drawn from this distribution.

```
temperature = 0.5
logits = [-2, 2, 0]
dist = RelaxedOneHotCategorical(temperature, logits=logits)
```

Creates a continuous distribution, which approximates a 3-class one-hot categorical distribution. Because the temperature is very low, samples from this distribution are almost discrete, with one component almost 1 and the others nearly 0. The 2nd class is the most likely to be the largest component in samples drawn from this distribution.

```
temperature = 1e-5
logits = [-2, 2, 0]
dist = RelaxedOneHotCategorical(temperature, logits=logits)
```

Creates a continuous distribution, which approximates a 3-class one-hot categorical distribution. Because the temperature is very high, samples from this distribution are usually close to the (1/3, 1/3, 1/3) vector. The 2nd class is still the most likely to be the largest component in samples drawn from this distribution.

```
temperature = 10
logits = [-2, 2, 0]
dist = RelaxedOneHotCategorical(temperature, logits=logits)
```

Eric Jang, Shixiang Gu, and Ben Poole. Categorical Reparameterization with Gumbel-Softmax. 2016.

Chris J. Maddison, Andriy Mnih, and Yee Whye Teh. The Concrete Distribution: A Continuous Relaxation of Discrete Random Variables. 2016.

## Properties

---

### **allow\_nan\_stats**

Python `bool` describing behavior when a stat is undefined.

Stats return +/- infinity when it makes sense. E.g., the variance of a Cauchy distribution is infinity. However, sometimes the statistic is undefined, e.g., if a distribution's pdf does not achieve a maximum within the support of the distribution, the mode is undefined. If the mean is undefined, then by definition the variance is undefined. E.g. the mean for Student's T for  $df = 1$  is undefined (no clear way to say it is either + or - infinity), so the variance =  $E[(X - \text{mean})^2]$  is also undefined.

Returns:

- `allow_nan_stats`: Python `bool`.

### **batch\_shape**

Shape of a single sample from a single event index as a `TensorShape`.

May be partially defined or unknown.

The batch dimensions are indexes into independent, non-identical parameterizations of this distribution.

Returns:

- `batch_shape`: `TensorShape`, possibly unknown.

### **bijector**

Function transforming  $x \Rightarrow y$ .

### **distribution**

Base distribution,  $p(x)$ .

### **dtype**

The `DType` of `Tensor` s handled by this `Distribution` .

## `event_shape`

Shape of a single sample from a single batch as a `TensorShape` .

May be partially defined or unknown.

Returns:

- `event_shape` : `TensorShape` , possibly unknown.

## `name`

Name prepended to all ops created by this `Distribution` .

## `parameters`

Dictionary of parameters used to instantiate this `Distribution` .

## `reparameterization_type`

Describes how samples from the distribution are reparameterized.

Currently this is one of the static instances `distributions.FULLY_REPARAMETERIZED` or `distributions.NOT_REPARAMETERIZED` .

Returns:

An instance of `ReparameterizationType` .

## `validate_args`

Python `bool` indicating possibly expensive checks are enabled.

## Methods

---

### `__init__`

```
__init__(
    temperature,
    logits=None,
    probs=None,
    dtype=tf.float32,
    validate_args=False,
    allow_nan_stats=True,
    name='RelaxedOneHotCategorical'
)
```

Initialize `RelaxedOneHotCategorical` using class log-probabilities.

Args:

- `temperature`: An 0-D **Tensor**, representing the temperature of a set of RelaxedOneHotCategorical distributions. The temperature should be positive.
- `logits`: An N-D **Tensor**, `N >= 1`, representing the log probabilities of a set of RelaxedOneHotCategorical distributions. The first `N - 1` dimensions index into a batch of independent distributions and the last dimension represents a vector of logits for each class. Only one of `logits` or `probs` should be passed in.
- `probs`: An N-D **Tensor**, `N >= 1`, representing the probabilities of a set of RelaxedOneHotCategorical distributions. The first `N - 1` dimensions index into a batch of independent distributions and the last dimension represents a vector of probabilities for each class. Only one of `logits` or `probs` should be passed in.
- `dtype`: The type of the event samples (default: float32).
- `validate_args`: Unused in this distribution.
- `allow_nan_stats`: Python **bool**, default **True**. If **False**, raise an exception if a statistic (e.g. mean/mode/etc...) is undefined for any batch member. If **True**, batch members with valid parameters leading to undefined statistics will return NaN for this statistic.
- `name`: A name for this distribution (optional).

## batch\_shape\_tensor

```
batch_shape_tensor(name='batch_shape_tensor')
```

Shape of a single sample from a single event index as a 1-D **Tensor**.

The batch dimensions are indexes into independent, non-identical parameterizations of this distribution.

Args:

- `name`: name to give to the op

Returns:

- `batch_shape`: **Tensor**.

## cdf

```
cdf(
    value,
    name='cdf'
)
```

Cumulative distribution function.

Given random variable **X**, the cumulative distribution function **cdf** is:

```
cdf(x) := P[X <= x]
```

Args:

- `value`: **float** or **double Tensor**.
- `name`: The name to give this op.

Returns:

- `cdf`: a `Tensor` of shape `sample_shape(x) + self.batch_shape` with values of type `self.dtype`.

## copy

```
copy(**override_parameters_kwargs)
```

Creates a deep copy of the distribution.

★ **Note:** the copy distribution may continue to depend on the original initialization arguments.

Args:

- `**override_parameters_kwargs`: String/value dictionary of initialization arguments to override with new values.

Returns:

- `distribution`: A new instance of `type(self)` initialized from the union of `self.parameters` and `override_parameters_kwargs`, i.e., `dict(self.parameters, **override_parameters_kwargs)`.

## covariance

```
covariance(name='covariance')
```

Covariance.

Covariance is (possibly) defined only for non-scalar-event distributions.

For example, for a length-`k`, vector-valued distribution, it is calculated as,

$$\text{Cov}[i, j] = \text{Covariance}(X_i, X_j) = E[(X_i - E[X_i]) (X_j - E[X_j])]$$

where `Cov` is a (batch of) `k x k` matrix,  $0 \leq (i, j) < k$ , and `E` denotes expectation.

Alternatively, for non-vector, multivariate distributions (e.g., matrix-valued, Wishart), `Covariance` shall return a (batch of) matrices under some vectorization of the events, i.e.,

$$\text{Cov}[i, j] = \text{Covariance}(\text{Vec}(X)_i, \text{Vec}(X)_j) = [\text{as above}]$$

where `Cov` is a (batch of) `k' x k'` matrices,  $0 \leq (i, j) < k' = \text{reduce\_prod}(\text{event\_shape})$ , and `Vec` is some function mapping indices of this distribution's event dimensions to indices of a length-`k'` vector.

Args:

- `name`: The name to give this op.

Returns:

- `covariance`: Floating-point `Tensor` with shape `[B1, ..., Bn, k', k']` where the first `n` dimensions are batch coordinates and `k' = reduce_prod(self.event_shape)`.

## entropy

```
entropy(name='entropy')
```

Shannon entropy in nats.

## event\_shape\_tensor

```
event_shape_tensor(name='event_shape_tensor')
```

Shape of a single sample from a single batch as a 1-D int32 **Tensor** .

Args:

- `name` : name to give to the op

Returns:

- `event_shape` : **Tensor** .

## is\_scalar\_batch

```
is_scalar_batch(name='is_scalar_batch')
```

Indicates that `batch_shape == []` .

Args:

- `name` : The name to give this op.

Returns:

- `is_scalar_batch` : **bool** scalar **Tensor** .

## is\_scalar\_event

```
is_scalar_event(name='is_scalar_event')
```

Indicates that `event_shape == []` .

Args:

- `name` : The name to give this op.

Returns:

- `is_scalar_event` : **bool** scalar **Tensor** .

## log\_cdf

```
log_cdf(
    value,
    name='log_cdf'
)
```

Log cumulative distribution function.

Given random variable  $X$ , the cumulative distribution function **cdf** is:

$$\text{log\_cdf}(x) := \text{Log}[P[X \leq x]]$$

Often, a numerical approximation can be used for **log\_cdf(x)** that yields a more accurate answer than simply taking the logarithm of the **cdf** when  $x \ll -1$ .

Args:

- **value**: **float** or **double Tensor**.
- **name**: The name to give this op.

Returns:

- **logcdf**: a **Tensor** of shape **sample\_shape(x) + self.batch\_shape** with values of type **self.dtype**.

## log\_prob

```
log_prob(
    value,
    name='log_prob'
)
```

Log probability density/mass function.

Args:

- **value**: **float** or **double Tensor**.
- **name**: The name to give this op.

Returns:

- **log\_prob**: a **Tensor** of shape **sample\_shape(x) + self.batch\_shape** with values of type **self.dtype**.

## log\_survival\_function

```
log_survival_function(
    value,
    name='log_survival_function'
)
```

Log survival function.

Given random variable  $X$ , the survival function is defined:

```
log_survival_function(x) = Log[ P[X > x] ]
                        = Log[ 1 - P[X <= x] ]
                        = Log[ 1 - cdf(x) ]
```

Typically, different numerical approximations can be used for the log survival function, which are more accurate than `1 - cdf(x)` when `x >> 1`.

Args:

- `value`: `float` or `double Tensor`.
- `name`: The name to give this op.

Returns:

`Tensor` of shape `sample_shape(x) + self.batch_shape` with values of type `self.dtype`.

## mean

```
mean(name='mean')
```

Mean.

## mode

```
mode(name='mode')
```

Mode.

## param\_shapes

```
param_shapes(
    cls,
    sample_shape,
    name='DistributionParamShapes'
)
```

Shapes of parameters given the desired shape of a call to `sample()`.

This is a class method that describes what key/value arguments are required to instantiate the given `Distribution` so that a particular shape is returned for that instance's call to `sample()`.

Subclasses should override class method `_param_shapes`.

Args:

- `sample_shape`: `Tensor` or python list/tuple. Desired shape of a call to `sample()`.
- `name`: name to prepend ops with.

Returns:

`dict` of parameter name to `Tensor` shapes.



## param\_static\_shapes

```
param_static_shapes(  
    cls,  
    sample_shape  
)
```

param\_shapes with static (i.e. `TensorShape`) shapes.

This is a class method that describes what key/value arguments are required to instantiate the given `Distribution` so that a particular shape is returned for that instance's call to `sample()`. Assumes that the sample's shape is known statically.

Subclasses should override class method `_param_shapes` to return constant-valued tensors when constant values are fed.

Args:

- `sample_shape`: `TensorShape` or python list/tuple. Desired shape of a call to `sample()`.

Returns:

`dict` of parameter name to `TensorShape`.

Raises:

- `ValueError`: if `sample_shape` is a `TensorShape` and is not fully defined.

## prob

```
prob(  
    value,  
    name='prob'  
)
```

Probability density/mass function.

Args:

- `value`: `float` or `double Tensor`.
- `name`: The name to give this op.

Returns:

- `prob`: a `Tensor` of shape `sample_shape(x) + self.batch_shape` with values of type `self.dtype`.

## quantile

```
quantile(  
    value,  
    name='quantile'  
)
```

Quantile function. Aka "inverse cdf" or "percent point function".

Given random variable  $X$  and  $p$  in  $[0, 1]$ , the **quantile** is:

```
quantile(p) := x such that  $P[X \leq x] = p$ 
```

Args:

- **value**: **float** or **double Tensor**.
- **name**: The name to give this op.

Returns:

- **quantile**: a **Tensor** of shape **sample\_shape(x) + self.batch\_shape** with values of type **self.dtype**.

## sample

```
sample(  
    sample_shape=(),  
    seed=None,  
    name='sample'  
)
```

Generate samples of the specified shape.

Note that a call to **sample()** without arguments will generate a single sample.

Args:

- **sample\_shape**: 0D or 1D **int32 Tensor**. Shape of the generated samples.
- **seed**: Python integer seed for RNG
- **name**: name to give to the op.

Returns:

- **samples**: a **Tensor** with prepended dimensions **sample\_shape**.

## stddev

```
stddev(name='stddev')
```

Standard deviation.

Standard deviation is defined as,

$$\text{stddev} = E[(X - E[X])**2]**0.5$$

where  $X$  is the random variable associated with this distribution,  $E$  denotes expectation, and **stddev.shape = batch\_shape + event\_shape**.

Args:

- **name**: The name to give this op.

Returns:

- `stddev`: Floating-point `Tensor` with shape identical to `batch_shape + event_shape`, i.e., the same shape as `self.mean()`.

## survival\_function

```
survival_function(  
    value,  
    name='survival_function'  
)
```

Survival function.

Given random variable `X`, the survival function is defined:

```
survival_function(x) = P[X > x]  
                    = 1 - P[X <= x]  
                    = 1 - cdf(x).
```

Args:

- `value`: `float` or `double Tensor`.
- `name`: The name to give this op.

Returns:

`Tensor` of shape `sample_shape(x) + self.batch_shape` with values of type `self.dtype`.

## variance

```
variance(name='variance')
```

Variance.

Variance is defined as,

```
Var = E[(X - E[X])**2]
```

where `X` is the random variable associated with this distribution, `E` denotes expectation, and `Var.shape = batch_shape + event_shape`.

Args:

- `name`: The name to give this op.

Returns:

- `variance`: Floating-point `Tensor` with shape identical to `batch_shape + event_shape`, i.e., the same shape as `self.mean()`.

Stay Connected

- Blog
- GitHub
- Twitter

Support

- Issue Tracker
- Release Notes
- Stack Overflow

English

Terms | Privacy