

## tf.contrib.distributions.VectorDiffeomixture

## Contents

Class VectorDiffeomixture

batch\_shape\_tensor

cdf

copy

Class **VectorDiffeomixture**Inherits From: **Distribution**Defined in `tensorflow/contrib/distributions/python/ops/vector_diffeomixture.py`.

VectorDiffeomixture distribution.

The VectorDiffeomixture is an approximation to a **compound distribution**, i.e.,
$$p(x) = \int_{\mathcal{X}} q(x | v) p(v) dv$$

$$= \lim_{Q \rightarrow \infty} \sum_{i=0}^{Q-1} \text{prob}[i] q(x | \text{loc} = \sum_k \lambda[k; i] \text{loc}[k], \text{scale} = \sum_k \lambda[k; i] \text{scale}[k])$$

$$: i=0, \dots, Q-1 \}$$

where  $q(x | v)$  is a vector version of the **distribution** argument and  $p(v)$  is a SoftmaxNormal parameterized by **mix\_loc** and **mix\_scale**. The vector-ization of **distribution** entails an affine transformation of iid samples from **distribution**. The **prob** term is from quadrature and  $\lambda[k] = \text{sigmoid}(\text{mix\_loc}[k] + \sqrt{2} \text{mix\_scale}[k] \text{grid}[k])$  where the **grid** points correspond to the **prob**s.

In the non-approximation case, a draw from the mixture distribution (the "prior") represents the convex weights for different affine transformations. I.e., draw a mixing vector  $v$  (from the  $K-1$ -simplex) and let the final sample be:  $y = (\sum_k \lambda[k] \text{scale}[k]) @ x + (\sum_k \lambda[k] \text{loc}[k])$  where  $@$  denotes matrix multiplication. However, the non-approximate distribution does not have an analytical probability density function (pdf). Therefore the **VectorDiffeomixture** class implements an approximation based on **Gauss-Hermite quadrature**. I.e., in Note: although the **VectorDiffeomixture** is approximately the **SoftmaxNormal-Distribution** compound distribution, it is itself a valid distribution. It possesses a **sample**, **log\_prob**, **mean**, **covariance** which are all mutually consistent.

## Intended Use

This distribution is noteworthy because it implements a mixture of **Vector**-ized distributions yet has samples differentiable in the distribution's parameters (aka "reparameterized"). It has an analytical density function with  $O(dKQ)$  complexity.  $d$  is the vector dimensionality,  $K$  is the number of components, and  $Q$  is the number of quadrature points. These properties make it well-suited for Bayesian Variational Inference, i.e., as a surrogate family for the posterior.

For large values of **mix\_scale**, the **VectorDistribution** behaves increasingly like a discrete mixture. (In most cases this limit is only achievable by also increasing the quadrature polynomial degree,  $Q$ .)

The term **Vector** is consistent with similar named Tensorflow **Distribution**s. For more details, see the "About **Vector** distributions in Tensorflow." section.

The term **Diffeomixture** is a portmanteau of [diffeomorphism](#) and [compound mixture](#). For more details, see the "About **Diffeomixture** s and reparametrization." section.

## Mathematical Details

The **VectorDiffeomixture** approximates a SoftmaxNormal-mixed ("prior") [compound distribution](#). Using variable-substitution and [Gauss-Hermite quadrature](#) we can redefine the distribution to be a parameter-less convex combination of **K** different affine combinations of a **d** iid samples from **distribution**.

That is, defined over  $\mathbb{R}^{d \times K}$  this distribution is parameterized by a (batch of) length-**K** **mix\_loc** and **mix\_scale** vectors, a length-**K** list of (a batch of) length-**d** **loc** vectors, and a length-**K** list of **scale** **LinearOperator** s each operating on a (batch of) length-**d** vector space. Finally, a **distribution** parameter specifies the underlying base distribution which is "lifted" to become multivariate ("lifting" is the same concept as in **TransformedDistribution**).

The probability density function (pdf) is,

```
pdf(y; mix_loc, mix_scale, loc, scale, phi)
  = sum{ prob[i] phi(f_inverse(x; i)) / abs(det(interp_scale[i]))
        : i=0, ..., Q-1 }
```

where, **phi** is the base distribution pdf, and,

```
f_inverse(x; i) = inv(interp_scale[i]) @ (x - interp_loc[i])
interp_loc[i]   = sum{ lambda[k; i] loc[k]   : k=0, ..., K-1 }
interp_scale[i] = sum{ lambda[k; i] scale[k] : k=0, ..., K-1 }
```

and,

```
grid, weight = np.polynomial.hermite.hermgauss(quadrature_polynomial_degree)
prob[k]       = weight[k] / sqrt(pi)
lambda[k; i]  = sigmoid(mix_loc[k] + sqrt(2) mix_scale[k] grid[i])
```

The distribution corresponding to **phi** must be a scalar-batch, scalar-event distribution. Typically it is reparameterized. If not, it must be a function of non-trainable parameters.

WARNING: If you backprop through a VectorDiffeomixture sample and the "base" distribution is both: not **FULLY\_REPARAMETERIZED** and a function of trainable variables, then the gradient is not guaranteed correct!

## About **Vector** distributions in TensorFlow.

The **VectorDiffeomixture** is a non-standard distribution that has properties particularly useful in [variational Bayesian methods](#).

Conditioned on a draw from the SoftmaxNormal, **Y|v** is a vector whose components are linear combinations of affine transformations, thus is itself an affine transformation. Therefore **Y|v** lives in the vector space generated by vectors of affine-transformed distributions.

★ **Note:** The marginals  $Y_1 | v, \dots, Y_d | v$  are *not* generally identical to some parameterization of **distribution**. This is due to the fact that the sum of draws from **distribution** are not generally itself the same **distribution**.

## About **Diffeomixtures** and reparameterization.

The **VectorDiffeomixture** is designed to be reparameterized, i.e., its parameters are only used to transform samples from a distribution which has no trainable parameters. This property is important because backprop stops at sources of stochasticity. That is, as long as the parameters are used *after* the underlying source of stochasticity, the computed

gradient is accurate.

Reparametrization means that we can use gradient-descent (via backprop) to optimize Monte-Carlo objectives. Such objectives are a finite-sample approximation of an expectation and arise throughout scientific computing.

## Examples

```
ds = tf.contrib.distributions
la = tf.contrib.linalg

# Create two batches of VectorDiffeomixtures, one with mix_loc=[0.] and
# another with mix_loc=[1]. In both cases, `K=2` and the affine
# transformations involve:
# k=0: loc=zeros(dims)  scale=LinearOperatorScaledIdentity
# k=1: loc=[2.]*dims    scale=LinOpDiag
dims = 5
vdm = ds.VectorDiffeomixture(
    mix_loc=[[0.], [1]],
    mix_scale=[1.],
    distribution=ds.Normal(loc=0., scale=1.),
    loc=[
        None, # Equivalent to `np.zeros(dims, dtype=np.float32)`.
        np.float32([2.]*dims),
    ],
    scale=[
        la.LinearOperatorScaledIdentity(
            num_rows=dims,
            multiplier=np.float32(1.1),
            is_positive_definite=True),
        la.LinearOperatorDiag(
            diag=np.linspace(2.5, 3.5, dims, dtype=np.float32),
            is_positive_definite=True),
    ],
    validate_args=True)
```

## Properties

`allow_nan_stats`

Python `bool` describing behavior when a stat is undefined.

Stats return +/- infinity when it makes sense. E.g., the variance of a Cauchy distribution is infinity. However, sometimes the statistic is undefined, e.g., if a distribution's pdf does not achieve a maximum within the support of the distribution, the mode is undefined. If the mean is undefined, then by definition the variance is undefined. E.g. the mean for Student's T for df = 1 is undefined (no clear way to say it is either + or - infinity), so the variance =  $E[(X - \text{mean})^2]$  is also undefined.

#### Returns:

\* `allow_nan_stats`: Python `bool`.

`batch_shape`

Shape of a single sample from a single event index as a `TensorShape`.

May be partially defined or unknown.

The batch dimensions are indexes into independent, non-identical parameterizations of this distribution.

#### Returns:

\* **<b>`batch\_shape`</b>**: ``TensorShape``, possibly unknown.

**<h3 id="distribution"><code>distribution</code></h3>**

Base scalar-event, scalar-batch distribution.

**<h3 id="dtype"><code>dtype</code></h3>**

The ``DType`` of ``Tensor``s handled by this ``Distribution``.

**<h3 id="endpoint\_affine"><code>endpoint\_affine</code></h3>**

Affine transformation for each of ``K`` components.

**<h3 id="event\_shape"><code>event\_shape</code></h3>**

Shape of a single sample from a single batch as a ``TensorShape``.

May be partially defined or unknown.

#### Returns:

\* **<b>`event\_shape`</b>**: ``TensorShape``, possibly unknown.

**<h3 id="interpolate\_weight"><code>interpolate\_weight</code></h3>**

Grid of mixing probabilities, one for each grid point.

**<h3 id="interpolated\_affine"><code>interpolated\_affine</code></h3>**

Affine transformation for each convex combination of ``K`` components.

**<h3 id="mixture\_distribution"><code>mixture\_distribution</code></h3>**

Distribution used to select a convex combination of affine transforms.

**<h3 id="name"><code>name</code></h3>**

Name prepended to all ops created by this ``Distribution``.

**<h3 id="parameters"><code>parameters</code></h3>**

Dictionary of parameters used to instantiate this ``Distribution``.

**<h3 id="quadrature\_polynomial\_degree"><code>quadrature\_polynomial\_degree</code></h3>**

Polynomial largest exponent used for Gauss-Hermite quadrature.

**<h3 id="reparameterization\_type"><code>reparameterization\_type</code></h3>**

Describes how samples from the distribution are reparameterized.

Currently this is one of the static instances  
``distributions.FULLY_REPARAMETERIZED``  
or ``distributions.NOT_REPARAMETERIZED``.

#### Returns:

An instance of ``ReparameterizationType``.

**<h3 id="validate\_args"><code>validate\_args</code></h3>**

Python ``bool`` indicating possibly expensive checks are enabled.

## Methods

```
<h3 id="__init__"><code>__init__</code></h3>
```

```
``` python
__init__(
    mix_loc,
    mix_scale,
    distribution,
    loc=None,
    scale=None,
    quadrature_polynomial_degree=8,
    validate_args=False,
    allow_nan_stats=True,
    name='VectorDiffeomixture'
)
```

Constructs the VectorDiffeomixture on  $\mathbf{R}^{K \times \mathbf{d}}$ .

Args:

- `mix_loc`: `float`-like `Tensor`. Represents the `location` parameter of the SoftmaxNormal used for selecting one of the `K` affine transformations.
- `mix_scale`: `float`-like `Tensor`. Represents the `scale` parameter of the SoftmaxNormal used for selecting one of the `K` affine transformations.
- `distribution`: `tf.Distribution`-like instance. Distribution from which `d` iid samples are used as input to the selected affine transformation. Must be a scalar-batch, scalar-event distribution. Typically `distribution.reparameterization_type = FULLY_REPARAMETERIZED` or it is a function of non-trainable parameters. WARNING: If you backprop through a VectorDiffeomixture sample and the `distribution` is not `FULLY_REPARAMETERIZED` yet is a function of trainable variables, then the gradient will be incorrect!
- `loc`: Length-`K` list of `float`-type `Tensor`s. The `k`-th element represents the `shift` used for the `k`-th affine transformation. If the `k`-th item is `None`, `loc` is implicitly `0`. When specified, must have shape `[B1, ..., Bb, d]` where `b >= 0` and `d` is the event size.
- `scale`: Length-`K` list of `LinearOperator`s. Each should be positive-definite and operate on a `d`-dimensional vector space. The `k`-th element represents the `scale` used for the `k`-th affine transformation. `LinearOperator`s must have shape `[B1, ..., Bb, d, d]`, `b >= 0`, i.e., characterizes `b`-batches of `d x d` matrices
- `quadrature_polynomial_degree`: Python `int`-like scalar.
- `validate_args`: Python `bool`, default `False`. When `True` distribution parameters are checked for validity despite possibly degrading runtime performance. When `False` invalid inputs may silently render incorrect outputs.
- `allow_nan_stats`: Python `bool`, default `True`. When `True`, statistics (e.g., mean, mode, variance) use the value "NaN" to indicate the result is undefined. When `False`, an exception is raised if one or more of the statistic's batch members are undefined.
- `name`: Python `str` name prefixed to Ops created by this class.

Raises:

- `ValueError`: if `not scale or len(scale) < 2`.
- `ValueError`: if `len(loc) != len(scale)`
- `ValueError`: if `quadrature_polynomial_degree < 1`.
- `ValueError`: if `validate_args` and any not `scale.is_positive_definite`.
- `TypeError`: if any `scale.dtype != scale[0].dtype`.
- `TypeError`: if any `loc.dtype != scale[0].dtype`.

- `NotImplementedError`: if `len(scale) != 2`.
- `ValueError`: if `not distribution.is_scalar_batch`.
- `ValueError`: if `not distribution.is_scalar_event`.

## batch\_shape\_tensor

```
batch_shape_tensor(name='batch_shape_tensor')
```

Shape of a single sample from a single event index as a 1-D **Tensor**.

The batch dimensions are indexes into independent, non-identical parameterizations of this distribution.

Args:

- `name`: name to give to the op

Returns:

- `batch_shape`: **Tensor**.

## cdf

```
cdf(
    value,
    name='cdf'
)
```

Cumulative distribution function.

Given random variable **X**, the cumulative distribution function **cdf** is:

```
cdf(x) := P[X <= x]
```

Args:

- `value`: **float** or **double Tensor**.
- `name`: The name to give this op.

Returns:

- `cdf`: a **Tensor** of shape `sample_shape(x) + self.batch_shape` with values of type `self.dtype`.

## copy

```
copy(**override_parameters_kwargs)
```

Creates a deep copy of the distribution.

★ **Note:** the copy distribution may continue to depend on the original initialization arguments.

Args:

- `**override_parameters_kwargs`: String/value dictionary of initialization arguments to override with new values.

Returns:

- `distribution`: A new instance of `type(self)` initialized from the union of `self.parameters` and `override_parameters_kwargs`, i.e., `dict(self.parameters, **override_parameters_kwargs)`.

## covariance

```
covariance(name='covariance')
```

Covariance.

Covariance is (possibly) defined only for non-scalar-event distributions.

For example, for a length-`k`, vector-valued distribution, it is calculated as,

```
Cov[i, j] = Covariance(X_i, X_j) = E[(X_i - E[X_i]) (X_j - E[X_j])]
```

where `Cov` is a (batch of) `k x k` matrix, `0 <= (i, j) < k`, and `E` denotes expectation.

Alternatively, for non-vector, multivariate distributions (e.g., matrix-valued, Wishart), `Covariance` shall return a (batch of) matrices under some vectorization of the events, i.e.,

```
Cov[i, j] = Covariance(Vec(X)_i, Vec(X)_j) = [as above]
```

where `Cov` is a (batch of) `k' x k'` matrices, `0 <= (i, j) < k' = reduce_prod(event_shape)`, and `Vec` is some function mapping indices of this distribution's event dimensions to indices of a length-`k'` vector.

Args:

- `name`: The name to give this op.

Returns:

- `covariance`: Floating-point `Tensor` with shape `[B1, ..., Bn, k', k']` where the first `n` dimensions are batch coordinates and `k' = reduce_prod(self.event_shape)`.

## entropy

```
entropy(name='entropy')
```

Shannon entropy in nats.

## event\_shape\_tensor

```
event_shape_tensor(name='event_shape_tensor')
```

Shape of a single sample from a single batch as a 1-D int32 `Tensor`.

Args:

- `name` : name to give to the op

Returns:

- `event_shape` : `Tensor` .

## **is\_scalar\_batch**

```
is_scalar_batch(name='is_scalar_batch')
```

Indicates that `batch_shape == []` .

Args:

- `name` : The name to give this op.

Returns:

- `is_scalar_batch` : `bool` scalar `Tensor` .

## **is\_scalar\_event**

```
is_scalar_event(name='is_scalar_event')
```

Indicates that `event_shape == []` .

Args:

- `name` : The name to give this op.

Returns:

- `is_scalar_event` : `bool` scalar `Tensor` .

## **log\_cdf**

```
log_cdf(  
    value,  
    name='log_cdf'  
)
```

Log cumulative distribution function.

Given random variable `X` , the cumulative distribution function `cdf` is:

```
log_cdf(x) := Log[ P[X <= x] ]
```

Often, a numerical approximation can be used for `log_cdf(x)` that yields a more accurate answer than simply taking the logarithm of the `cdf` when `x << -1` .



Args:

- `value`: `float` or `double Tensor` .
- `name` : The name to give this op.

Returns:

- `logcdf`: a `Tensor` of shape `sample_shape(x) + self.batch_shape` with values of type `self.dtype` .

## log\_prob

```
log_prob(  
    value,  
    name='log_prob'  
)
```

Log probability density/mass function.

Args:

- `value`: `float` or `double Tensor` .
- `name` : The name to give this op.

Returns:

- `log_prob`: a `Tensor` of shape `sample_shape(x) + self.batch_shape` with values of type `self.dtype` .

## log\_survival\_function

```
log_survival_function(  
    value,  
    name='log_survival_function'  
)
```

Log survival function.

Given random variable `X`, the survival function is defined:

```
log_survival_function(x) = Log[ P[X > x] ]  
                        = Log[ 1 - P[X <= x] ]  
                        = Log[ 1 - cdf(x) ]
```

Typically, different numerical approximations can be used for the log survival function, which are more accurate than `1 - cdf(x)` when `x >> 1` .

Args:

- `value`: `float` or `double Tensor` .
- `name` : The name to give this op.

Returns:

**Tensor** of shape `sample_shape(x) + self.batch_shape` with values of type `self.dtype`.

## mean

```
mean(name='mean')
```

Mean.

## mode

```
mode(name='mode')
```

Mode.

## param\_shapes

```
param_shapes(  
    cls,  
    sample_shape,  
    name='DistributionParamShapes'  
)
```

Shapes of parameters given the desired shape of a call to `sample()`.

This is a class method that describes what key/value arguments are required to instantiate the given **Distribution** so that a particular shape is returned for that instance's call to `sample()`.

Subclasses should override class method `_param_shapes`.

Args:

- `sample_shape`: **Tensor** or python list/tuple. Desired shape of a call to `sample()`.
- `name`: name to prepend ops with.

Returns:

**dict** of parameter name to **Tensor** shapes.

## param\_static\_shapes

```
param_static_shapes(  
    cls,  
    sample_shape  
)
```

`param_shapes` with static (i.e. **TensorShape**) shapes.

This is a class method that describes what key/value arguments are required to instantiate the given **Distribution** so that a particular shape is returned for that instance's call to `sample()`. Assumes that the sample's shape is known statically.

Subclasses should override class method `_param_shapes` to return constant-valued tensors when constant values are fed.

Args:

- `sample_shape`: `TensorShape` or python list/tuple. Desired shape of a call to `sample()`.

Returns:

`dict` of parameter name to `TensorShape`.

Raises:

- `ValueError`: if `sample_shape` is a `TensorShape` and is not fully defined.

## prob

```
prob(  
    value,  
    name='prob'  
)
```

Probability density/mass function.

Args:

- `value`: `float` or `double Tensor`.
- `name`: The name to give this op.

Returns:

- `prob`: a `Tensor` of shape `sample_shape(x) + self.batch_shape` with values of type `self.dtype`.

## quantile

```
quantile(  
    value,  
    name='quantile'  
)
```

Quantile function. Aka "inverse cdf" or "percent point function".

Given random variable `X` and `p in [0, 1]`, the `quantile` is:

```
quantile(p) := x such that P[X <= x] == p
```

Args:

- `value`: `float` or `double Tensor`.
- `name`: The name to give this op.

Returns:

- `quantile`: a `Tensor` of shape `sample_shape(x) + self.batch_shape` with values of type `self.dtype`.

## sample

```
sample(  
    sample_shape=(),  
    seed=None,  
    name='sample'  
)
```

Generate samples of the specified shape.

Note that a call to `sample()` without arguments will generate a single sample.

Args:

- `sample_shape`: 0D or 1D `int32 Tensor`. Shape of the generated samples.
- `seed`: Python integer seed for RNG
- `name`: name to give to the op.

Returns:

- `samples`: a `Tensor` with prepended dimensions `sample_shape`.

## stddev

```
stddev(name='stddev')
```

Standard deviation.

Standard deviation is defined as,

$$\text{stddev} = E[(X - E[X])**2]**0.5$$

where `X` is the random variable associated with this distribution, `E` denotes expectation, and `stddev.shape = batch_shape + event_shape`.

Args:

- `name`: The name to give this op.

Returns:

- `stddev`: Floating-point `Tensor` with shape identical to `batch_shape + event_shape`, i.e., the same shape as `self.mean()`.

## survival\_function

```
survival_function(  
    value,  
    name='survival_function'  
)
```

Survival function.

Given random variable `X`, the survival function is defined:

```
survival_function(x) = P[X > x]
                    = 1 - P[X <= x]
                    = 1 - cdf(x).
```

Args:

- `value`: `float` or `double Tensor`.
- `name`: The name to give this op.

Returns:

`Tensor` of shape `sample_shape(x) + self.batch_shape` with values of type `self.dtype`.

## variance

```
variance(name='variance')
```

Variance.

Variance is defined as,

```
Var = E[(X - E[X])**2]
```

where `X` is the random variable associated with this distribution, `E` denotes expectation, and `Var.shape = batch_shape + event_shape`.

Args:

- `name`: The name to give this op.

Returns:

- `variance`: Floating-point `Tensor` with shape identical to `batch_shape + event_shape`, i.e., the same shape as `self.mean()`.

---

Except as otherwise noted, the content of this page is licensed under the [Creative Commons Attribution 3.0 License](#), and code samples are licensed under the [Apache 2.0 License](#). For details, see our [Site Policies](#). Java is a registered trademark of Oracle and/or its affiliates.

Last updated November 2, 2017.

### Stay Connected

[Blog](#)

[GitHub](#)

[Twitter](#)

### Support

[Issue Tracker](#)

[Release Notes](#)

[Stack Overflow](#)

**English**

[Terms](#) | [Privacy](#)