

tf.contrib.distributions.VectorSinhArcsinhDiag

Contents

Class VectorSinhArcsinhDiag

Properties

allow_nan_stats

batch_shape

Class **VectorSinhArcsinhDiag**Inherits From: [TransformedDistribution](#)Defined in [tensorflow/contrib/distributions/python/ops/vector_sinh_arcsinh_diag.py](#).The (diagonal) SinhArcsinh transformation of a distribution on \mathbf{R}^k .

This distribution models a random vector $\mathbf{Y} = (Y_1, \dots, Y_k)$, making use of a **SinhArcsinh** transformation (which has adjustable tailweight and skew), a rescaling, and a shift.

The **SinhArcsinh** transformation of the Normal is described in great depth in [Sinh-arcsinh distributions](#). Here we use a slightly different parameterization, in terms of **tailweight** and **skewness**. Additionally we allow for distributions other than Normal, and control over **scale** as well as a "shift" parameter **loc**.

Mathematical Details

Given iid random vector $\mathbf{Z} = (Z_1, \dots, Z_k)$, we define the VectorSinhArcsinhDiag transformation of \mathbf{Z} , \mathbf{Y} , parameterized by **(loc, scale, skewness, tailweight)**, via the relation (with \odot denoting matrix multiplication):

$$\begin{aligned} \mathbf{Y} &:= \text{loc} + \text{scale} \odot \mathbf{F}(\mathbf{Z}) * (2 / \mathbf{F}(\mathbf{Z})) \\ \mathbf{F}(\mathbf{Z}) &:= \text{Sinh}(\text{Arcsinh}(\mathbf{Z}) + \text{skewness}) * \text{tailweight} \end{aligned}$$

This distribution is similar to the location-scale transformation $\mathbf{L}(\mathbf{Z}) := \text{loc} + \text{scale} \odot \mathbf{Z}$ in the following ways:

- If **skewness** = 0 and **tailweight** = 1 (the defaults), $\mathbf{F}(\mathbf{Z}) = \mathbf{Z}$, and then $\mathbf{Y} = \mathbf{L}(\mathbf{Z})$ exactly.
- **loc** is used in both to shift the result by a constant factor.
- Our definition of **C** ensures that $\mathbf{P}[\mathbf{Y} - \text{loc} \leq 2 * \text{scale}] = \mathbf{P}[\mathbf{L}(\mathbf{Z}) - \text{loc} \leq 2 * \text{scale}]$. Thus it can be said that the weights in the tails of \mathbf{Y} and $\mathbf{L}(\mathbf{Z})$ beyond $\text{loc} + 2 * \text{scale}$ are the same.

This distribution is different than $\text{loc} + \text{scale} \odot \mathbf{Z}$ due to the reshaping done by **F**:

- Positive (negative) **skewness** leads to positive (negative) skew.
- positive skew means, the mode of $\mathbf{F}(\mathbf{Z})$ is "tilted" to the right.
- positive skew means positive values of $\mathbf{F}(\mathbf{Z})$ become more likely, and negative values become less likely.
- Larger (smaller) **tailweight** leads to fatter (thinner) tails.
- Fatter tails mean larger values of $|\mathbf{F}(\mathbf{Z})|$ become more likely.
- **tailweight** < 1 leads to a distribution that is "flat" around $\mathbf{Y} = \text{loc}$, and a very steep drop-off in the tails.

- `tailweight > 1` leads to a distribution more peaked at the mode with heavier tails.

To see the argument about the tails, note that for $|Z| \gg 1$ and $|Z| \gg (|\text{skewness}| * \text{tailweight})^{1/\text{tailweight}}$, we have $Y \approx 0.5 Z^{2/\text{tailweight}} e^{(\text{sign}(Z) \text{ skewness} * \text{tailweight})}$.

To see the argument about `C` and quantiles, note that

$$\begin{aligned} P[(Y - \text{loc}) / \text{scale} \leq 2] &= P[F(Z) \leq 2 * \text{scale} / C] \\ &= P[Z \leq F^{-1}(2 * \text{scale} / C)] \\ &= P[Z \leq 2]. \end{aligned}$$

Properties

`allow_nan_stats`

Python `bool` describing behavior when a stat is undefined.

Stats return +/- infinity when it makes sense. E.g., the variance of a Cauchy distribution is infinity. However, sometimes the statistic is undefined, e.g., if a distribution's pdf does not achieve a maximum within the support of the distribution, the mode is undefined. If the mean is undefined, then by definition the variance is undefined. E.g. the mean for Student's T for $df = 1$ is undefined (no clear way to say it is either + or - infinity), so the variance = $E[(X - \text{mean})^2]$ is also undefined.

Returns:

- `allow_nan_stats`: Python `bool`.

`batch_shape`

Shape of a single sample from a single event index as a `TensorShape`.

May be partially defined or unknown.

The batch dimensions are indexes into independent, non-identical parameterizations of this distribution.

Returns:

- `batch_shape`: `TensorShape`, possibly unknown.

`bijector`

Function transforming $x \Rightarrow y$.

`distribution`

Base distribution, $p(x)$.

`dtype`

The `DType` of `Tensor`s handled by this `Distribution`.

`event_shape`

Shape of a single sample from a single batch as a `TensorShape`.

May be partially defined or unknown.

Returns:

- `event_shape`: `TensorShape`, possibly unknown.

loc

The `loc` in $\text{Y} := \text{loc} + \text{scale} @ \text{F}(\text{Z}) * (2 / \text{F}(2))$.

name

Name prepended to all ops created by this `Distribution`.

parameters

Dictionary of parameters used to instantiate this `Distribution`.

reparameterization_type

Describes how samples from the distribution are reparameterized.

Currently this is one of the static instances `distributions.FULLY_REPARAMETERIZED` or `distributions.NOT_REPARAMETERIZED`.

Returns:

An instance of `ReparameterizationType`.

scale

The `LinearOperator` `scale` in $\text{Y} := \text{loc} + \text{scale} @ \text{F}(\text{Z}) * (2 / \text{F}(2))$.

skewness

Controls the skewness. `Skewness > 0` means right skew.

tailweight

Controls the tail decay. `tailweight > 1` means faster than Normal.

validate_args

Python `bool` indicating possibly expensive checks are enabled.

Methods

`__init__`

```

__init__(
    loc=None,
    scale_diag=None,
    scale_identity_multiplier=None,
    skewness=None,
    tailweight=None,
    distribution=None,
    validate_args=False,
    allow_nan_stats=True,
    name='MultivariateNormalLinearOperator'
)

```

Construct VectorSinhArcsinhDiag distribution on \mathbb{R}^k .

The arguments `scale_diag` and `scale_identity_multiplier` combine to define the diagonal `scale` referred to in this class docstring:

```
scale = diag(scale_diag + scale_identity_multiplier * ones(k))
```

The `batch_shape` is the broadcast shape between `loc` and `scale` arguments.

The `event_shape` is given by last dimension of the matrix implied by `scale`. The last dimension of `loc` (if provided) must broadcast with this

Additional leading dimensions (if any) will index batches.

Args:

- `loc`: Floating-point **Tensor**. If this is set to `None`, `loc` is implicitly `0`. When specified, may have shape `[B1, ..., Bb, k]` where `b >= 0` and `k` is the event size.
- `scale_diag`: Non-zero, floating-point **Tensor** representing a diagonal matrix added to `scale`. May have shape `[B1, ..., Bb, k]`, `b >= 0`, and characterizes `b`-batches of `k x k` diagonal matrices added to `scale`. When both `scale_identity_multiplier` and `scale_diag` are `None` then `scale` is the **Identity**.
- `scale_identity_multiplier`: Non-zero, floating-point **Tensor** representing a scale-identity-matrix added to `scale`. May have shape `[B1, ..., Bb]`, `b >= 0`, and characterizes `b`-batches of scale `k x k` identity matrices added to `scale`. When both `scale_identity_multiplier` and `scale_diag` are `None` then `scale` is the **Identity**.
- `skewness`: Skewness parameter. floating-point **Tensor** with shape broadcastable with `event_shape`.
- `tailweight`: Tailweight parameter. floating-point **Tensor** with shape broadcastable with `event_shape`.
- `distribution`: **tf.Distribution**-like instance. Distribution from which `k` iid samples are used as input to transformation `F`. Default is `ds.Normal(0., 1.)`. Must be a scalar-batch, scalar-event distribution. Typically `distribution.reparameterization_type = FULLY_REPARAMETERIZED` or it is a function of non-trainable parameters. WARNING: If you backprop through a VectorSinhArcsinhDiag sample and `distribution` is not **FULLY_REPARAMETERIZED** yet is a function of trainable variables, then the gradient will be incorrect!
- `validate_args`: Python **bool**, default `False`. When `True` distribution parameters are checked for validity despite possibly degrading runtime performance. When `False` invalid inputs may silently render incorrect outputs.
- `allow_nan_stats`: Python **bool**, default `True`. When `True`, statistics (e.g., mean, mode, variance) use the value `"NaN"` to indicate the result is undefined. When `False`, an exception is raised if one or more of the statistic's batch members are undefined.
- `name`: Python **str** name prefixed to Ops created by this class.

Raises:

- **ValueError**: if at most `scale_identity_multiplier` is specified.

batch_shape_tensor

```
batch_shape_tensor(name='batch_shape_tensor')
```

Shape of a single sample from a single event index as a 1-D **Tensor** .

The batch dimensions are indexes into independent, non-identical parameterizations of this distribution.

Args:

- `name` : name to give to the op

Returns:

- `batch_shape` : **Tensor** .

cdf

```
cdf(  
    value,  
    name='cdf'  
)
```

Cumulative distribution function.

Given random variable `x`, the cumulative distribution function `cdf` is:

```
cdf(x) := P[X <= x]
```

Args:

- `value` : **float** or **double Tensor** .
- `name` : The name to give this op.

Returns:

- `cdf` : a **Tensor** of shape `sample_shape(x) + self.batch_shape` with values of type `self.dtype` .

copy

```
copy(**override_parameters_kwargs)
```

Creates a deep copy of the distribution.

★ **Note:** the copy distribution may continue to depend on the original initialization arguments.

Args:

- `**override_parameters_kwargs` : String/value dictionary of initialization arguments to override with new values.

Returns:

- `distribution`: A new instance of `type(self)` initialized from the union of `self.parameters` and `override_parameters_kwargs`, i.e., `dict(self.parameters, **override_parameters_kwargs)`.

covariance

```
covariance(name='covariance')
```

Covariance.

Covariance is (possibly) defined only for non-scalar-event distributions.

For example, for a length-`k`, vector-valued distribution, it is calculated as,

```
Cov[i, j] = Covariance(X_i, X_j) = E[(X_i - E[X_i]) (X_j - E[X_j])]
```

where `Cov` is a (batch of) `k x k` matrix, $0 \leq (i, j) < k$, and `E` denotes expectation.

Alternatively, for non-vector, multivariate distributions (e.g., matrix-valued, Wishart), `Covariance` shall return a (batch of) matrices under some vectorization of the events, i.e.,

```
Cov[i, j] = Covariance(Vec(X)_i, Vec(X)_j) = [as above]
```

where `Cov` is a (batch of) `k' x k'` matrices, $0 \leq (i, j) < k' = \text{reduce_prod}(\text{event_shape})$, and `Vec` is some function mapping indices of this distribution's event dimensions to indices of a length-`k'` vector.

Args:

- `name`: The name to give this op.

Returns:

- `covariance`: Floating-point `Tensor` with shape `[B1, ..., Bn, k', k']` where the first `n` dimensions are batch coordinates and $k' = \text{reduce_prod}(\text{self.event_shape})$.

entropy

```
entropy(name='entropy')
```

Shannon entropy in nats.

event_shape_tensor

```
event_shape_tensor(name='event_shape_tensor')
```

Shape of a single sample from a single batch as a 1-D int32 `Tensor`.

Args:

- `name`: name to give to the op

Returns:

- `event_shape`: `Tensor`.

is_scalar_batch

```
is_scalar_batch(name='is_scalar_batch')
```

Indicates that `batch_shape == []`.

Args:

- `name`: The name to give this op.

Returns:

- `is_scalar_batch`: `bool` scalar `Tensor`.

is_scalar_event

```
is_scalar_event(name='is_scalar_event')
```

Indicates that `event_shape == []`.

Args:

- `name`: The name to give this op.

Returns:

- `is_scalar_event`: `bool` scalar `Tensor`.

log_cdf

```
log_cdf(  
    value,  
    name='log_cdf'  
)
```

Log cumulative distribution function.

Given random variable `X`, the cumulative distribution function `cdf` is:

```
log_cdf(x) := Log[ P[X <= x] ]
```

Often, a numerical approximation can be used for `log_cdf(x)` that yields a more accurate answer than simply taking the logarithm of the `cdf` when `x << -1`.

Args:

- `value`: `float` or `double Tensor`.
- `name`: The name to give this op.

Returns:

- `logcdf`: a `Tensor` of shape `sample_shape(x) + self.batch_shape` with values of type `self.dtype`.

log_prob

```
log_prob(  
    value,  
    name='log_prob'  
)
```

Log probability density/mass function.

Args:

- `value`: `float` or `double Tensor`.
- `name`: The name to give this op.

Returns:

- `log_prob`: a `Tensor` of shape `sample_shape(x) + self.batch_shape` with values of type `self.dtype`.

log_survival_function

```
log_survival_function(  
    value,  
    name='log_survival_function'  
)
```

Log survival function.

Given random variable `x`, the survival function is defined:

```
log_survival_function(x) = Log[ P[X > x] ]  
                        = Log[ 1 - P[X <= x] ]  
                        = Log[ 1 - cdf(x) ]
```

Typically, different numerical approximations can be used for the log survival function, which are more accurate than `1 - cdf(x)` when `x >> 1`.

Args:

- `value`: `float` or `double Tensor`.
- `name`: The name to give this op.

Returns:

`Tensor` of shape `sample_shape(x) + self.batch_shape` with values of type `self.dtype`.

mean

```
mean(name='mean')
```

Mean.

mode


```
mode(name='mode')
```

Mode.

param_shapes

```
param_shapes(  
    cls,  
    sample_shape,  
    name='DistributionParamShapes'  
)
```

Shapes of parameters given the desired shape of a call to `sample()`.

This is a class method that describes what key/value arguments are required to instantiate the given `Distribution` so that a particular shape is returned for that instance's call to `sample()`.

Subclasses should override class method `_param_shapes`.

Args:

- `sample_shape`: `Tensor` or python list/tuple. Desired shape of a call to `sample()`.
- `name`: name to prepend ops with.

Returns:

`dict` of parameter name to `Tensor` shapes.

param_static_shapes

```
param_static_shapes(  
    cls,  
    sample_shape  
)
```

`param_shapes` with static (i.e. `TensorShape`) shapes.

This is a class method that describes what key/value arguments are required to instantiate the given `Distribution` so that a particular shape is returned for that instance's call to `sample()`. Assumes that the sample's shape is known statically.

Subclasses should override class method `_param_shapes` to return constant-valued tensors when constant values are fed.

Args:

- `sample_shape`: `TensorShape` or python list/tuple. Desired shape of a call to `sample()`.

Returns:

`dict` of parameter name to `TensorShape`.

Raises:

- `ValueError`: if `sample_shape` is a `TensorShape` and is not fully defined.

prob

```
prob(
    value,
    name='prob'
)
```

Probability density/mass function.

Args:

- `value`: `float` or `double Tensor`.
- `name`: The name to give this op.

Returns:

- `prob`: a `Tensor` of shape `sample_shape(x) + self.batch_shape` with values of type `self.dtype`.

quantile

```
quantile(
    value,
    name='quantile'
)
```

Quantile function. Aka "inverse cdf" or "percent point function".

Given random variable `X` and `p in [0, 1]`, the `quantile` is:

```
quantile(p) := x such that P[X <= x] == p
```

Args:

- `value`: `float` or `double Tensor`.
- `name`: The name to give this op.

Returns:

- `quantile`: a `Tensor` of shape `sample_shape(x) + self.batch_shape` with values of type `self.dtype`.

sample

```
sample(
    sample_shape=(),
    seed=None,
    name='sample'
)
```

Generate samples of the specified shape.

Note that a call to `sample()` without arguments will generate a single sample.

Args:

- `sample_shape` : 0D or 1D `int32 Tensor` . Shape of the generated samples.
- `seed` : Python integer seed for RNG
- `name` : name to give to the op.

Returns:

- `samples` : a `Tensor` with prepended dimensions `sample_shape` .

stddev

```
stddev(name='stddev')
```

Standard deviation.

Standard deviation is defined as,

$$\text{stddev} = E[(X - E[X])**2]**0.5$$

where X is the random variable associated with this distribution, E denotes expectation, and `stddev.shape = batch_shape + event_shape` .

Args:

- `name` : The name to give this op.

Returns:

- `stddev` : Floating-point `Tensor` with shape identical to `batch_shape + event_shape` , i.e., the same shape as `self.mean()` .

survival_function

```
survival_function(  
    value,  
    name='survival_function'  
)
```

Survival function.

Given random variable X , the survival function is defined:

$$\begin{aligned}\text{survival_function}(x) &= P[X > x] \\ &= 1 - P[X \leq x] \\ &= 1 - \text{cdf}(x).\end{aligned}$$

Args:

- `value` : `float` or `double Tensor` .
- `name` : The name to give this op.

Returns:

`Tensor` of shape `sample_shape(x) + self.batch_shape` with values of type `self.dtype`.

variance

```
variance(name='variance')
```

Variance.

Variance is defined as,

$$\text{Var} = E[(X - E[X])**2]$$

where `X` is the random variable associated with this distribution, `E` denotes expectation, and `Var.shape = batch_shape + event_shape`.

Args:

- `name`: The name to give this op.

Returns:

- `variance`: Floating-point `Tensor` with shape identical to `batch_shape + event_shape`, i.e., the same shape as `self.mean()`.

Except as otherwise noted, the content of this page is licensed under the [Creative Commons Attribution 3.0 License](#), and code samples are licensed under the [Apache 2.0 License](#). For details, see our [Site Policies](#). Java is a registered trademark of Oracle and/or its affiliates.

Last updated November 2, 2017.

Stay Connected

[Blog](#)

[GitHub](#)

[Twitter](#)

Support

[Issue Tracker](#)

[Release Notes](#)

[Stack Overflow](#)

English

[Terms](#) | [Privacy](#)