

tf.contrib.kernel_methods.KernelLinearClassifier

Contents

Class KernelLinearClassifier

Properties

config

model_dir

Class **KernelLinearClassifier**

Defined in [tensorflow/contrib/kernel_methods/python/kernel_estimators.py](#).

Linear classifier using kernel methods as feature preprocessing.

It trains a linear model after possibly mapping initial input features into a mapped space using explicit kernel mappings. Due to the kernel mappings, training a linear classifier in the mapped (output) space can detect non-linearities in the input space.

The user can provide a list of kernel mappers to be applied to all or a subset of existing feature_columns. This way, the user can effectively provide 2 types of feature columns:

- those passed as elements of feature_columns in the classifier's constructor
- those appearing as a key of the kernel_mappers dict.

If a column appears in feature_columns only, no mapping is applied to it. If it appears as a key in kernel_mappers, the corresponding kernel mappers are applied to it. Note that it is possible that a column appears in both places. Currently kernel_mappers are supported for _RealValuedColumns only.

Example usage:

```

real_column_a = real_valued_column(name='real_column_a',...)
sparse_column_b = sparse_column_with_hash_bucket(...)
kernel_mappers = {real_column_a : [RandomFourierFeatureMapper(...)]}
optimizer = ...

# real_column_a is used as a feature in both its initial and its transformed
# (mapped) form. sparse_column_b is not affected by kernel mappers.
kernel_classifier = KernelLinearClassifier(
    feature_columns=[real_column_a, sparse_column_b],
    model_dir=...,
    optimizer=optimizer,
    kernel_mappers=kernel_mappers)

# real_column_a is used as a feature in its transformed (mapped) form only.
# sparse_column_b is not affected by kernel mappers.
kernel_classifier = KernelLinearClassifier(
    feature_columns=[sparse_column_b],
    model_dir=...,
    optimizer=optimizer,
    kernel_mappers=kernel_mappers)

# Input builders
def train_input_fn: # returns x, y
    ...
def eval_input_fn: # returns x, y
    ...

kernel_classifier.fit(input_fn=train_input_fn)
kernel_classifier.evaluate(input_fn=eval_input_fn)
kernel_classifier.predict(...)

```

Input of `fit` and `evaluate` should have following features, otherwise there will be a `KeyError` :

- if `weight_column_name` is not `None`, a feature with `key=weight_column_name` whose value is a `Tensor` .
- for each `column` in `feature_columns` :
- if `column` is a `SparseColumn` , a feature with `key=column.name` whose `value` is a `SparseTensor` .
- if `column` is a `WeightedSparseColumn` , two features: the first with `key` the id column name, the second with `key` the weight column name. Both features' `value` must be a `SparseTensor` .
- if `column` is a `RealValuedColumn` , a feature with `key=column.name` whose `value` is a `Tensor` .

Properties

`config`

`model_dir`

Methods

`__init__`

```

__init__(
    feature_columns=None,
    model_dir=None,
    n_classes=2,
    weight_column_name=None,
    optimizer=None,
    kernel_mappers=None,
    config=None
)

```

Construct a **KernelLinearClassifier** estimator object.

Args:

- **feature_columns** : An iterable containing all the feature columns used by the model. All items in the set should be instances of classes derived from **FeatureColumn**.
- **model_dir** : Directory to save model parameters, graph etc. This can also be used to load checkpoints from the directory into an estimator to continue training a previously saved model.
- **n_classes** : number of label classes. Default is binary classification. Note that class labels are integers representing the class index (i.e. values from 0 to n_classes-1). For arbitrary label values (e.g. string labels), convert to class indices first.
- **weight_column_name** : A string defining feature column name representing weights. It is used to down weight or boost examples during training. It will be multiplied by the loss of the example.
- **optimizer** : The optimizer used to train the model. If specified, it should be an instance of **tf.Optimizer**. If **None**, the Ftrl optimizer is used by default.
- **kernel_mappers** : Dictionary of kernel mappers to be applied to the input features before training a (linear) model. Keys are feature columns and values are lists of mappers to be applied to the corresponding feature column. Currently only **RealValuedColumns** are supported and therefore all mappers should conform to the **DenseKernelMapper** interface (see `./mappers/dense_kernel_mapper.py`).
- **config** : **RunConfig** object to configure the runtime settings.

Returns:

A **KernelLinearClassifier** estimator.

Raises:

- **ValueError** : if n_classes < 2.
- **ValueError** : if neither feature_columns nor kernel_mappers are provided.
- **ValueError** : if mappers provided as kernel_mappers values are invalid.

evaluate

```

evaluate(
    x=None,
    y=None,
    input_fn=None,
    feed_fn=None,
    batch_size=None,
    steps=None,
    metrics=None,
    name=None,
    checkpoint_path=None,
    hooks=None,
    log_progress=True
)

```

See **Evaluable** . (deprecated arguments)

SOME ARGUMENTS ARE DEPRECATED. They will be removed after 2016-12-01. Instructions for updating: Estimator is decoupled from Scikit Learn interface by moving into separate class SKCompat. Arguments x, y and batch_size are only available in the SKCompat class, Estimator will only accept input_fn. Example conversion: est = Estimator(...) -> est = SKCompat(Estimator(...))

Raises:

- **ValueError** : If at least one of **x** or **y** is provided, and at least one of **input_fn** or **feed_fn** is provided. Or if **metrics** is not **None** or **dict** .

export

```

export(
    export_dir,
    input_fn=export._default_input_fn,
    input_feature_key=None,
    use_deprecated_input_fn=True,
    signature_fn=None,
    prediction_key=None,
    default_batch_size=1,
    exports_to_keep=None,
    checkpoint_path=None
)

```

Exports inference graph into given dir. (deprecated)

THIS FUNCTION IS DEPRECATED. It will be removed after 2017-03-25. Instructions for updating: Please use Estimator.export_savedmodel() instead.

Args:

- **export_dir** : A string containing a directory to write the exported graph and checkpoints.
- **input_fn** : If **use_deprecated_input_fn** is true, then a function that given **Tensor** of **Example** strings, parses it into features that are then passed to the model. Otherwise, a function that takes no argument and returns a tuple of (features, labels), where features is a dict of string key to **Tensor** and labels is a **Tensor** that's currently not used (and so can be **None**).
- **input_feature_key** : Only used if **use_deprecated_input_fn** is false. String key into the features dict returned by **input_fn** that corresponds to a the raw **Example** strings **Tensor** that the exported model will take as input. Can only be **None** if you're using a custom **signature_fn** that does not use the first arg (examples).
- **use_deprecated_input_fn** : Determines the signature format of **input_fn** .

- `signature_fn`: Function that returns a default signature and a named signature map, given `Tensor` of `Example` strings, `dict` of `Tensor` s for features and `Tensor` or `dict` of `Tensor` s for predictions.
- `prediction_key`: The key for a tensor in the `predictions` dict (output from the `model_fn`) to use as the `predictions` input to the `signature_fn`. Optional. If `None`, predictions will pass to `signature_fn` without filtering.
- `default_batch_size`: Default batch size of the `Example` placeholder.
- `exports_to_keep`: Number of exports to keep.
- `checkpoint_path`: the checkpoint path of the model to be exported. If it is `None` (which is default), will use the latest checkpoint in `export_dir`.

Returns:

The string path to the exported directory. NB: this functionality was added ca. 2016/09/25; clients that depend on the return value may need to handle the case where this function returns `None` because subclasses are not returning a value.

export_savedmodel

```
export_savedmodel(
    export_dir_base,
    serving_input_fn,
    default_output_alternative_key=None,
    assets_extra=None,
    as_text=False,
    checkpoint_path=None,
    graph_rewrite_specs=(GraphRewriteSpec((tag_constants.SERVING, ), ()),)
)
```

Exports inference graph as a SavedModel into given dir.

Args:

- `export_dir_base`: A string containing a directory to write the exported graph and checkpoints.
- `serving_input_fn`: A function that takes no argument and returns an `InputFnOps`.
- `default_output_alternative_key`: the name of the head to serve when none is specified. Not needed for single-headed models.
- `assets_extra`: A dict specifying how to populate the `assets.extra` directory within the exported SavedModel. Each key should give the destination path (including the filename) relative to the `assets.extra` directory. The corresponding value gives the full path of the source file to be copied. For example, the simple case of copying a single file without renaming it is specified as `{ 'my_asset_file.txt': '/path/to/my_asset_file.txt' }`.
- `as_text`: whether to write the SavedModel proto in text format.
- `checkpoint_path`: The checkpoint path to export. If `None` (the default), the most recent checkpoint found within the model directory is chosen.
- `graph_rewrite_specs`: an iterable of `GraphRewriteSpec`. Each element will produce a separate MetaGraphDef within the exported SavedModel, tagged and rewritten as specified. Defaults to a single entry using the default serving tag ("serve") and no rewriting.

Returns:

The string path to the exported directory.

Raises:

- `ValueError` : if an unrecognized `export_type` is requested.

fit

```
fit(
    x=None,
    y=None,
    input_fn=None,
    steps=None,
    batch_size=None,
    monitors=None,
    max_steps=None
)
```

See `Trainable` . (deprecated arguments)

SOME ARGUMENTS ARE DEPRECATED. They will be removed after 2016-12-01. Instructions for updating: Estimator is decoupled from Scikit Learn interface by moving into separate class SKCompat. Arguments x, y and batch_size are only available in the SKCompat class, Estimator will only accept input_fn. Example conversion: `est = Estimator(...)` -> `est = SKCompat(Estimator(...))`

Raises:

- `ValueError` : If `x` or `y` are not `None` while `input_fn` is not `None` .
- `ValueError` : If both `steps` and `max_steps` are not `None` .

get_params

```
get_params(deep=True)
```

Get parameters for this estimator.

Args:

- `deep` : boolean, optional
If `True` , will return the parameters for this estimator and contained subobjects that are estimators.

Returns:

- `params` : mapping of string to any Parameter names mapped to their values.

get_variable_names

```
get_variable_names()
```

Returns list of all variable names in this model.

Returns:

List of names.

get_variable_value

```
get_variable_value(name)
```

Returns value of the variable given by name.

Args:

- `name` : string, name of the tensor.

Returns:

Numpy array - value of the tensor.

partial_fit

```
partial_fit(  
    x=None,  
    y=None,  
    input_fn=None,  
    steps=1,  
    batch_size=None,  
    monitors=None  
)
```

Incremental fit on a batch of samples. (deprecated arguments)

SOME ARGUMENTS ARE DEPRECATED. They will be removed after 2016-12-01. Instructions for updating: Estimator is decoupled from Scikit Learn interface by moving into separate class SKCompat. Arguments `x`, `y` and `batch_size` are only available in the SKCompat class, Estimator will only accept `input_fn`. Example conversion: `est = Estimator(...)` -> `est = SKCompat(Estimator(...))`

This method is expected to be called several times consecutively on different or the same chunks of the dataset. This either can implement iterative training or out-of-core/online training.

This is especially useful when the whole dataset is too big to fit in memory at the same time. Or when model is taking long time to converge, and you want to split up training into subparts.

Args:

- `x` : Matrix of shape `[n_samples, n_features...]`. Can be iterator that returns arrays of features. The training input samples for fitting the model. If set, `input_fn` must be `None`.
- `y` : Vector or matrix `[n_samples]` or `[n_samples, n_outputs]`. Can be iterator that returns array of labels. The training label values (class labels in classification, real numbers in regression). If set, `input_fn` must be `None`.
- `input_fn` : Input function. If set, `x`, `y`, and `batch_size` must be `None`.
- `steps` : Number of steps for which to train model. If `None`, train forever.
- `batch_size` : minibatch size to use on the input, defaults to first dimension of `x`. Must be `None` if `input_fn` is provided.
- `monitors` : List of `BaseMonitor` subclass instances. Used for callbacks inside the training loop.

Returns:

`self`, for chaining.

Raises:

- `ValueError` : If at least one of `x` and `y` is provided, and `input_fn` is provided.

predict

```
predict(  
    x=None,  
    input_fn=None,  
    batch_size=None,  
    outputs=None,  
    as_iterable=True  
)
```

Returns predictions for given features. (deprecated arguments)

SOME ARGUMENTS ARE DEPRECATED. They will be removed after 2016-12-01. Instructions for updating: Estimator is decoupled from Scikit Learn interface by moving into separate class SKCompat. Arguments `x`, `y` and `batch_size` are only available in the SKCompat class, Estimator will only accept `input_fn`. Example conversion: `est = Estimator(...)` -> `est = SKCompat(Estimator(...))`

Args:

- `x` : Matrix of shape `[n_samples, n_features...]`. Can be iterator that returns arrays of features. The training input samples for fitting the model. If set, `input_fn` must be `None`.
- `input_fn` : Input function. If set, `x` and 'batch_size' must be `None`.
- `batch_size` : Override default batch size. If set, 'input_fn' must be 'None'.
- `outputs` : list of `str`, name of the output to predict. If `None`, returns all.
- `as_iterable` : If True, return an iterable which keeps yielding predictions for each example until inputs are exhausted. Note: The inputs must terminate if you want the iterable to terminate (e.g. be sure to pass `num_epochs=1` if you are using something like `read_batch_features`).

Returns:

A numpy array of predicted classes or regression values if the constructor's `model_fn` returns a `Tensor` for `predictions` or a `dict` of numpy arrays if `model_fn` returns a `dict`. Returns an iterable of predictions if `as_iterable` is True.

Raises:

- `ValueError` : If `x` and `input_fn` are both provided or both `None`.

predict_classes

```
predict_classes(input_fn=None)
```

Runs inference to determine the predicted class per instance.

Args:

- `input_fn` : The input function providing features.

Returns:

A generator of predicted classes for the features provided by `input_fn`. Each predicted class is represented by its class index (i.e. integer from 0 to `n_classes-1`)

predict_proba

```
predict_proba(input_fn=None)
```

Runs inference to determine the class probability predictions.

Args:

- `input_fn` : The input function providing features.

Returns:

A generator of predicted class probabilities for the features provided by `input_fn`.

set_params

```
set_params(**params)
```

Set the parameters of this estimator.

The method works on simple estimators as well as on nested objects (such as pipelines). The former have parameters of the form `<component>__<parameter>` so that it's possible to update each component of a nested object.

Args:

- `**params` : Parameters.

Returns:

`self`

Raises:

- `ValueError` : If params contain invalid names.

Except as otherwise noted, the content of this page is licensed under the [Creative Commons Attribution 3.0 License](#), and code samples are licensed under the [Apache 2.0 License](#). For details, see our [Site Policies](#). Java is a registered trademark of Oracle and/or its affiliates.

Last updated November 2, 2017.

Stay Connected

[Blog](#)

[GitHub](#)

[Twitter](#)

Support

[Issue Tracker](#)

[Release Notes](#)

[Stack Overflow](#)

English

[Terms](#) | [Privacy](#)