

tf.contrib.layers.layer_norm

```
layer_norm(  
    inputs,  
    center=True,  
    scale=True,  
    activation_fn=None,  
    reuse=None,  
    variables_collections=None,  
    outputs_collections=None,  
    trainable=True,  
    begin_norm_axis=1,  
    begin_params_axis=-1,  
    scope=None  
)
```

Defined in [tensorflow/contrib/layers/python/layers/layers.py](#).

See the guide: [Layers \(contrib\)](#) > Higher level ops for building neural network layers

Adds a Layer Normalization layer.

Based on the paper:

"Layer Normalization"

Jimmy Lei Ba, Jamie Ryan Kiros, Geoffrey E. Hinton

<https://arxiv.org/abs/1607.06450>.

Can be used as a normalizer function for conv2d and fully_connected.

Given a tensor **inputs** of rank **R**, moments are calculated and normalization is performed over axes **begin_norm_axis ... R - 1**. Scaling and centering, if requested, is performed over axes **begin_shift_axis .. R - 1**.

By default, **begin_norm_axis = 1** and **begin_params_axis = -1**, meaning that normalization is performed over all but the first axis (the **HWC** if **inputs** is **NHWC**), while the **beta** and **gamma** trainable parameters are calculated for the rightmost axis (the **C** if **inputs** is **NHWC**). Scaling and recentering is performed via broadcast of the **beta** and **gamma** parameters with the normalized tensor.

The shapes of **beta** and **gamma** are **inputs.shape[begin_params_axis:]**, and this part of the inputs' shape must be fully defined.

Args:

- **inputs**: A tensor having rank **R**. The normalization is performed over axes **begin_norm_axis ... R - 1** and centering and scaling parameters are calculated over **begin_params_axis ... R - 1**.
- **center**: If True, add offset of **beta** to normalized tensor. If False, **beta** is ignored.
- **scale**: If True, multiply by **gamma**. If False, **gamma** is not used. When the next layer is linear (also e.g. **nn.relu**), this can be disabled since the scaling can be done by the next layer.
- **activation_fn**: Activation function, default set to None to skip it and maintain a linear activation.
- **reuse**: Whether or not the layer and its variables should be reused. To be able to reuse the layer scope must be

given.

- `variables_collections` : Optional collections for the variables.
- `outputs_collections` : Collections to add the outputs.
- `trainable` : If `True` also add variables to the graph collection `GraphKeys.TRAINABLE_VARIABLES` (see `tf.Variable`).
- `begin_norm_axis` : The first normalization dimension: normalization will be performed along dimensions `begin_norm_axis : rank(inputs)`
- `begin_params_axis` : The first parameter (beta, gamma) dimension: scale and centering parameters will have dimensions `begin_params_axis : rank(inputs)` and will be broadcast with the normalized inputs accordingly.
- `scope` : Optional scope for `variable_scope`.

Returns:

A `Tensor` representing the output of the operation, having the same shape and dtype as `inputs`.

Raises:

- `ValueError` : If the rank of `inputs` is not known at graph build time, or if `inputs.shape[begin_params_axis:]` is not fully defined at graph build time.

Except as otherwise noted, the content of this page is licensed under the [Creative Commons Attribution 3.0 License](#), and code samples are licensed under the [Apache 2.0 License](#). For details, see our [Site Policies](#). Java is a registered trademark of Oracle and/or its affiliates.

Last updated November 2, 2017.

Stay Connected

[Blog](#)

[GitHub](#)

[Twitter](#)

Support

[Issue Tracker](#)

[Release Notes](#)

[Stack Overflow](#)

English

[Terms](#) | [Privacy](#)