

tf.contrib.legacy_seq2seq.attention_decoder

```
attention_decoder(  
    decoder_inputs,  
    initial_state,  
    attention_states,  
    cell,  
    output_size=None,  
    num_heads=1,  
    loop_function=None,  
    dtype=None,  
    scope=None,  
    initial_state_attention=False  
)
```

Defined in [tensorflow/contrib/legacy_seq2seq/python/ops/seq2seq.py](#).

RNN decoder with attention for the sequence-to-sequence model.

In this context "attention" means that, during decoding, the RNN can look up information in the additional tensor `attention_states`, and it does this by focusing on a few entries from the tensor. This model has proven to yield especially good results in a number of sequence-to-sequence tasks. This implementation is based on <http://arxiv.org/abs/1412.7449> (see below for details). It is recommended for complex sequence-to-sequence tasks.

Args:

- `decoder_inputs`: A list of 2D Tensors [batch_size x input_size].
- `initial_state`: 2D Tensor [batch_size x cell.state_size].
- `attention_states`: 3D Tensor [batch_size x attn_length x attn_size].
- `cell`: `tf.nn.rnn_cell.RNNCell` defining the cell function and size.
- `output_size`: Size of the output vectors; if None, we use `cell.output_size`.
- `num_heads`: Number of attention heads that read from `attention_states`.
- `loop_function`: If not None, this function will be applied to *i*-th output in order to generate *i*+1-th input, and `decoder_inputs` will be ignored, except for the first element ("GO" symbol). This can be used for decoding, but also for training to emulate <http://arxiv.org/abs/1506.03099>. Signature – `loop_function(prev, i) = next`
 - `prev` is a 2D Tensor of shape [batch_size x output_size],
 - `i` is an integer, the step number (when advanced control is needed),
 - `next` is a 2D Tensor of shape [batch_size x input_size].
- `dtype`: The dtype to use for the RNN initial state (default: `tf.float32`).
- `scope`: `VariableScope` for the created subgraph; default: "attention_decoder".
- `initial_state_attention`: If False (default), initial attentions are zero. If True, initialize the attentions from the initial state and attention states – useful when we wish to resume decoding from a previously stored decoder state and attention states.

Returns:

A tuple of the form (outputs, state), where: **outputs** : A list of the same length as `decoder_inputs` of 2D Tensors of shape `[batch_size x output_size]`. These represent the generated outputs. Output *i* is computed from input *i* (which is either the *i*-th element of `decoder_inputs` or `loop_function(output {i-1}, i)`) as follows. First, we run the cell on a combination of the input and previous attention masks: `cell_output, new_state = cell(linear(input, prev_attn), prev_state)`. Then, we calculate new attention masks: `new_attn = softmax(V^T * tanh(W * attention_states + U * new_state))` and then we calculate the output: `output = linear(cell_output, new_attn)`. **state** : The state of each decoder cell the final time-step. It is a 2D Tensor of shape `[batch_size x cell.state_size]`.

Raises:

- **ValueError** : when `num_heads` is not positive, there are no inputs, shapes of `attention_states` are not set, or input size cannot be inferred from the input.

Except as otherwise noted, the content of this page is licensed under the [Creative Commons Attribution 3.0 License](#), and code samples are licensed under the [Apache 2.0 License](#). For details, see our [Site Policies](#). Java is a registered trademark of Oracle and/or its affiliates.

Last updated November 2, 2017.

Stay Connected

Blog
GitHub
Twitter

Support

Issue Tracker
Release Notes
Stack Overflow

English

[Terms](#) | [Privacy](#)