

tf.contrib.legacy_seq2seq.sequence_loss

```
sequence_loss(  
    logits,  
    targets,  
    weights,  
    average_across_timesteps=True,  
    average_across_batch=True,  
    softmax_loss_function=None,  
    name=None  
)
```

Defined in [tensorflow/contrib/legacy_seq2seq/python/ops/seq2seq.py](#).

Weighted cross-entropy loss for a sequence of logits, batch-collapsed.

Args:

- **logits**: List of 2D Tensors of shape [batch_size x num_decoder_symbols].
- **targets**: List of 1D batch-sized int32 Tensors of the same length as logits.
- **weights**: List of 1D batch-sized float-Tensors of the same length as logits.
- **average_across_timesteps**: If set, divide the returned cost by the total label weight.
- **average_across_batch**: If set, divide the returned cost by the batch size.
- **softmax_loss_function**: Function (labels, logits) -> loss-batch to be used instead of the standard softmax (the default if this is None). **Note that to avoid confusion, it is required for the function to accept named arguments.**
- **name**: Optional name for this operation, defaults to "sequence_loss".

Returns:

A scalar float Tensor: The average log-perplexity per symbol (weighted).

Raises:

- **ValueError**: If len(logits) is different from len(targets) or len(weights).

Except as otherwise noted, the content of this page is licensed under the [Creative Commons Attribution 3.0 License](#), and code samples are licensed under the [Apache 2.0 License](#). For details, see our [Site Policies](#). Java is a registered trademark of Oracle and/or its affiliates.

Last updated November 2, 2017.

Stay Connected

[Blog](#)

[GitHub](#)

[Twitter](#)

Support

[Issue Tracker](#)

[Release Notes](#)

[Stack Overflow](#)

English

[Terms](#) | [Privacy](#)