tf.contrib.legacy_seq2seq.sequence_loss_by_example

```
sequence_loss_by_example(
    logits,
    targets,
    weights,
    average_across_timesteps=True,
    softmax_loss_function=None,
    name=None
)
```

Defined in `tensorflow/contrib/legacy_seq2seq/python/ops/seq2seq.py`.

Weighted cross-entropy loss for a sequence of logits (per example).

Args:

- `logits`: List of 2D Tensors of shape [batch_size x num_decoder_symbols].
- `targets`: List of 1D batch-sized int32 Tensors of the same length as logits.
- `weights`: List of 1D batch-sized float-Tensors of the same length as logits.
- `average_across_timesteps`: If set, divide the returned cost by the total label weight.
- `softmax_loss_function`: Function (labels, logits) -> loss-batch to be used instead of the standard softmax (the default if this is None). **Note that to avoid confusion, it is required for the function to accept named arguments.**
- `name`: Optional name for this operation, default: "sequence_loss_by_example".

Returns:

1D batch-sized float Tensor: The log-perplexity for each sequence.

Raises:

- `ValueError`: If len(logits) is different from len(targets) or len(weights).

Last updated November 2, 2017.