

## tf.contrib.linear\_optimizer.SDCAOptimizer

### Contents

Class SDCAOptimizer

### Properties

example\_id\_column

num\_loss\_partitions

## Class SDCAOptimizer

Defined in [tensorflow/contrib/linear\\_optimizer/python/sdca\\_optimizer.py](#).

Wrapper class for SDCA optimizer.

The wrapper is currently meant for use as an optimizer within a tf.learn Estimator.

Example usage:

```
real_feature_column = real_valued_column(...)
sparse_feature_column = sparse_column_with_hash_bucket(...)
sdca_optimizer = linear.SDCAOptimizer(example_id_column='example_id',
                                      num_loss_partitions=1,
                                      num_table_shards=1,
                                      symmetric_l2_regularization=2.0)

classifier = tf.contrib.learn.LinearClassifier(
    feature_columns=[real_feature_column, sparse_feature_column],
    weight_column_name=...,
    optimizer=sdca_optimizer)
classifier.fit(input_fn_train, steps=50)
classifier.evaluate(input_fn=input_fn_eval)
```

Here the expectation is that the `input_fn*` functions passed to train and evaluate return a pair (dict, label\_tensor) where dict has `example_id_column` as **key** whose value is a **Tensor** of shape [batch\_size] and dtype string. `num_loss_partitions` defines the number of partitions of the global loss function and should be set to **(#concurrent train ops/per worker) x (#workers)**. Convergence of (global) loss is guaranteed if `num_loss_partitions` is larger or equal to the above product. Larger values for `num_loss_partitions` lead to slower convergence. The recommended value for `num_loss_partitions` in **tf.learn** (where currently there is one process per worker) is the number of workers running the train steps. It defaults to 1 (single machine). `num_table_shards` defines the number of shards for the internal state table, typically set to match the number of parameter servers for large data sets.

## Properties

**example\_id\_column**

**num\_loss\_partitions**

**num\_table\_shards**

**symmetric\_l1\_regularization**

**symmetric\_l2\_regularization**

## Methods

---

### **\_\_init\_\_**

```
__init__(
    example_id_column,
    num_loss_partitions=1,
    num_table_shards=None,
    symmetric_l1_regularization=0.0,
    symmetric_l2_regularization=1.0
)
```

### **get\_name**

```
get_name()
```

### **get\_train\_step**

```
get_train_step(
    columns_to_variables,
    weight_column_name,
    loss_type,
    features,
    targets,
    global_step
)
```

Returns the training operation of an SdcaModel optimizer.

---

Except as otherwise noted, the content of this page is licensed under the [Creative Commons Attribution 3.0 License](#), and code samples are licensed under the [Apache 2.0 License](#). For details, see our [Site Policies](#). Java is a registered trademark of Oracle and/or its affiliates.

*Last updated November 2, 2017.*

#### **Stay Connected**

Blog

GitHub

Twitter

#### **Support**

Issue Tracker

Release Notes

Stack Overflow

