

tf.contrib.metrics.streaming_recall_at_k

```
streaming_recall_at_k(  
    predictions,  
    labels,  
    k,  
    weights=None,  
    metrics_collections=None,  
    updates_collections=None,  
    name=None  
)
```

Defined in [tensorflow/contrib/metrics/python/ops/metric_ops.py](#).

See the guide: [Metrics \(contrib\)](#) > [Metric Ops](#)

Computes the recall@k of the predictions with respect to dense labels. (deprecated)

THIS FUNCTION IS DEPRECATED. It will be removed after 2016-11-08. Instructions for updating: Please use [streaming_sparse_recall_at_k](#), and reshape labels from [batch_size] to [batch_size, 1].

The [streaming_recall_at_k](#) function creates two local variables, [total](#) and [count](#), that are used to compute the recall@k frequency. This frequency is ultimately returned as [recall_at_k](#): an idempotent operation that simply divides [total](#) by [count](#).

For estimation of the metric over a stream of data, the function creates an [update_op](#) operation that updates these variables and returns the [recall_at_k](#). Internally, an [in_top_k](#) operation computes a [Tensor](#) with shape [batch_size] whose elements indicate whether or not the corresponding label is in the top [k](#) [predictions](#). Then [update_op](#) increments [total](#) with the reduced sum of [weights](#) where [in_top_k](#) is [True](#), and it increments [count](#) with the reduced sum of [weights](#).

If [weights](#) is [None](#), weights default to 1. Use weights of 0 to mask values.

Args:

- [predictions](#): A float [Tensor](#) of dimension [batch_size, num_classes].
- [labels](#): A [Tensor](#) of dimension [batch_size] whose type is in [int32](#), [int64](#).
- [k](#): The number of top elements to look at for computing recall.
- [weights](#): [Tensor](#) whose rank is either 0, or the same rank as [labels](#), and must be broadcastable to [labels](#) (i.e., all dimensions must be either [1](#), or the same as the corresponding [labels](#) dimension).
- [metrics_collections](#): An optional list of collections that [recall_at_k](#) should be added to.
- [updates_collections](#): An optional list of collections [update_op](#) should be added to.
- [name](#): An optional variable_scope name.

Returns:

- [recall_at_k](#): A [Tensor](#) representing the recall@k, the fraction of labels which fall into the top [k](#) predictions.
- [update_op](#): An operation that increments the [total](#) and [count](#) variables appropriately and whose value matches [recall_at_k](#).

Raises:

- `ValueError`: If `predictions` and `labels` have mismatched shapes, or if `weights` is not `None` and its shape doesn't match `predictions`, or if either `metrics_collections` or `updates_collections` are not a list or tuple.

Except as otherwise noted, the content of this page is licensed under the [Creative Commons Attribution 3.0 License](#), and code samples are licensed under the [Apache 2.0 License](#). For details, see our [Site Policies](#). Java is a registered trademark of Oracle and/or its affiliates.

Last updated November 2, 2017.

Stay Connected

[Blog](#)

[GitHub](#)

[Twitter](#)

Support

[Issue Tracker](#)

[Release Notes](#)

[Stack Overflow](#)

English

[Terms](#) | [Privacy](#)