

tf.contrib.rnn.AttentionCellWrapper

Contents

Class `AttentionCellWrapper`

Properties

`activity_regularizer`

`dtype`

Class `AttentionCellWrapper`

Inherits From: `RNNCell`

Defined in `tensorflow/contrib/rnn/python/ops/rnn_cell.py`.

See the guide: [RNN and Cells \(contrib\) > Core RNN Cell wrappers \(RNNCells that wrap other RNNCells\)](#)

Basic attention cell wrapper.

Implementation based on <https://arxiv.org/abs/1409.0473>.

Properties

`activity_regularizer`

Optional regularizer function for the output of this layer.

`dtype`

`graph`

`input`

Retrieves the input tensor(s) of a layer.

Only applicable if the layer has exactly one input, i.e. if it is connected to one incoming layer.

Returns:

Input tensor or list of input tensors.

Raises:

- `AttributeError`: if the layer is connected to more than one incoming layers.

Raises:

- `RuntimeError` : If called in Eager mode.
- `AttributeError` : If no inbound nodes are found.

input_shape

Retrieves the input shape(s) of a layer.

Only applicable if the layer has exactly one input, i.e. if it is connected to one incoming layer, or if all inputs have the same shape.

Returns:

Input shape, as an integer shape tuple (or list of shape tuples, one tuple per input tensor).

Raises:

- `AttributeError` : if the layer has no defined `input_shape`.
- `RuntimeError` : if called in Eager mode.

losses

name

non_trainable_variables

non_trainable_weights

output

Retrieves the output tensor(s) of a layer.

Only applicable if the layer has exactly one output, i.e. if it is connected to one incoming layer.

Returns:

Output tensor or list of output tensors.

Raises:

- `AttributeError` : if the layer is connected to more than one incoming layers.
- `RuntimeError` : if called in Eager mode.

output_shape

Retrieves the output shape(s) of a layer.

Only applicable if the layer has one output, or if all outputs have the same shape.

Returns:

Output shape, as an integer shape tuple (or list of shape tuples, one tuple per output tensor).

Raises:

- `AttributeError` : if the layer has no defined output shape.
- `RuntimeError` : if called in Eager mode.

output_size

scope_name

state_size

trainable_variables

trainable_weights

updates

variables

Returns the list of all layer variables/weights.

Returns:

A list of variables.

weights

Returns the list of all layer variables/weights.

Returns:

A list of variables.

Methods

`__init__`

```
__init__(
    cell,
    attn_length,
    attn_size=None,
    attn_vec_size=None,
    input_size=None,
    state_is_tuple=True,
    reuse=None
)
```

Create a cell with attention.

Args:

- `cell` : an `RNNCell`, an attention is added to it.

- `attn_length` : integer, the size of an attention window.
- `attn_size` : integer, the size of an attention vector. Equal to `cell.output_size` by default.
- `attn_vec_size` : integer, the number of convolutional features calculated on attention state and a size of the hidden layer built from base cell state. Equal `attn_size` to by default.
- `input_size` : integer, the size of a hidden linear layer, built from inputs and attention. Derived from the input tensor by default.
- `state_is_tuple` : If True, accepted and returned states are n-tuples, where `n = len(cells)` . By default (False), the states are all concatenated along the column axis.
- `reuse` : (optional) Python boolean describing whether to reuse variables in an existing scope. If not `True` , and the existing scope already has the given variables, an error is raised.

Raises:

- `TypeError` : if cell is not an `RNNCell`.
- `ValueError` : if cell returns a state tuple but the flag `state_is_tuple` is `False` or if `attn_length` is zero or less.

`__call__`

```
__call__(
    inputs,
    state,
    scope=None
)
```

Run this RNN cell on inputs, starting from the given state.

Args:

- `inputs` : **2-D** tensor with shape `[batch_size x input_size]` .
- `state` : if `self.state_size` is an integer, this should be a **2-D Tensor** with shape `[batch_size x self.state_size]` . Otherwise, if `self.state_size` is a tuple of integers, this should be a tuple with shapes `[batch_size x s]` for `s in self.state_size` .
- `scope` : `VariableScope` for the created subgraph; defaults to class name.

Returns:

A pair containing:

- Output: A **2-D** tensor with shape `[batch_size x self.output_size]` .
- New state: Either a single **2-D** tensor, or a tuple of tensors matching the arity and shapes of `state` .

`__deepcopy__`

```
__deepcopy__(memo)
```

`add_loss`

```
add_loss(  
    losses,  
    inputs=None  
)
```

Add loss tensor(s), potentially dependent on layer inputs.

Some losses (for instance, activity regularization losses) may be dependent on the inputs passed when calling a layer. Hence, when reusing a same layer on different inputs **a** and **b**, some entries in **layer.losses** may be dependent on **a** and some on **b**. This method automatically keeps track of dependencies.

The **get_losses_for** method allows to retrieve the losses relevant to a specific set of inputs.

Arguments:

- **losses** : Loss tensor, or list/tuple of tensors.
- **inputs** : Optional input tensor(s) that the loss(es) depend on. Must match the **inputs** argument passed to the **__call__** method at the time the losses are created. If **None** is passed, the losses are assumed to be unconditional, and will apply across all dataflows of the layer (e.g. weight regularization losses).

Raises:

- **RuntimeError** : If called in Eager mode.

add_update

```
add_update(  
    updates,  
    inputs=None  
)
```

Add update op(s), potentially dependent on layer inputs.

Weight updates (for instance, the updates of the moving mean and variance in a BatchNormalization layer) may be dependent on the inputs passed when calling a layer. Hence, when reusing a same layer on different inputs **a** and **b**, some entries in **layer.updates** may be dependent on **a** and some on **b**. This method automatically keeps track of dependencies.

The **get_updates_for** method allows to retrieve the updates relevant to a specific set of inputs.

This call is ignored in Eager mode.

Arguments:

- **updates** : Update op, or list/tuple of update ops.
- **inputs** : Optional input tensor(s) that the update(s) depend on. Must match the **inputs** argument passed to the **__call__** method at the time the updates are created. If **None** is passed, the updates are assumed to be unconditional, and will apply across all dataflows of the layer.

add_variable

```
add_variable(  
    name,  
    shape,  
    dtype=None,  
    initializer=None,  
    regularizer=None,  
    trainable=True,  
    constraint=None  
)
```

Adds a new variable to the layer, or gets an existing one; returns it.

Arguments:

- `name` : variable name.
- `shape` : variable shape.
- `dtype` : The type of the variable. Defaults to `self.dtype` or `float32`.
- `initializer` : initializer instance (callable).
- `regularizer` : regularizer instance (callable).
- `trainable` : whether the variable should be part of the layer's "trainable_variables" (e.g. variables, biases) or "non_trainable_variables" (e.g. BatchNorm mean, stddev).
- `constraint` : constraint instance (callable).

Returns:

The created variable.

Raises:

- `RuntimeError` : If called in Eager mode with regularizers.

apply

```
apply(  
    inputs,  
    *args,  
    **kwargs  
)
```

Apply the layer on a input.

This simply wraps `self.__call__`.

Arguments:

- `inputs` : Input tensor(s).
- `*args` : additional positional arguments to be passed to `self.call`.
- `**kwargs` : additional keyword arguments to be passed to `self.call`.

Returns:

Output tensor(s).

build

```
build(_)
```

call

```
call(  
    inputs,  
    state  
)
```

Long short-term memory cell with attention (LSTMA).

count_params

```
count_params()
```

Count the total number of scalars composing the weights.

Returns:

An integer count.

Raises:

- `ValueError` : if the layer isn't yet built (in which case its weights aren't yet defined).

get_input_at

```
get_input_at(node_index)
```

Retrieves the input tensor(s) of a layer at a given node.

Arguments:

- `node_index` : Integer, index of the node from which to retrieve the attribute. E.g. `node_index=0` will correspond to the first time the layer was called.

Returns:

A tensor (or list of tensors if the layer has multiple inputs).

Raises:

- `RuntimeError` : If called in Eager mode.

get_input_shape_at

```
get_input_shape_at(node_index)
```

Retrieves the input shape(s) of a layer at a given node.

Arguments:

- `node_index` : Integer, index of the node from which to retrieve the attribute. E.g. `node_index=0` will correspond to the first time the layer was called.

Returns:

A shape tuple (or list of shape tuples if the layer has multiple inputs).

Raises:

- `RuntimeError` : If called in Eager mode.

get_losses_for

```
get_losses_for(inputs)
```

Retrieves losses relevant to a specific set of inputs.

Arguments:

- `inputs` : Input tensor or list/tuple of input tensors. Must match the `inputs` argument passed to the `__call__` method at the time the losses were created. If you pass `inputs=None`, unconditional losses are returned, such as weight regularization losses.

Returns:

List of loss tensors of the layer that depend on `inputs`.

Raises:

- `RuntimeError` : If called in Eager mode.

get_output_at

```
get_output_at(node_index)
```

Retrieves the output tensor(s) of a layer at a given node.

Arguments:

- `node_index` : Integer, index of the node from which to retrieve the attribute. E.g. `node_index=0` will correspond to the first time the layer was called.

Returns:

A tensor (or list of tensors if the layer has multiple outputs).

Raises:

- `RuntimeError` : If called in Eager mode.

get_output_shape_at

```
get_output_shape_at(node_index)
```

Retrieves the output shape(s) of a layer at a given node.

Arguments:

- `node_index` : Integer, index of the node from which to retrieve the attribute. E.g. `node_index=0` will correspond to the first time the layer was called.

Returns:

A shape tuple (or list of shape tuples if the layer has multiple outputs).

Raises:

- `RuntimeError` : If called in Eager mode.

get_updates_for

```
get_updates_for(inputs)
```

Retrieves updates relevant to a specific set of inputs.

Arguments:

- `inputs` : Input tensor or list/tuple of input tensors. Must match the `inputs` argument passed to the `__call__` method at the time the updates were created. If you pass `inputs=None`, unconditional updates are returned.

Returns:

List of update ops of the layer that depend on `inputs`.

Raises:

- `RuntimeError` : If called in Eager mode.

zero_state

```
zero_state(  
    batch_size,  
    dtype  
)
```

Return zero-filled state tensor(s).

Args:

- `batch_size` : int, float, or unit Tensor representing the batch size.
- `dtype` : the data type to use for the state.

Returns:

If `state_size` is an int or TensorShape, then the return value is a **N-D** tensor of shape `[batch_size x state_size]` filled with zeros.

If `state_size` is a nested list or tuple, then the return value is a nested list or tuple (of the same structure) of **2-D** tensors with the shapes `[batch_size x s]` for each s in `state_size`.

Except as otherwise noted, the content of this page is licensed under the [Creative Commons Attribution 3.0 License](#), and code samples are licensed under the [Apache 2.0 License](#). For details, see our [Site Policies](#). Java is a registered trademark of Oracle and/or its affiliates.

Last updated November 2, 2017.

Stay Connected

Blog

GitHub

Twitter

Support

Issue Tracker

Release Notes

Stack Overflow

English

[Terms](#) | [Privacy](#)