tf.contrib.timeseries.NumpyReader

Contents

Class `NumpyReader`

Defined in `tensorflow/contrib/timeseries/python/timeseries/input_pipeline.py`.

A time series parser for feeding Numpy arrays to a `TimeSeriesInputFn`.

Avoids embedding data in the graph as constants.

Methods

`__init__`

```
__init__(
    data,
    read_num_records_hint=4096
)
```

Numpy array input for a `TimeSeriesInputFn`.

Args:

- `data`: A dictionary mapping feature names to Numpy arrays, with two possible shapes (requires keys `TrainEvalFeatures.TIMES` and `TrainEvalFeatures.VALUES`): Univariate; `TIMES` and `VALUES` are both vectors of shape [series length] Multivariate; `TIMES` is a vector of shape [series length], `VALUES` has shape [series length x number of features]. In any case, `VALUES` and any exogenous features must have their shapes prefixed by the shape of the value corresponding to the `TIMES` key.
- `read_num_records_hint`: The maximum number of samples to read at one time, for efficiency.

`check_dataset_size`

```
check_dataset_size(minimum_dataset_size)
```

Raise an error if the dataset is too small.

`read`

```
read()
```

Returns a large chunk of the Numpy arrays for later re-chunking.

`read_full`

```
read_full()
```

Returns `Tensor` versions of the full Numpy arrays.

Last updated November 2, 2017.