tf.contrib.training.train

```
train(
    train_op,
    logdir,
    master='',
    is_chief=True,
    scaffold=None,
    hooks=None,
    chief_only_hooks=None,
    save_checkpoint_secs=600,
    save_summaries_steps=100,
    config=None
)
```

Defined in `tensorflow/contrib/training/python/training/training.py`.

Runs the training loop.

Args:

- `train_op`: A `Tensor` that, when executed, will apply the gradients and return the loss value.
- `logdir`: The directory where the graph and checkpoints are saved.
- `master`: The URL of the master.
- `is_chief`: Specifies whether or not the training is being run by the primary replica during replica training.
- `scaffold`: An tf.train.Scaffold instance.
- `hooks`: List of `tf.train.SessionRunHook` callbacks which are run inside the training loop.
- `chief_only_hooks`: List of `tf.train.SessionRunHook` instances which are run inside the training loop for the chief trainer only.
- `save_checkpoint_secs`: The frequency, in seconds, that a checkpoint is saved using a default checkpoint saver. If `save_checkpoint_secs` is set to `None`, then the default checkpoint saver isn't used.
- `save_summaries_steps`: The frequency, in number of global steps, that the summaries are written to disk using a default summary saver. If `save_summaries_steps` is set to `None`, then the default summary saver isn't used.
- `config`: An instance of `tf.ConfigProto`.

Returns:

the value of the loss function after training.

Raises:

- `ValueError`: if `logdir` is `None` and either `save_checkpoint_secs` or `save_summaries_steps` are `None.

Last updated November 2, 2017.

Stay Connected

Blog

GitHub

Twitter

Support

Issue Tracker

Release Notes

Stack Overflow

English

Terms | **Privacy**