tf.layers.dropout

```
dropout(
    inputs,
    rate=0.5,
    noise_shape=None,
    seed=None,
    training=False,
    name=None
)
```

Defined in `tensorflow/python/layers/core.py`.

See the guide: Reading data > Multiple input pipelines

Applies Dropout to the input.

Dropout consists in randomly setting a fraction `rate` of input units to 0 at each update during training time, which helps prevent overfitting. The units that are kept are scaled by `1 / (1 - rate)`, so that their sum is unchanged at training time and inference time.

Arguments:

- `inputs` : Tensor input.
- `rate` : The dropout rate, between 0 and 1. E.g. "rate=0.1" would drop out 10% of input units.
- `noise_shape` : 1D tensor of type `int32` representing the shape of the binary dropout mask that will be multiplied with the input. For instance, if your inputs have shape `(batch_size, timesteps, features)`, and you want the dropout mask to be the same for all timesteps, you can use `noise_shape=[batch_size, 1, features]`.
- `seed` : A Python integer. Used to create random seeds. See `tf.set_random_seed` for behavior.
- `training` : Either a Python boolean, or a TensorFlow boolean scalar tensor (e.g. a placeholder). Whether to return the output in training mode (apply dropout) or in inference mode (return the input untouched).
- `name` : The name of the layer (string).

Returns:

Output tensor.

Stay Connected

Blog

GitHub

Twitter

Support

Issue Tracker

Release Notes

Stack Overflow

English

Terms | **Privacy**