tf.nn.fused_batch_norm

```
fused_batch_norm(
    x,
    scale,
    offset,
    mean=None,
    variance=None,
    epsilon=0.001,
    data_format='NHWC',
    is_training=True,
    name=None
)
```

Defined in `tensorflow/python/ops/nn_impl.py`.

See the guide: Neural Network > Normalization

Batch normalization.

As described in http://arxiv.org/abs/1502.03167.

Args:

- `x`: Input `Tensor` of 4 dimensions.
- `scale`: A `Tensor` of 1 dimension for scaling.
- `offset`: A `Tensor` of 1 dimension for bias.
- `mean`: A `Tensor` of 1 dimension for population mean used for inference.
- `variance`: A `Tensor` of 1 dimension for population variance used for inference.
- `epsilon`: A small float number added to the variance of x.
- `data_format`: The data format for x. Either "NHWC" (default) or "NCHW".
- `is_training`: A bool value to specify if the operation is used for training or inference.
- `name`: A name for this operation (optional).

Returns:

- `y`: A 4D Tensor for the normalized, scaled, offsetted x.
- `batch_mean`: A 1D Tensor for the mean of x.
- `batch_var`: A 1D Tensor for the variance of x.

Raises:

- `ValueError`: If mean or variance is not None when is_training is True.

Last updated November 2, 2017.