

tf.reduce_max

```
reduce_max(  
    input_tensor,  
    axis=None,  
    keep_dims=False,  
    name=None,  
    reduction_indices=None  
)
```

Defined in [tensorflow/python/ops/math_ops.py](#).

See the guide: [Math > Reduction](#)

Computes the maximum of elements across dimensions of a tensor.

Reduces `input_tensor` along the dimensions given in `axis`. Unless `keep_dims` is true, the rank of the tensor is reduced by 1 for each entry in `axis`. If `keep_dims` is true, the reduced dimensions are retained with length 1.

If `axis` has no entries, all dimensions are reduced, and a tensor with a single element is returned.

Args:

- `input_tensor`: The tensor to reduce. Should have numeric type.
- `axis`: The dimensions to reduce. If `None` (the default), reduces all dimensions. Must be in the range `[-rank(input_tensor), rank(input_tensor))`.
- `keep_dims`: If true, retains reduced dimensions with length 1.
- `name`: A name for the operation (optional).
- `reduction_indices`: The old (deprecated) name for axis.

Returns:

The reduced tensor.

numpy compatibility

Equivalent to `np.max`

Except as otherwise noted, the content of this page is licensed under the [Creative Commons Attribution 3.0 License](#), and code samples are licensed under the [Apache 2.0 License](#). For details, see our [Site Policies](#). Java is a registered trademark of Oracle and/or its affiliates.

Last updated November 2, 2017.

Stay Connected

[Blog](#)

[GitHub](#)

Twitter

Support

Issue Tracker

Release Notes

Stack Overflow

English

Terms | Privacy