

- Contents
 - Class Session
 - Properties
 - graph
 - graph_def

Defined in `tensorflow/python/client/session.py`.

See the guides: [Running Graphs > Session management](#), [Running Graphs](#)

A class for running TensorFlow operations.

A **Session** object encapsulates the environment in which **Operation** objects are executed, and **Tensor** objects are evaluated. For example:

```
# Build a graph.  
a = tf.constant(5.0)  
b = tf.constant(6.0)  
c = a * b  
  
# Launch the graph in a session.  
sess = tf.Session()  
  
# Evaluate the tensor `c`.  
print(sess.run(c))
```

A session may own resources, such as `tf.Variable`, `tf.QueueBase`, and `tf.ReaderBase`. It is important to release these resources when they are no longer required. To do this, either invoke the `tf.Session.close` method on the session, or use the session as a context manager. The following two examples are equivalent:

```
# Using the `close()` method.
sess = tf.Session()
sess.run(...)
sess.close()

# Using the context manager.
with tf.Session() as sess:
    sess.run(...)
```

The `ConfigProto` protocol buffer exposes various configuration options for a session. For example, to create a session that uses soft constraints for device placement, and log the resulting placement decisions, create a session as follows:

[illegible]

Properties

graph

The graph that was launched in this session.

graph_def

A serializable version of the underlying TensorFlow graph.

Returns:

A `graph_pb2.GraphDef` proto containing nodes for all of the Operations in the underlying TensorFlow graph.

sess_str

Methods

__init__

```
__init__(
    target='',
    graph=None,
    config=None
)
```

Creates a new TensorFlow session.

If no `graph` argument is specified when constructing the session, the default graph will be launched in the session. If you are using more than one graph (created with `tf.Graph()` in the same process, you will have to use different sessions for each graph, but each graph can be used in multiple sessions. In this case, it is often clearer to pass the graph to be launched explicitly to the session constructor.

Args:

- `target`: (Optional.) The execution engine to connect to. Defaults to using an in-process engine. See [Distributed TensorFlow](#) for more examples.
- `graph`: (Optional.) The `Graph` to be launched (described above).
- `config`: (Optional.) A `ConfigProto` protocol buffer with configuration options for the session.

__enter__

```
__enter__()
```

__exit__

```
__exit__(
    exec_type,
    exec_value,
    exec_tb
)
```

as_default

```
as_default()
```

Returns a context manager that makes this object the default session.

Use with the `with` keyword to specify that calls to `tf.Operation.run` or `tf.Tensor.eval` should be executed in this session.

```
c = tf.constant(..)
sess = tf.Session()

with sess.as_default():
    assert tf.get_default_session() is sess
    print(c.eval())
```

To get the current default session, use `tf.get_default_session`.

N.B. The `as_default` context manager *does not* close the session when you exit the context, and you must close the session explicitly.

```
c = tf.constant(...)
sess = tf.Session()
with sess.as_default():
    print(c.eval())
# ...
with sess.as_default():
    print(c.eval())

sess.close()
```

Alternatively, you can use `with tf.Session():` to create a session that is automatically closed on exiting the context, including when an uncaught exception is raised.

N.B. The default session is a property of the current thread. If you create a new thread, and wish to use the default session in that thread, you must explicitly add a `with sess.as_default():` in that thread's function.

N.B. Entering a `with sess.as_default():` block does not affect the current default graph. If you are using multiple graphs, and `sess.graph` is different from the value of `tf.get_default_graph`, you must explicitly enter a `with sess.graph.as_default():` block to make `sess.graph` the default graph.

Returns:

A context manager using this session as the default session.

close

```
close()
```

Closes this session.

Calling this method frees all resources associated with the session.

Raises:

- `tf.errors.OpError` : Or one of its subclasses if an error occurs while closing the TensorFlow session.

list_devices

```
list_devices()
```

Lists available devices in this session.

```
devices = sess.list_devices()
for d in devices:
    print(d.name)
```

Each element in the list has the following properties: - **name** : A string with the full name of the device. ex: `/job:worker/replica:0/task:3/device:CPU:0` - **device_type** : The type of the device (e.g. `CPU`, `GPU`, `TPU`.) - **memory_limit** : The maximum amount of memory available on the device. Note: depending on the device, it is possible the usable memory could be substantially less.

Raises:

- `tf.errors.OpError` : If it encounters an error (e.g. session is in an invalid state, or network errors occur).

Returns:

A list of devices in the session.

make_callable

```
make_callable(
    fetches,
    feed_list=None,
    accept_options=False
)
```

Returns a Python callable that runs a particular step.

The returned callable will take `len(feed_list)` arguments whose types must be compatible feed values for the respective elements of `feed_list`. For example, if element `i` of `feed_list` is a `tf.Tensor`, the `i`th argument to the returned callable must be a numpy ndarray (or something convertible to an ndarray) with matching element type and shape. See [tf.Session.run](#) for details of the allowable feed key and value types.

The returned callable will have the same return type as `tf.Session.run(fetches, ...)`. For example, if `fetches` is a `tf.Tensor`, the callable will return a numpy ndarray; if `fetches` is a `tf.Operation`, it will return `None`.

Args:

- `fetches` : A value or list of values to fetch. See [tf.Session.run](#) for details of the allowable fetch types.
- `feed_list` : (Optional.) A list of `feed_dict` keys. See [tf.Session.run](#) for details of the allowable feed key types.
- `accept_options` : (Optional.) If `True`, the returned `Callable` will be able to accept [tf.RunOptions](#) and [tf.RunMetadata](#) as optional keyword arguments `options` and `run_metadata`, respectively, with the same syntax and semantics as [tf.Session.run](#), which is useful for certain use cases (profiling and debugging) but will result in measurable slowdown of the `Callable`'s performance. Default: `False`.

Returns:

A function that when called will execute the step defined by `feed_list` and `fetches` in this session.

Raises:

- `TypeError`: If `fetches` or `feed_list` cannot be interpreted as arguments to `tf.Session.run`.

`partial_run`

```
partial_run(  
    handle,  
    fetches,  
    feed_dict=None  
)
```

Continues the execution with more feeds and fetches.

This is EXPERIMENTAL and subject to change.

To use partial execution, a user first calls `partial_run_setup()` and then a sequence of `partial_run()`. `partial_run_setup` specifies the list of feeds and fetches that will be used in the subsequent `partial_run` calls.

The optional `feed_dict` argument allows the caller to override the value of tensors in the graph. See `run()` for more information.

Below is a simple example:

```
a = array_ops.placeholder(dtypes.float32, shape=[])  
b = array_ops.placeholder(dtypes.float32, shape=[])  
c = array_ops.placeholder(dtypes.float32, shape=[])  
r1 = math_ops.add(a, b)  
r2 = math_ops.multiply(r1, c)  
  
h = sess.partial_run_setup([r1, r2], [a, b, c])  
res = sess.partial_run(h, r1, feed_dict={a: 1, b: 2})  
res = sess.partial_run(h, r2, feed_dict={c: res})
```

Args:

- `handle`: A handle for a sequence of partial runs.
- `fetches`: A single graph element, a list of graph elements, or a dictionary whose values are graph elements or lists of graph elements (see documentation for `run`).
- `feed_dict`: A dictionary that maps graph elements to values (described above).

Returns:

Either a single value if `fetches` is a single graph element, or a list of values if `fetches` is a list, or a dictionary with the same keys as `fetches` if that is a dictionary (see documentation for `run`).

Raises:

- `tf.errors.OpError`: Or one of its subclasses on error.

`partial_run_setup`

```
partial_run_setup(  
    fetches,  
    feeds=None  
)
```

Sets up a graph with feeds and fetches for partial run.

This is EXPERIMENTAL and subject to change.

Note that contrary to `run`, `feeds` only specifies the graph elements. The tensors will be supplied by the subsequent `partial_run` calls.

Args:

- `fetches`: A single graph element, or a list of graph elements.
- `feeds`: A single graph element, or a list of graph elements.

Returns:

A handle for partial run.

Raises:

- `RuntimeError`: If this `Session` is in an invalid state (e.g. has been closed).
- `TypeError`: If `fetches` or `feed_dict` keys are of an inappropriate type.
- `tf.errors.OpError`: Or one of its subclasses if a TensorFlow error happens.

reset

```
@staticmethod  
reset(  
    target,  
    containers=None,  
    config=None  
)
```

Resets resource containers on `target`, and close all connected sessions.

A resource container is distributed across all workers in the same cluster as `target`. When a resource container on `target` is reset, resources associated with that container will be cleared. In particular, all Variables in the container will become undefined: they lose their values and shapes.

NOTE: (i) `reset()` is currently only implemented for distributed sessions. (ii) Any sessions on the master named by `target` will be closed.

If no resource containers are provided, all containers are reset.

Args:

- `target`: The execution engine to connect to.
- `containers`: A list of resource container name strings, or `None` if all of all the containers are to be reset.
- `config`: (Optional.) Protocol buffer with configuration options.

Raises:

- `tf.errors.OpError` : Or one of its subclasses if an error occurs while resetting containers.

run

```
run(  
    fetches,  
    feed_dict=None,  
    options=None,  
    run_metadata=None  
)
```

Runs operations and evaluates tensors in `fetches`.

This method runs one "step" of TensorFlow computation, by running the necessary graph fragment to execute every `Operation` and evaluate every `Tensor` in `fetches`, substituting the values in `feed_dict` for the corresponding input values.

The `fetches` argument may be a single graph element, or an arbitrarily nested list, tuple, namedtuple, dict, or OrderedDict containing graph elements at its leaves. A graph element can be one of the following types:

- An `tf.Operation`. The corresponding fetched value will be `None`.
- A `tf.Tensor`. The corresponding fetched value will be a numpy ndarray containing the value of that tensor.
- A `tf.SparseTensor`. The corresponding fetched value will be a `tf.SparseTensorValue` containing the value of that sparse tensor.
- A `get_tensor_handle` op. The corresponding fetched value will be a numpy ndarray containing the handle of that tensor.
- A `string` which is the name of a tensor or operation in the graph.

The value returned by `run()` has the same shape as the `fetches` argument, where the leaves are replaced by the corresponding values returned by TensorFlow.

Example:

```
a = tf.constant([10, 20])  
b = tf.constant([1.0, 2.0])  
# 'fetches' can be a singleton  
v = session.run(a)  
# v is the numpy array [10, 20]  
# 'fetches' can be a list.  
v = session.run([a, b])  
# v is a Python list with 2 numpy arrays: the 1-D array [10, 20] and the  
# 1-D array [1.0, 2.0]  
# 'fetches' can be arbitrary lists, tuples, namedtuple, dicts:  
MyData = collections.namedtuple('MyData', ['a', 'b'])  
v = session.run({'k1': MyData(a, b), 'k2': [b, a]})  
# v is a dict with  
# v['k1'] is a MyData namedtuple with 'a' (the numpy array [10, 20]) and  
# 'b' (the numpy array [1.0, 2.0])  
# v['k2'] is a list with the numpy array [1.0, 2.0] and the numpy array  
# [10, 20].
```

The optional `feed_dict` argument allows the caller to override the value of tensors in the graph. Each key in `feed_dict` can be one of the following types:

- If the key is a `tf.Tensor`, the value may be a Python scalar, string, list, or numpy ndarray that can be converted to the same `dtype` as that tensor. Additionally, if the key is a `tf.placeholder`, the shape of the value will be checked

for compatibility with the placeholder.

- If the key is a `tf.SparseTensor`, the value should be a `tf.SparseTensorValue`.
- If the key is a nested tuple of `Tensor`s or `SparseTensor`s, the value should be a nested tuple with the same structure that maps to their corresponding values as above.

Each value in `feed_dict` must be convertible to a numpy array of the dtype of the corresponding key.

The optional `options` argument expects a `[RunOptions]` proto. The options allow controlling the behavior of this particular step (e.g. turning tracing on).

The optional `run_metadata` argument expects a `[RunMetadata]` proto. When appropriate, the non-Tensor output of this step will be collected there. For example, when users turn on tracing in `options`, the profiled info will be collected into this argument and passed back.

Args:

- `fetches`: A single graph element, a list of graph elements, or a dictionary whose values are graph elements or lists of graph elements (described above).
- `feed_dict`: A dictionary that maps graph elements to values (described above).
- `options`: A `[RunOptions]` protocol buffer
- `run_metadata`: A `[RunMetadata]` protocol buffer

Returns:

Either a single value if `fetches` is a single graph element, or a list of values if `fetches` is a list, or a dictionary with the same keys as `fetches` if that is a dictionary (described above).

Raises:

- `RuntimeError`: If this `Session` is in an invalid state (e.g. has been closed).
- `TypeError`: If `fetches` or `feed_dict` keys are of an inappropriate type.
- `ValueError`: If `fetches` or `feed_dict` keys are invalid or refer to a `Tensor` that doesn't exist.

Except as otherwise noted, the content of this page is licensed under the [Creative Commons Attribution 3.0 License](#), and code samples are licensed under the [Apache 2.0 License](#). For details, see our [Site Policies](#). Java is a registered trademark of Oracle and/or its affiliates.

Last updated November 2, 2017.

Stay Connected

[Blog](#)

[GitHub](#)

[Twitter](#)

Support

[Issue Tracker](#)

[Release Notes](#)

[Stack Overflow](#)

