

tf.sparse_tensor_dense_matmul

```
sparse_tensor_dense_matmul(
    sp_a,
    b,
    adjoint_a=False,
    adjoint_b=False,
    name=None
)
```

Defined in [tensorflow/python/ops/sparse_ops.py](#).

See the guide: [Sparse Tensors > Math Operations](#)

Multiply SparseTensor (of rank 2) "A" by dense matrix "B".

No validity checking is performed on the indices of **A**. However, the following input format is recommended for optimal behavior:

- If **adjoint_a == false**: **A** should be sorted in lexicographically increasing order. Use **sparse_reorder** if you're not sure.
- If **adjoint_a == true**: **A** should be sorted in order of increasing dimension 1 (i.e., "column major" order instead of "row major" order).

Using **tf.nn.embedding_lookup_sparse** for sparse multiplication:

It's not obvious but you can consider **embedding_lookup_sparse** as another sparse and dense multiplication. In some situations, you may prefer to use **embedding_lookup_sparse** even though you're not dealing with embeddings.

There are two questions to ask in the decision process: Do you need gradients computed as sparse too? Is your sparse data represented as two **SparseTensor** s: ids and values? There is more explanation about data format below. If you answer any of these questions as yes, consider using **tf.nn.embedding_lookup_sparse**.

Following explains differences between the expected SparseTensors: For example if dense form of your sparse data has shape **[3, 5]** and values:

```
[[ a      ]
 [b      c]
 [      d  ]]
```

SparseTensor format expected by **sparse_tensor_dense_matmul**: **sp_a** (indices, values):

```
[0, 1]: a
[1, 0]: b
[1, 4]: c
[2, 2]: d
```

SparseTensor format expected by **embedding_lookup_sparse**: **sp_ids** **sp_weights**

```
[0, 0]: 1      [0, 0]: a
[1, 0]: 0      [1, 0]: b
[1, 1]: 4      [1, 1]: c
[2, 0]: 2      [2, 0]: d
```

Deciding when to use `sparse_tensor_dense_matmul` vs. `matmul` (`a_is_sparse=True`):

There are a number of questions to ask in the decision process, including:

- Will the SparseTensor `A` fit in memory if densified?
- Is the column count of the product large (> 1)?
- Is the density of `A` larger than approximately 15%?

If the answer to several of these questions is yes, consider converting the `SparseTensor` to a dense one and using `tf.matmul` with `a_is_sparse=True`.

This operation tends to perform well when `A` is more sparse, if the column size of the product is small (e.g. matrix-vector multiplication), if `sp_a.dense_shape` takes on large values.

Below is a rough speed comparison between `sparse_tensor_dense_matmul`, labeled 'sparse', and `matmul` (`a_is_sparse=True`), labeled 'dense'. For purposes of the comparison, the time spent converting from a `SparseTensor` to a dense `Tensor` is not included, so it is overly conservative with respect to the time ratio.

Benchmark system: CPU: Intel Ivybridge with HyperThreading (6 cores) dL1:32KB dL2:256KB dL3:12MB GPU: NVidia Tesla k40c

Compiled with: `-c opt --config=cuda --copt=-mavx`

```
tensorflow/python/sparse_tensor_dense_matmul_op_test --benchmarks
A sparse [m, k] with % nonzero values between 1% and 80%
B dense [k, n]
```

% nnz	n	gpu	m	k	dt(dense)	dt(sparse)	dt(sparse)/dt(dense)
0.01	1	True	100	100	0.000221166	0.00010154	0.459112
0.01	1	True	100	1000	0.00033858	0.000109275	0.322745
0.01	1	True	1000	100	0.000310557	9.85661e-05	0.317385
0.01	1	True	1000	1000	0.0008721	0.000100875	0.115669
0.01	1	False	100	100	0.000208085	0.000107603	0.51711
0.01	1	False	100	1000	0.000327112	9.51118e-05	0.290762
0.01	1	False	1000	100	0.000308222	0.00010345	0.335635
0.01	1	False	1000	1000	0.000865721	0.000101397	0.117124
0.01	10	True	100	100	0.000218522	0.000105537	0.482958
0.01	10	True	100	1000	0.000340882	0.000111641	0.327506
0.01	10	True	1000	100	0.000315472	0.000117376	0.372064
0.01	10	True	1000	1000	0.000905493	0.000123263	0.136128
0.01	10	False	100	100	0.000221529	9.82571e-05	0.44354
0.01	10	False	100	1000	0.000330552	0.000112615	0.340687
0.01	10	False	1000	100	0.000341277	0.000114097	0.334324
0.01	10	False	1000	1000	0.000819944	0.000120982	0.147549
0.01	25	True	100	100	0.000207806	0.000105977	0.509981
0.01	25	True	100	1000	0.000322879	0.00012921	0.400181
0.01	25	True	1000	100	0.00038262	0.00014158	0.370035
0.01	25	True	1000	1000	0.000865438	0.000202083	0.233504
0.01	25	False	100	100	0.000209401	0.000104696	0.499979
0.01	25	False	100	1000	0.000321161	0.000130737	0.407076
0.01	25	False	1000	100	0.000377012	0.000136801	0.362856
0.01	25	False	1000	1000	0.000861125	0.00020272	0.235413
0.2	1	True	100	100	0.000206952	9.69219e-05	0.46833
0.2	1	True	100	1000	0.000348674	0.000147475	0.422959
0.2	1	True	1000	100	0.000336908	0.00010122	0.300439
0.2	1	True	1000	1000	0.001022	0.000203274	0.198898
0.2	1	False	100	100	0.000207532	9.5412e-05	0.459746
0.2	1	False	100	1000	0.000356127	0.000146824	0.41228
0.2	1	False	1000	100	0.000322664	0.000100918	0.312764
0.2	1	False	1000	1000	0.000998987	0.000203442	0.203648
0.2	10	True	100	100	0.000211692	0.000109903	0.519165
0.2	10	True	100	1000	0.000372819	0.000164321	0.440753
0.2	10	True	1000	100	0.000338651	0.000144806	0.427596

0.2	10	True	1000	1000	0.00108312	0.000758876	0.70064
0.2	10	False	100	100	0.000215727	0.000110502	0.512231
0.2	10	False	100	1000	0.000375419	0.0001613	0.429653
0.2	10	False	1000	100	0.000336999	0.000145628	0.432132
0.2	10	False	1000	1000	0.00110502	0.000762043	0.689618
0.2	25	True	100	100	0.000218705	0.000129913	0.594009
0.2	25	True	100	1000	0.000394794	0.00029428	0.745402
0.2	25	True	1000	100	0.000404483	0.0002693	0.665788
0.2	25	True	1000	1000	0.0012002	0.00194494	1.62052
0.2	25	False	100	100	0.000221494	0.0001306	0.589632
0.2	25	False	100	1000	0.000396436	0.000297204	0.74969
0.2	25	False	1000	100	0.000409346	0.000270068	0.659754
0.2	25	False	1000	1000	0.00121051	0.00193737	1.60046
0.5	1	True	100	100	0.000214981	9.82111e-05	0.456836
0.5	1	True	100	1000	0.000415328	0.000223073	0.537101
0.5	1	True	1000	100	0.000358324	0.00011269	0.314492
0.5	1	True	1000	1000	0.00137612	0.000437401	0.317851
0.5	1	False	100	100	0.000224196	0.000101423	0.452386
0.5	1	False	100	1000	0.000400987	0.000223286	0.556841
0.5	1	False	1000	100	0.000368825	0.00011224	0.304318
0.5	1	False	1000	1000	0.00136036	0.000429369	0.31563
0.5	10	True	100	100	0.000222125	0.000112308	0.505608
0.5	10	True	100	1000	0.000461088	0.00032357	0.701753
0.5	10	True	1000	100	0.000394624	0.000225497	0.571422
0.5	10	True	1000	1000	0.00158027	0.00190898	1.20801
0.5	10	False	100	100	0.000232083	0.000114978	0.495418
0.5	10	False	100	1000	0.000454574	0.000324632	0.714146
0.5	10	False	1000	100	0.000379097	0.000227768	0.600817
0.5	10	False	1000	1000	0.00160292	0.00190168	1.18638
0.5	25	True	100	100	0.00023429	0.000151703	0.647501
0.5	25	True	100	1000	0.000497462	0.000598873	1.20386
0.5	25	True	1000	100	0.000460778	0.000557038	1.20891
0.5	25	True	1000	1000	0.00170036	0.00467336	2.74845
0.5	25	False	100	100	0.000228981	0.000155334	0.678371
0.5	25	False	100	1000	0.000496139	0.000620789	1.25124
0.5	25	False	1000	100	0.00045473	0.000551528	1.21287
0.5	25	False	1000	1000	0.00171793	0.00467152	2.71927
0.8	1	True	100	100	0.000222037	0.000105301	0.47425
0.8	1	True	100	1000	0.000410804	0.000329327	0.801664
0.8	1	True	1000	100	0.000349735	0.000131225	0.375212
0.8	1	True	1000	1000	0.00139219	0.000677065	0.48633
0.8	1	False	100	100	0.000214079	0.000107486	0.502085
0.8	1	False	100	1000	0.000413746	0.000323244	0.781261
0.8	1	False	1000	100	0.000348983	0.000131983	0.378193
0.8	1	False	1000	1000	0.00136296	0.000685325	0.50282
0.8	10	True	100	100	0.000229159	0.00011825	0.516017
0.8	10	True	100	1000	0.000498845	0.000532618	1.0677
0.8	10	True	1000	100	0.000383126	0.00029935	0.781336
0.8	10	True	1000	1000	0.00162866	0.00307312	1.88689
0.8	10	False	100	100	0.000230783	0.000124958	0.541452
0.8	10	False	100	1000	0.000493393	0.000550654	1.11606
0.8	10	False	1000	100	0.000377167	0.000298581	0.791642
0.8	10	False	1000	1000	0.00165795	0.00305103	1.84024
0.8	25	True	100	100	0.000233496	0.000175241	0.75051
0.8	25	True	100	1000	0.00055654	0.00102658	1.84458
0.8	25	True	1000	100	0.000463814	0.000783267	1.68875
0.8	25	True	1000	1000	0.00186905	0.00755344	4.04132
0.8	25	False	100	100	0.000240243	0.000175047	0.728625
0.8	25	False	100	1000	0.000578102	0.00104499	1.80763
0.8	25	False	1000	100	0.000485113	0.000776849	1.60138
0.8	25	False	1000	1000	0.00211448	0.00752736	3.55992

Args:

- `sp_a` : SparseTensor A, of rank 2.
- `b` : A dense Matrix with the same dtype as `sp_a`.
- `adjoint_a` : Use the adjoint of A in the matrix multiply. If A is complex, this is `transpose(conj(A))`. Otherwise it's `transpose(A)`.
- `adjoint_b` : Use the adjoint of B in the matrix multiply. If B is complex, this is `transpose(conj(B))`. Otherwise it's `transpose(B)`.
- `name` : A name prefix for the returned tensors (optional)

Returns:

A dense matrix (pseudo-code in dense np.matrix notation): `A = A.H if adjoint_a else A` `B = B.H if adjoint_b else B` `return A*B`

Except as otherwise noted, the content of this page is licensed under the [Creative Commons Attribution 3.0 License](#), and code samples are licensed under the [Apache 2.0 License](#). For details, see our [Site Policies](#). Java is a registered trademark of Oracle and/or its affiliates.

Last updated November 2, 2017.

Stay Connected

[Blog](#)

[GitHub](#)

[Twitter](#)

Support

[Issue Tracker](#)

[Release Notes](#)

[Stack Overflow](#)

English

[Terms](#) | [Privacy](#)