

## tf.train.FtrlOptimizer

## Contents

Class FtrlOptimizer

## Methods

`__init__``apply_gradients`Class **FtrlOptimizer**Inherits From: [Optimizer](#)Defined in [tensorflow/python/training/ftrl.py](#).See the guide: [Training > Optimizers](#)

Optimizer that implements the FTRL algorithm.

See this [paper](#). This version has support for both online L2 (the L2 penalty given in the paper above) and shrinkage-type L2 (which is the addition of an L2 penalty to the loss function).

## Methods

**`__init__`**

```
__init__(
    learning_rate,
    learning_rate_power=-0.5,
    initial_accumulator_value=0.1,
    l1_regularization_strength=0.0,
    l2_regularization_strength=0.0,
    use_locking=False,
    name='Ftrl',
    accum_name=None,
    linear_name=None,
    l2_shrinkage_regularization_strength=0.0
)
```

Construct a new FTRL optimizer.

## Args:

- `learning_rate`: A float value or a constant float [Tensor](#).
- `learning_rate_power`: A float value, must be less or equal to zero.
- `initial_accumulator_value`: The starting value for accumulators. Only positive values are allowed.
- `l1_regularization_strength`: A float value, must be greater than or equal to zero.

- `l2_regularization_strength` : A float value, must be greater than or equal to zero.
- `use_locking` : If `True` use locks for update operations.
- `name` : Optional name prefix for the operations created when applying gradients. Defaults to "Ftrl".
- `accum_name` : The suffix for the variable that keeps the gradient squared accumulator. If not present, defaults to `name`.
- `linear_name` : The suffix for the variable that keeps the linear gradient accumulator. If not present, defaults to `name + "_1"`.
- `l2_shrinkage_regularization_strength` : A float value, must be greater than or equal to zero. This differs from L2 above in that the L2 above is a stabilization penalty, whereas this L2 shrinkage is a magnitude penalty. The FTRL formulation can be written as:  $w_{t+1} = \text{argmin}_w (\hat{g}_{1:t} w + L1 \|w\|_1 + L2 \|w\|_2^2)$ , where  $\hat{g} = g + (2L2\_shrinkage w)$ , and  $g$  is the gradient of the loss function w.r.t. the weights  $w$ . Specifically, in the absence of L1 regularization, it is equivalent to the following update rule:  $w_{t+1} = w_t - lr_t / (1 + 2L2lr_t) * g_t - 2L2\_shrinkage lr_t / (1 + 2L2lr_t) * w_t$  where  $lr_t$  is the learning rate at  $t$ . When input is sparse shrinkage will only happen on the active weights.

Raises:

- `ValueError` : If one of the arguments is invalid.

## apply\_gradients

```
apply_gradients(
    grads_and_vars,
    global_step=None,
    name=None
)
```

Apply gradients to variables.

This is the second part of `minimize()` . It returns an `Operation` that applies gradients.

Args:

- `grads_and_vars` : List of (gradient, variable) pairs as returned by `compute_gradients()` .
- `global_step` : Optional `Variable` to increment by one after the variables have been updated.
- `name` : Optional name for the returned operation. Default to the name passed to the `Optimizer` constructor.

Returns:

An `Operation` that applies the specified gradients. If `global_step` was not None, that operation also increments `global_step` .

Raises:

- `TypeError` : If `grads_and_vars` is malformed.
- `ValueError` : If none of the variables have gradients.

## compute\_gradients

```
compute_gradients(
    loss,
    var_list=None,
    gate_gradients=GATE_OP,
    aggregation_method=None,
    colocate_gradients_with_ops=False,
    grad_loss=None
)
```

Compute gradients of `loss` for the variables in `var_list`.

This is the first part of `minimize()`. It returns a list of (gradient, variable) pairs where "gradient" is the gradient for "variable". Note that "gradient" can be a `Tensor`, an `IndexedSlices`, or `None` if there is no gradient for the given variable.

Args:

- `loss`: A `Tensor` containing the value to minimize.
- `var_list`: Optional list or tuple of `tf.Variable` to update to minimize `loss`. Defaults to the list of variables collected in the graph under the key `GraphKey.TRAINABLE_VARIABLES`.
- `gate_gradients`: How to gate the computation of gradients. Can be `GATE_NONE`, `GATE_OP`, or `GATE_GRAPH`.
- `aggregation_method`: Specifies the method used to combine gradient terms. Valid values are defined in the class `AggregationMethod`.
- `colocate_gradients_with_ops`: If True, try colocating gradients with the corresponding op.
- `grad_loss`: Optional. A `Tensor` holding the gradient computed for `loss`.

Returns:

A list of (gradient, variable) pairs. Variable is always present, but gradient can be `None`.

Raises:

- `TypeError`: If `var_list` contains anything else than `Variable` objects.
- `ValueError`: If some arguments are invalid.

## get\_name

```
get_name()
```

## get\_slot

```
get_slot(
    var,
    name
)
```

Return a slot named `name` created for `var` by the Optimizer.

Some `Optimizer` subclasses use additional variables. For example `Momentum` and `Adagrad` use variables to accumulate updates. This method gives access to these `Variable` objects if for some reason you need them.

Use `get_slot_names()` to get the list of slot names created by the `Optimizer`.

Args:

- `var` : A variable passed to `minimize()` or `apply_gradients()` .
- `name` : A string.

Returns:

The `Variable` for the slot if it was created, `None` otherwise.

## get\_slot\_names

```
get_slot_names()
```

Return a list of the names of slots created by the `Optimizer` .

See `get_slot()` .

Returns:

A list of strings.

## minimize

```
minimize(  
    loss,  
    global_step=None,  
    var_list=None,  
    gate_gradients=GATE_OP,  
    aggregation_method=None,  
    colocate_gradients_with_ops=False,  
    name=None,  
    grad_loss=None  
)
```

Add operations to minimize `loss` by updating `var_list` .

This method simply combines calls `compute_gradients()` and `apply_gradients()` . If you want to process the gradient before applying them call `compute_gradients()` and `apply_gradients()` explicitly instead of using this function.

Args:

- `loss` : A `Tensor` containing the value to minimize.
- `global_step` : Optional `Variable` to increment by one after the variables have been updated.
- `var_list` : Optional list or tuple of `Variable` objects to update to minimize `loss` . Defaults to the list of variables collected in the graph under the key `GraphKeys.TRAINABLE_VARIABLES` .
- `gate_gradients` : How to gate the computation of gradients. Can be `GATE_NONE` , `GATE_OP` , or `GATE_GRAPH` .
- `aggregation_method` : Specifies the method used to combine gradient terms. Valid values are defined in the class `AggregationMethod` .
- `colocate_gradients_with_ops` : If True, try colocating gradients with the corresponding op.
- `name` : Optional name for the returned operation.
- `grad_loss` : Optional. A `Tensor` holding the gradient computed for `loss` .

Returns:

An Operation that updates the variables in `var_list` . If `global_step` was not `None` , that operation also increments `global_step` .

Raises:

- `ValueError` : If some of the variables are not `Variable` objects.

## Class Members

---

**GATE\_GRAPH**

**GATE\_NONE**

**GATE\_OP**

---

*Except as otherwise noted, the content of this page is licensed under the [Creative Commons Attribution 3.0 License](#), and code samples are licensed under the [Apache 2.0 License](#). For details, see our [Site Policies](#). Java is a registered trademark of Oracle and/or its affiliates.*

*Last updated November 2, 2017.*

### Stay Connected

Blog

GitHub

Twitter

### Support

Issue Tracker

Release Notes

Stack Overflow

English

[Terms](#) | [Privacy](#)